

COLLECTIVE VARIABLES MODULE

Reference manual for LAMMPS

Code version: 2022-05-24

Updated versions of this manual: [[GROMACS](#)] [[LAMMPS](#)] [[NAMD](#)] [[VMD](#)]



Alejandro Bernardin, Haochuan Chen, Jeffrey R. Comer, Giacomo Fiorin, Haohao Fu, Jérôme Hénin, Axel Kohlmeyer, Fabrizio Marinelli, Joshua V. Vermaas, Andrew D. White
([HTML version](#))

Contents

1	Overview	4
2	Writing a Colvars configuration: a crash course	5
3	Enabling and controlling the Colvars module in LAMMPS	7
3.1	Units in the Colvars module	7
3.2	LAMMPS keywords	7
3.3	Configuration syntax used by the Colvars module	8
3.4	Global keywords	9
3.5	Input state file	10
3.6	Output files	11
3.7	File formats	11
3.7.1	Configuration and state files.	11
3.7.2	Index (NDX) files	11
3.7.3	XYZ coordinate files	12
3.7.4	Grid files: multicolumn text format	12
3.7.5	Output trajectory files	14
4	Defining collective variables	16
4.1	Choosing a function	16
4.2	Distances	18
4.2.1	distance: center-of-mass distance between two groups.	18
4.2.2	distanceZ: projection of a distance vector on an axis.	19
4.2.3	distanceXY: modulus of the projection of a distance vector on a plane.	20
4.2.4	distanceVec: distance vector between two groups.	20
4.2.5	distanceDir: distance unit vector between two groups.	20
4.2.6	distanceInv: mean distance between two groups of atoms.	21
4.3	Angles	21
4.3.1	angle: angle between three groups.	22
4.3.2	dipoleAngle: angle between two groups and dipole of a third group.	22
4.3.3	dihedral: torsional angle between four groups.	22
4.3.4	polarTheta: polar angle in spherical coordinates.	23
4.3.5	polarPhi: azimuthal angle in spherical coordinates.	23
4.4	Contacts	24
4.4.1	coordNum: coordination number between two groups.	24
4.4.2	selfCoordNum: coordination number between atoms within a group.	25
4.4.3	hBond: hydrogen bond between two atoms.	26
4.5	Collective metrics	26
4.5.1	rmsd: root mean square displacement (RMSD) from reference positions.	27
4.5.2	Advanced usage of the rmsd component.	28
4.5.3	eigenvector: projection of the atomic coordinates on a vector.	28
4.5.4	gyration: radius of gyration of a group of atoms.	30
4.5.5	inertia: total moment of inertia of a group of atoms.	30
4.5.6	dipoleMagnitude: dipole magnitude of a group of atoms.	30
4.5.7	inertiaZ: total moment of inertia of a group of atoms around a chosen axis.	30
4.6	Rotations	31

4.6.1	orientation: orientation from reference coordinates.	32
4.6.2	orientationAngle: angle of rotation from reference coordinates.	33
4.6.3	orientationProj: cosine of the angle of rotation from reference coordinates.	33
4.6.4	spinAngle: angle of rotation around a given axis.	33
4.6.5	tilt: cosine of the rotation orthogonal to a given axis.	34
4.6.6	eulerPhi: Roll angle from references coordinates.	34
4.6.7	eulerTheta: Pitch angle from references coordinates.	35
4.6.8	eulerPsi: Yaw angle from references coordinates.	36
4.7	Raw data: building blocks for custom functions	36
4.7.1	cartesian: vector of atomic Cartesian coordinates.	36
4.7.2	distancePairs: set of pairwise distances between two groups.	36
4.8	Geometric path collective variables	37
4.8.1	gspath: progress along a path defined in atomic Cartesian coordinate space.	37
4.8.2	gzpath: distance from a path defined in atomic Cartesian coordinate space.	38
4.8.3	linearCombination: Helper CV to define a linear combination of other CVs	39
4.8.4	customColvar: Helper CV to define a mathematical expression as CV from other CVs	39
4.8.5	gspathCV: progress along a path defined in CV space.	40
4.8.6	gzpathCV: distance from a path defined in CV space.	40
4.9	Arithmetic path collective variables	42
4.9.1	aspathCV: progress along a path defined in CV space.	42
4.9.2	azpathCV: distance from a path defined in CV space.	43
4.10	Dense neural network in CV space	44
4.11	Shared keywords for all components	46
4.12	Periodic components	46
4.13	Non-scalar components	47
4.13.1	Calculating total forces	48
4.14	Linear and polynomial combinations of components	48
4.15	Custom functions	49
4.16	Defining grid parameters	50
4.17	Trajectory output	52
4.18	Extended Lagrangian	53
4.19	Multiple time-step variables	54
4.20	Backward-compatibility	54
4.21	Statistical analysis	55
5	Selecting atoms	57
5.1	Atom selection keywords	57
5.2	Moving frame of reference.	58
5.3	Treatment of periodic boundary conditions.	60
5.4	Performance of a Colvars calculation based on group size.	61
6	Biasing and analysis methods	62
6.1	Thermodynamic integration	63
6.2	Adaptive Biasing Force	64
6.2.1	ABF requirements on collective variables	65
6.2.2	Parameters for ABF	66
6.2.3	Output files	68

6.2.4	Multidimensional free energy surfaces	68
6.3	Extended-system Adaptive Biasing Force (eABF)	70
6.3.1	CZAR estimator of the free energy	71
6.4	Metadynamics	72
6.4.1	Treatment of the PMF boundaries	73
6.4.2	Basic configuration keywords	74
6.4.3	Output files	75
6.4.4	Performance optimization	76
6.4.5	Ensemble-Biased Metadynamics	77
6.4.6	Well-tempered metadynamics	79
6.4.7	Multiple-walker metadynamics	79
6.5	Harmonic restraints	81
6.5.1	Moving restraints: steered molecular dynamics	83
6.5.2	Moving restraints: umbrella sampling	84
6.5.3	Changing force constant	84
6.6	Computing the work of a changing restraint	85
6.7	Harmonic wall restraints	85
6.8	Linear restraints	88
6.9	Adaptive Linear Bias/Experiment Directed Simulation	89
6.10	Multidimensional histograms	91
6.10.1	Defining grids for multidimensional histograms	92
6.10.2	Output options for multi-dimensional histograms	92
6.10.3	Histogramming vector variables	93
6.11	Probability distribution-restraints	94
7	Syntax changes from older versions	96
8	Compilation notes	98

1 Overview

In molecular dynamics simulations, it is often useful to reduce the large number of degrees of freedom of a physical system into few parameters whose statistical distributions can be analyzed individually, or used to define biasing potentials to alter the dynamics of the system in a controlled manner. These have been called ‘order parameters’, ‘collective variables’, ‘(surrogate) reaction coordinates’, and many other terms.

Here we use primarily the term ‘collective variable’, often shortened to *colvar*, to indicate any differentiable function of atomic Cartesian coordinates, x_i , with i between 1 and N , the total number of atoms:

$$\xi(t) = \xi(X(t)) = \xi(x_i(t), x_j(t), x_k(t), \dots), \quad 1 \leq i, j, k \dots \leq N \quad (1)$$

This manual documents the collective variables module (**Colvars**), a software that provides an implementation for the functions $\xi(X)$ with a focus on flexibility, robustness and high performance. The module is designed to perform multiple tasks concurrently during or after a simulation, the most common of which are:

- apply restraints or biasing potentials to multiple variables, tailored on the system by choosing from a wide set of basis functions, without limitations on their number or on the number of atoms involved;
- calculate potentials of mean force (PMFs) along any set of variables, using different enhanced sampling methods, such as Adaptive Biasing Force (ABF), metadynamics, steered MD and umbrella sampling; variants of these methods that make use of an ensemble of replicas are supported as well;
- calculate statistical properties of the variables, such as running averages and standard deviations, correlation functions of pairs of variables, and multidimensional histograms: this can be done either at run-time without the need to save very large trajectory files, or after a simulation has been completed (post-processing).

Detailed explanations of the design of the Colvars module are provided in reference [1]. Please cite this reference whenever publishing work that makes use of this module, alongside any other publications for specific features being, according to the usage summary printed when running a Colvars-enabled MD simulation or analysis.

2 Writing a Colvars configuration: a crash course

The Colvars configuration is a plain text file or string that defines collective variables, biases, and general parameters of the Colvars module. It is passed to the module using back-end-specific commands documented in section 3.

Example: steering two groups of atoms away from each other. Now let us look at a complete, non-trivial configuration. Suppose that we want to run a steered MD experiment where a small molecule is pulled away from a protein binding site. In Colvars terms, this is done by applying a moving restraint to the distance between the two objects. The configuration will contain two blocks, one defining the distance variable (see section 4 and 4.2.1), and the other the moving harmonic restraint (6.5).

```
indexFile index.ndx
colvar {
  name dist
  distance {
    group1 { atomNumbersRange 42-55 }
    group2 { indexGroup C-alpha_15-30 }
  }
}

harmonic {
  colvars dist
  forceConstant 20.0
  centers 4.0      # initial distance
  targetCenters 15.0 # final distance
  targetNumSteps 500000
}
```

Reading this input in plain English: the variable here named *dist* consists in a distance function between the centers of two groups: the ligand (atoms 42 to 55) and the α -carbon atoms of residues 15 to 30 in the protein (selected from the index group “C-alpha_15-30”). To the “*dist*” variable, we apply a harmonic potential of force constant 20 *energy unit/length unit*², initially centered around a value of 4 *length unit*, which will increase to 15 *length unit* over 500,000 simulation steps.

The atom selection keywords are detailed in section 5. If the selection is too complex to implement only via internal keywords, an external index file may be created following the NDX format used in GROMACS (see 3.7.2) or by using the `group2ndx` LAMMPS command.

Example: using multiple variables and multiple biasing/analysis methods together. A more complex example configuration is included below, showing how a variable may be constructed by combining multiple existing functions, and how multiple variables or multiple biases may be used concurrently. The colvar indicated below as “*d*” is defined as the difference between two distances (see 4.2): the first distance (d_1) is taken between the center of mass of atoms 1 and 2 and that of atoms 3 to 5, the second (d_2) between atom 7 and the center of mass of atoms 8 to 10 (see 5). The difference $d = d_1 - d_2$ is obtained by multiplying the two by a coefficient $C = +1$ or $C = -1$, respectively (see 4.14). The colvar called “*c*” is the coordination

number calculated between atoms 1 to 10 and atoms 11 to 20. A harmonic restraint (see 6.5) is applied to both d and c : to allow using the same force constant K , both d and c are scaled by their respective fluctuation widths w_d and w_c . The values of “ c ” are also recorded throughout the simulation as a joint 2-dimensional histogram (see 6.10).

```
colvar {
  # difference of two distances
  name d
  width 0.2 # estimated fluctuation width
  distance {
    componentCoeff 1.0
    group1 { atomNumbers 1 2 }
    group2 { atomNumbers 3 4 5 }
  }
  distance {
    componentCoeff -1.0
    group1 { atomNumbers 7 }
    group2 { atomNumbers 8 9 10 }
  }
}

colvar {
  name c
  coordNum {
    cutoff 6.0
    group1 { atomNumbersRange 1-10 }
    group2 { atomNumbersRange 11-20 }
    tolerance 1.0e-6
    pairListFrequency 1000
  }
}

harmonic {
  colvars d c
  centers 3.0 4.0
  forceConstant 5.0
}

histogram {
  colvars c
}
```

3 Enabling and controlling the Colvars module in LAMMPS

Here, we document the syntax of the commands and parameters used to set up and use the Colvars module in LAMMPS [2]. One of these parameters is the configuration file or the configuration text for the module itself, whose syntax is described in 3.3 and in the following sections.

3.1 Units in the Colvars module

The “internal units” of the Colvars module are the units in which values are expected to be in the configuration file, and in which collective variable values, energies, etc. are expressed in the output and colvars trajectory files. Generally **the Colvars module uses internally the same units as its back-end MD engine, with the exception of VMD**, where different unit sets are supported to allow for easy setup, visualization and analysis of Colvars simulations performed with any simulation engine.

Note that **angles** are expressed in degrees, and derived quantities such as force constants are based on degrees as well. Some colvar components have default values, expressed in Ångström in this documentation. They are converted to the current length unit as needed. Atomic coordinates read from **XYZ files** (and PDB files where applicable) are expected to be expressed in Ångström, no matter what unit system is in use by the back-end (LAMMPS) or the Colvars Module. They are converted internally to the current length unit as needed. Note that force constants in harmonic and harmonicWalls biases (6.5) are rescaled according to the **width** parameter of colvars, so that they are formally in energy units, although if **width** is given its default value of 1.0, force constants are effectively expressed in $\text{energy unit}/(\text{colvar unit})^2$.

To avoid errors due to reading configuration files written in a different unit system, it can be specified within the input:

- **Keyword units** <Unit system to be used>

Context: global

Acceptable values: string

Description: A string defining the units to be used internally by Colvars. Supported values are the following LAMMPS unit styles: *real* (Å, kcal/mol), *metal* (Å, eV), and *electron* (Bohr, Hartree). However only the unit system currently in use by LAMMPS is accepted: other values will trigger an error.

3.2 LAMMPS keywords

To enable a Colvars-based calculation, the following line must be added to the LAMMPS configuration file:

```
fix Colvars all colvars configfile [keyword value pairs ...]
```

where the fix ID is here set to the string “Colvars”, because there can only be one instance of `fix colvars` in a LAMMPS run. The value of *configfile* is the name of the configuration file for the Colvars module. Other optional keywords may follow:

- **Keyword input** \langle Name or prefix of the input state file \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: string
Description: If a value is provided, it is interpreted as either the name of the input state file, or as the prefix of the file named `input.colvars.state`. This file contains information needed to continue a previous collective variables-based calculation, including the number of the last computed step (useful for time-dependent biases). The same information is also stored in the binary restart files written by LAMMPS, so this option is not needed when continuing a calculation from a LAMMPS restart.
- **Keyword output** \langle Prefix of the output state file \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: string
Default value: “out”
Description: If a value is provided, it is interpreted as the prefix to all output files that will be written by the Colvars module (see 3.6).
- **Keyword unwrap** \langle Whether to unwrap coordinates passed to the Colvars module \rangle
Context: keyword of the `fix colvars` command
Acceptable values: “yes” or “no”
Default value: “yes”
Description: This keyword controls whether wrapped or unwrapped coordinates are passed to the Colvars module for calculation of the collective variables and of the resulting forces. The default is to use the image flags to reconstruct the absolute atom positions: under this convention, centers of mass and centers of geometry are calculated as a weighted vector sum (see 5.3). Setting this to *no* will use the current local coordinates that are wrapped back into the simulation cell at each re-neighboring instead.
- **Keyword seed** \langle Seed for the random number generator \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: positive integer
Default value: 1966
Description: If defined, the value of this keyword is provided as seed to the random number generator. This is only meaningful when the `extendedLangevinDamping` keyword is used (see 4.18).
- **Keyword tstat** \langle Thermostating fix \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: string
Default value: NULL
Description: This keyword provides the *ID* of an applicable thermostating fix command. This will be used to provide the Colvars module with the current thermostat target temperature when using a method that needs this information.

3.3 Configuration syntax used by the Colvars module

Configuration for the Colvars module is passed using an external file. Configuration lines follow the format “keyword value” or “keyword { ... }”, where the keyword and its value must be separated by one or more space characters. The following formatting rules apply:

- *Keywords are case-insensitive*; for example, `upperBoundary` is the same as `upperboundary` and `UPPERBOUNDARY`; note that their string values are however still case-sensitive (e.g. names of variables, file names).
- A long value, or a list of multiple values, can be distributed across multiple lines by using *curly braces*, “{” and “}”: the opening brace “{” must occur on the same line as the keyword, following at least one space character; the closing brace “}” may be at any position after that; any keywords following the closing brace on the same line are not valid (they should appear instead on a different line).
- Many keywords are nested, and are only meaningful within the specific context of another keyword; for example, the keyword `name` is available inside the block of the keyword `colvar { . . . }`, but not outside of it; for every keyword documented in the following, the “*parent*” keyword that defines such context is also indicated.
- If a keyword requiring a boolean value (`yes|on|true` or `no|off|false`) is provided without an explicit value, it defaults to ‘`yes|on|true`’; for example, ‘`outputAppliedForce`’ may be used as shorthand for ‘`outputAppliedForce on`’.
- The hash character “#” indicates a *comment*: all text in the same line following this character will be ignored.
- Outside of comments, *only ASCII characters* are allowed for defining keywords, and the only white-space characters supported are spaces, tabs and newlines: a warning will be printed upon detection of non-ASCII characters in a configuration line, which include both characters that are visibly “special”, as well as those with a very similar appearance to ASCII ones (for instance, **non-breaking spaces**); common ways to identify/remove non-ASCII characters are using the Emacs text editor, or using `LC_ALL=C vi`.

3.4 Global keywords

The following keywords are available in the global context of the Colvars configuration, i.e. they are not nested inside other keywords:

- **Keyword** `colvarsTrajFrequency` ⟨Colvar value trajectory frequency⟩
Context: global
Acceptable values: positive integer
Default value: 100
Description: The values of each colvar (and of other related quantities, if requested) are written to the file `output.colvars.traj` every these many steps throughout the simulation. If the value is 0, such trajectory file is not written. For optimization the output is buffered, and synchronized with the disk only when the restart file is being written.
- **Keyword** `colvarsRestartFrequency` ⟨Colvar module restart frequency⟩
Context: global
Acceptable values: positive integer
Default value: LAMMPS restart frequency when defined
Description: The state file and any other output files produced by Colvars are written every these many steps (the trajectory file is still written every **colvarsTrajFrequency** steps). It is generally

a good idea to leave this parameter at its default value, unless needed for special cases or to disable automatic writing of output files altogether. Writing can still be invoked at any time via the command `cv save`.

- **Keyword** `indexFile` ⟨Index file for atom selection (GROMACS “ndx” format)⟩

Context: global

Acceptable values: UNIX filename

Description: This option reads an index file (usually with a `.ndx` extension) as produced by the `make_ndx` tool of GROMACS. This keyword may be repeated to load multiple index files. A group with the same name may appear multiple times, as long as it contains the same indices in identical order each time: an error is raised otherwise. In LAMMPS, the `group2ndx` command can be used to generate such file from existing groups. Note that the Colvars module reads the indices of atoms from the index file: therefore, the LAMMPS groups do not need to remain active during the simulation, and could be deleted right after issuing `group2ndx`. The names of index groups contained in this file can then be used to define atom groups with the `indexGroup` keyword. Other supported methods to select atoms are described in 5.

- **Keyword** `smp` ⟨Whether SMP parallelism should be used⟩

Context: global

Acceptable values: boolean

Default value: on

Description: If this flag is enabled (default), SMP parallelism over threads will be used to compute variables and biases, provided that this is supported by the LAMMPS build in use.

3.5 Input state file

Because many of the methods implemented in Colvars are history-dependent, a *state file* is often needed to continue a long simulation over consecutive runs. Such state file is written automatically at the end of any simulation with Colvars, and contains data accumulated during that simulation along with the step number at the end of it. The step number read from the state file is then used to control such time-dependent biases: because of this essential role, the step number internal to Colvars may not always match the step number reported by the MD program that carried during the simulation (which may instead restart from zero each time).

Depending on the configuration, a state file may need to be loaded issued at the beginning of a new simulation when time-dependent biasing methods are applied (moving restraints, metadynamics, ABF, ...). When the Colvars fix is defined in LAMMPS, the keyword `input` can be used to load the state file, although it is typically easier to use the LAMMPS `read_restart` to re-initialize Colvars together with other fixes. After initialization, a state file may be loaded at any time with the `fix_modify <fix-ID> load <filename>` LAMMPS command.

It is possible to load a state file even if the configuration has changed: for example, new variables may be defined or restraints be added in between consecutive runs. For each newly defined variable or bias, no information will be read from the state file if this is unavailable: such new objects will remain uninitialized until the first compute step. Conversely, any information that the state file has about variables or biases that are not defined any longer is silently ignored. *Because these checks are done by the names of variables or biases, it is the user's responsibility to ensure that these are consistent between runs.*

3.6 Output files

During a simulation with collective variables defined, the following three output files are written:

- A *state file*, named `output.colvars.state`; this file is in ASCII (plain text) format. This file is written at the end of the specified run.
This is the only Colvars output file needed to continue a simulation.
- If the parameter `colvarsRestartFrequency` is larger than zero, a *restart file* is written every that many steps: this file is fully equivalent to the final state file. The name of this file is `restart.colvars.state`.
- If the parameter `colvarsTrajFrequency` is greater than 0 (default value: 100 steps), a *trajectory file*, named `output.colvars.traj`, is written during the simulation. Unlike the state file, this file is not needed to restart a simulation, but can be used for post-processing and analysis. The format of this file is described in sec. 3.7.5.
- Additionally, certain features, when enabled, can emit output files with a specific purpose: for example, potentials of mean force can be written to file to be analyzed or plotted. These files are described in the respective sections, but as a general rule they all use names beginning with the *output* prefix.

Other output files may also be written by specific methods, e.g. the ABF or metadynamics methods (6.2, 6.4). Like the trajectory file, they are needed only for analyzing, not continuing a simulation. All such files' names also begin with the prefix *output*.

3.7 File formats

3.7.1 Configuration and state files.

Configuration files are text files that are generally read as input by LAMMPS. Starting from version 2017-02-01, changes in newline encodings are handled transparently, i.e. it is possible to typeset a configuration file in Windows (CR-LF newlines) and then use it with Linux or macOS (LF-only newlines).

State files, although not written manually, follow otherwise the same text format as configuration files.

3.7.2 Index (NDX) files

For atom selections that cannot be specified only by using internal Colvars keywords, external *index files* may also be used following the **NDX format** used in GROMACS:

```
[ group_1_name ]
  i1 i2 i3 i4 ...
  ...           ... iN
```

```
[ group_2_name ]  
...
```

where iI through iN are 1-based indices. Each group name may not contain spaces or tabs: otherwise, a parsing error will be raised.

Multiple index files may be provided to Colvars, each using the keyword `indexFile`. Two index files may contain groups with the same names, however these must also represent identical atom selections, i.e. the same sequence of indices including order.

Other than with GROMACS, an index group may also be generated from the VMD command-line interface, using the helper function `write_index_group` provided in the `colvartools` folder:

```
source colvartools/write_index_group.tcl  
set sel [atomselect top "resname XXX and not hydrogen"]  
write_index_group indexfile.ndx $sel "Ligand"
```

In LAMMPS, NDX files may also be generated from internal groups via the `group2ndx` command.

3.7.3 XYZ coordinate files

XYZ coordinate files are text files with the extension “.xyz”. They are read by the Colvars module using an internal reader, and expect the following format:

```
 $N$   
Comment line  
 $E_1$        $x_1$    $y_1$    $z_1$   
 $E_2$        $x_2$    $y_2$    $z_2$   
...  
 $E_N$        $x_N$    $y_N$    $z_N$ 
```

where N is the number of atomic coordinates in the file and E_i is the chemical element of the i -th atom. Because E_i is not used in Colvars, any string that does not contain tabs or spaces is acceptable.

An XYZ file may contain either one of the following scenarios:

- (i) The file contains as many coordinates as the atoms that they are being read for: all coordinates will be read from the file following the same order as the atoms appear in the selection generated using the keywords listed in section 5. (Note that the order is guaranteed only if a single type of selection keyword is used one or more times, and not guaranteed when different types of selection keywords are used.)
- (ii) The file contains more coordinates than needed, and it is assumed to contain coordinates for the entire system: only coordinates that match the numeric indices of the selected atoms are read, in order of increasing number.

XYZ-file coordinates are read directly by Colvars and stored internally as double-precision floating point numbers.

3.7.4 Grid files: multicolumn text format

Many simulation methods and analysis tools write files that contain functions of the collective variables tabulated on a grid (e.g. potentials of mean force or multidimensional histograms) for the purpose of analyzing results. Such files are produced by ABF (6.2), metadynamics (6.4), multidimensional histograms (6.10), as well as any restraint with optional thermodynamic integration support (6.1).

In some cases, these files may also be read as input of a new simulation. Suitable input files for this purpose are typically generated as output files of previous simulations, or directly by the user in the specific case of ensemble-biased metadynamics (6.4.5). This section explains the “multicolumn” format used by these files. For a multidimensional function $f(\xi_1, \xi_2, \dots)$ the multicolumn grid format is defined as follows:

```
# Ncv
# min( $\xi_1$ )    width( $\xi_1$ )    npoints( $\xi_1$ )    periodic( $\xi_1$ )
# min( $\xi_2$ )    width( $\xi_2$ )    npoints( $\xi_2$ )    periodic( $\xi_2$ )
# ...        ...        ...        ...
# min( $\xi_{N_{cv}}$ ) width( $\xi_{N_{cv}}$ ) npoints( $\xi_{N_{cv}}$ ) periodic( $\xi_{N_{cv}}$ )

 $\xi_1^1$          $\xi_2^1$         ...         $\xi_{N_{cv}}^1$          $f(\xi_1^1, \xi_2^1, \dots, \xi_{N_{cv}}^1)$ 
 $\xi_1^2$          $\xi_2^2$         ...         $\xi_{N_{cv}}^2$          $f(\xi_1^2, \xi_2^2, \dots, \xi_{N_{cv}}^2)$ 
...          ...          ...          ...          ...
```

Lines beginning with the character “#” are the header of the file. N_{cv} is the number of collective variables sampled by the grid. For each variable ξ_i , $\min(\xi_i)$ is the lowest value sampled by the grid (i.e. the left-most boundary of the grid along ξ_i), $\text{width}(\xi_i)$ is the width of each grid step along ξ_i , $\text{npoints}(\xi_i)$ is the number of points and $\text{periodic}(\xi_i)$ is a flag whose value is 1 or 0 depending on whether the grid is periodic along ξ_i . In most situations:

- $\min(\xi_i)$ is given by the **lowerBoundary** keyword of the variable ξ_i ;
- $\text{width}(\xi_i)$ is given by the **width** keyword;
- $\text{npoints}(\xi_i)$ is calculated from the two above numbers and the **upperBoundary** keyword;
- $\text{periodic}(\xi_i)$ is set to 1 if and only if ξ_i is periodic and the grids’ boundaries cover its period.

Exception: there is a slightly different header in PMF files computed by ABF (6.2) or by other biases with an optional thermodynamic integration (TI) estimator (6.1). In this case, free-energy gradients are accumulated on an (npoints) -long grid along each variable ξ : after these gradients are integrated, the resulting PMF is discretized on a grid with $(\text{npoints}+1)$ points along ξ . Therefore, the edges of the PMF’s grid extend $\text{width}/2$ above and below the original boundaries (unless these are periodic). The format of the file’s header is otherwise unchanged.

After the header, the rest of the file contains values of the tabulated function $f(\xi_1, \xi_2, \dots, \xi_{N_{cv}})$, one for each line. The first N_{cv} columns contain values of $\xi_1, \xi_2, \dots, \xi_{N_{cv}}$ and the last column contains the value of the function f . Points are sorted in ascending order with the fastest-changing values at the right (“C-style” order). Each sweep of the right-most variable $\xi_{N_{cv}}$ is terminated by an empty line. For two dimensional grid files, this allows quick visualization by programs such as GNUplot.

Example 1: multicolumn text file for a one-dimensional histogram with **lowerBoundary** = 15, **upperBoundary** = 48 and **width** = 0.1.

```
# 1
# 15      0.1      330  0

15.05  6.14012e-07
15.15  7.47644e-07
...
47.85  1.65944e-06
47.95  1.46712e-06
```

Example 2: multicolumn text file for a two-dimensional histogram of two dihedral angles (periodic interval with 6° bins):

```
# 2
# -180.0  6.0      30      1
# -180.0  6.0      30      1

-177.0  -177.0  8.97117e-06
-177.0  -171.0  1.53525e-06
...
-177.0   177.0  2.442956-06

-171.0  -177.0  2.04702e-05
...      ...      ...
```

3.7.5 Output trajectory files

The Colvars trajectory file (with a suffix `.colvars.traj`) is a plain text file (scientific notation with 14-digit precision) whose columns represent quantities such as colvar values, applied forces, or individual restraints' energies. Under most scenarios, plotting or analyzing this file is straightforward: for example, the following contains a variable “A” and the energy of a restraint “rA”:

```
#      step  A      E_rA
      0  1.42467449615693e+01  6.30982865292123e+02
     100  1.42282559728026e+01  6.20640585041317e+02
...      ...
```

Occasionally, if the Colvars configuration is changed mid-run certain quantities may be added or removed, changing the column layout. Labels in comment lines can assist in such cases: for example, consider the trajectory above with the addition of a second variable, “B”, after 10,000 steps:

```
#      step  A      E_rA
      0  1.42467449615693e+01  6.30982865292123e+02
     100  1.42282559728026e+01  6.20640585041317e+02
...
#      step  A      B      E_rA
    10000  1.38136915830229e+01  9.99574098859265e-01  4.11184644791030e+02
```

```
10100    1.36437184346326e+01  9.99574091957314e-01  3.37726286543895e+02
```

Analyzing the above file with standard tools is possible, but laborious: a **convenience script** is provided for this and related purposes. It may be used either as a command-line tool or as a Python module:

```
>>> from plot_colvars_traj import Colvars_traj
>>> traj = Colvars_traj('test.colvars.traj')
>>> print(traj['A'].steps, traj['A'].values)
[   0  100  ... 10000 10100] [14.246745 14.228256 ... 13.813692 13.643718]
>>> print(traj['B'].steps, traj['B'].values)
[10000 10100] [0.999574  0.9995741]
```


4 Defining collective variables

A collective variable is defined by the keyword `colvar` followed by its configuration options contained within curly braces:

```
colvar {  
  name xi  
  <other options>  
  function_name {  
    <parameters>  
    <atom selection>  
  }  
}
```

There are multiple ways of defining a variable:

- The *simplest and most common way* is using one of the precompiled functions (here called “components”), which are listed in section 4.1. For example, using the keyword `rmsd` (section 4.5.1) defines the variable as the root mean squared deviation (RMSD) of the selected atoms.
- A new variable may also be constructed as a linear or polynomial combination of the components listed in section 4.1 (see 4.14 for details).
- A user-defined mathematical function of the existing components (see list in section 4.1), or of the atomic coordinates directly (see the `cartesian` keyword in 4.7.1). The function is defined through the keyword `customFunction` (see 4.15 for details).

Choosing a component (function) is the only parameter strictly required to define a collective variable. It is also highly recommended to specify a name for the variable:

- **Keyword name** `<Name of this colvar>`
Context: `colvar`
Acceptable values: `string`
Default value: “`colvar`” + numeric id
Description: The name is an unique case-sensitive string which allows the Colvars module to identify this colvar unambiguously; it is also used in the trajectory file to label the columns corresponding to this colvar.

4.1 Choosing a function

In this context, the function that computes a colvar is called a *component*. A component’s choice and definition consists of including in the variable’s configuration a keyword indicating the type of function (e.g. `rmsd`), followed by a definition block specifying the atoms involved (see 5) and any additional parameters (cutoffs, “reference” values, ...). *At least one component must be chosen to define a variable*: if none of the keywords listed below is found, an error is raised.

The following components implement functions with a scalar value (i.e. a real number):

- **distance**: distance between two groups;
- **distanceZ**: projection of a distance vector on an axis;
- **distanceXY**: projection of a distance vector on a plane;
- **distanceInv**: mean distance between two groups of atoms (e.g. NOE-based distance);
- **angle**: angle between three groups;
- **dihedral**: torsional (dihedral) angle between four groups;
- **dipoleAngle**: angle between two groups and dipole of a third group;
- **dipoleMagnitude**: magnitude of the dipole of a group of atoms;
- **polarTheta**: polar angle of a group in spherical coordinates;
- **polarPhi**: azimuthal angle of a group in spherical coordinates;
- **eulerPhi**: Roll angle of rotation from references coordinates;
- **eulerTheta**: Pitch angle of rotation from references coordinates;
- **eulerPsi**: Yaw angle of rotation from references coordinates;
- **coordNum**: coordination number between two groups;
- **selfCoordNum**: coordination number of atoms within a group;
- **hBond**: hydrogen bond between two atoms;
- **rmsd**: root mean square deviation (RMSD) from a set of reference coordinates;
- **eigenvector**: projection of the atomic coordinates on a vector;
- **orientationAngle**: angle of the best-fit rotation from a set of reference coordinates;
- **orientationProj**: cosine of **orientationProj**;
- **spinAngle**: projection orthogonal to an axis of the best-fit rotation from a set of reference coordinates;
- **tilt**: projection on an axis of the best-fit rotation from a set of reference coordinates;
- **gyration**: radius of gyration of a group of atoms;
- **inertia**: moment of inertia of a group of atoms;
- **inertiaZ**: moment of inertia of a group of atoms around a chosen axis;

Some components do not return scalar, but vector values:

- **distanceVec**: distance vector between two groups (length: 3);
- **distanceDir**: unit vector parallel to distanceVec (length: 3);
- **cartesian**: vector of atomic Cartesian coordinates (length: N times the number of Cartesian components requested, X, Y or Z);

- **distancePairs**: vector of mutual distances (length: $N_1 \times N_2$);
- **orientation**: best-fit rotation, expressed as a unit quaternion (length: 4).

The types of components used in a colvar (scalar or not) determine the properties of that colvar, and particularly which biasing or analysis methods can be applied.

What if “X” is not listed? If a function type is not available on this list, it may be possible to define it as a polynomial superposition of existing ones (see 4.14), a custom function (see 4.15).

In the rest of this section, all available component types are listed, along with their physical units and the ranges of values, if limited. Such limiting values can be used to define **lowerBoundary** and **upperBoundary** in the parent colvar.

For each type of component, the available configurations keywords are listed: when two components share certain keywords, the second component references to the documentation of the first one that uses that keyword. The very few keywords that are available for all types of components are listed in a separate section 4.11.

4.2 Distances

4.2.1 distance: center-of-mass distance between two groups.

The distance `{...}` block defines a distance component between the two atom groups, group1 and group2.

List of keywords (see also 4.14 for additional options):

- **Keyword** group1 `<First group of atoms>`
Context: distance
Acceptable values: Block group1 `{...}`
Description: First group of atoms.
- **Keyword** group2: analogous to group1
- **Keyword** forceNoPBC `<Calculate absolute rather than minimum-image distance?>`
Context: distance
Acceptable values: boolean
Default value: no
Description: By default, in calculations with periodic boundary conditions, the distance component returns the distance according to the minimum-image convention. If this parameter is set to yes, PBC will be ignored and the distance between the coordinates as maintained internally will be used. This is only useful in a limited number of special cases, e.g. to describe the distance between remote points of a single macromolecule, which cannot be split across periodic cell boundaries, and for which the minimum-image distance might give the wrong result because of a relatively small periodic cell.
- **Keyword** oneSiteTotalForce `<Measure total force on group 1 only?>`
Context: angle, dipoleAngle, dihedral
Acceptable values: boolean

Default value: no

Description: If this is set to yes, the total force is measured along a vector field (see equation (20) in section 6.2) that only involves atoms of group1. This option is only useful for ABF, or custom biases that compute total forces. See section 6.2 for details.

The value returned is a positive number (in *length unit*), ranging from 0 to the largest possible interatomic distance within the chosen boundary conditions (with PBCs, the minimum image convention is used unless the forceNoPBC option is set).

4.2.2 distanceZ: projection of a distance vector on an axis.

The distanceZ {...} block defines a distance projection component, which can be seen as measuring the distance between two groups projected onto an axis, or the position of a group along such an axis. The axis can be defined using either one reference group and a constant vector, or dynamically based on two reference groups. One of the groups can be set to a dummy atom to allow the use of an absolute Cartesian coordinate.

List of keywords (see also 4.14 for additional options):

- **Keyword** main ⟨Main group of atoms⟩
Context: distanceZ
Acceptable values: Block main {...}
Description: Group of atoms whose position r is measured.
- **Keyword** ref ⟨Reference group of atoms⟩
Context: distanceZ
Acceptable values: Block ref {...}
Description: Reference group of atoms. The position of its center of mass is noted r_1 below.
- **Keyword** ref2 ⟨Secondary reference group⟩
Context: distanceZ
Acceptable values: Block ref2 {...}
Default value: none
Description: Optional group of reference atoms, whose position r_2 can be used to define a dynamic projection axis: $e = (\|r_2 - r_1\|)^{-1} \times (r_2 - r_1)$. In this case, the origin is $r_m = 1/2(r_1 + r_2)$, and the value of the component is $e \cdot (r - r_m)$.
- **Keyword** axis ⟨Projection axis⟩
Context: distanceZ
Acceptable values: (x, y, z) triplet
Default value: (0.0, 0.0, 1.0)
Description: This vector will be normalized to define a projection axis e for the distance vector $r - r_1$ joining the centers of groups ref and main. The value of the component is then $e \cdot (r - r_1)$. The vector should be written as three components separated by commas and enclosed in parentheses.
- **Keyword** forceNoPBC: see definition of forceNoPBC (distance component)
- **Keyword** oneSiteTotalForce: see definition of oneSiteTotalForce (distance component)

This component returns a number (in *length unit*) whose range is determined by the chosen boundary conditions. For instance, if the z axis is used in a simulation with periodic boundaries, the returned value ranges between $-b_z/2$ and $b_z/2$, where b_z is the box length along z (this behavior is disabled if `forceNoPBC` is set).

4.2.3 distanceXY: modulus of the projection of a distance vector on a plane.

The `distanceXY { ... }` block defines a distance projected on a plane, and accepts the same keywords as the component `distanceZ`, i.e. `main`, `ref`, either `ref2` or `axis`, and `oneSiteTotalForce`. It returns the norm of the projection of the distance vector between `main` and `ref` onto the plane orthogonal to the axis. The axis is defined using the `axis` parameter or as the vector joining `ref` and `ref2` (see `distanceZ` above).

List of keywords (see also 4.14 for additional options):

- **Keyword** `main`: see definition of `main` (`distanceZ` component)
- **Keyword** `ref`: see definition of `ref` (`distanceZ` component)
- **Keyword** `ref2`: see definition of `ref2` (`distanceZ` component)
- **Keyword** `axis`: see definition of `axis` (`distanceZ` component)
- **Keyword** `forceNoPBC`: see definition of `forceNoPBC` (distance component)
- **Keyword** `oneSiteTotalForce`: see definition of `oneSiteTotalForce` (distance component)

4.2.4 distanceVec: distance vector between two groups.

The `distanceVec { ... }` block defines a distance vector component, which accepts the same keywords as the component `distance`: `group1`, `group2`, and `forceNoPBC`. Its value is the 3-vector joining the centers of mass of `group1` and `group2`.

List of keywords (see also 4.14 for additional options):

- **Keyword** `group1`: see definition of `group1` (distance component)
- **Keyword** `group2`: analogous to `group1`
- **Keyword** `forceNoPBC`: see definition of `forceNoPBC` (distance component)
- **Keyword** `oneSiteTotalForce`: see definition of `oneSiteTotalForce` (distance component)

4.2.5 distanceDir: distance unit vector between two groups.

The distanceDir {...} block defines a distance unit vector component, which accepts the same keywords as the component distance: group1, group2, and forceNoPBC. It returns a 3-dimensional unit vector $\mathbf{d} = (d_x, d_y, d_z)$, with $|\mathbf{d}| = 1$.

List of keywords (see also 4.14 for additional options):

- **Keyword** group1: see definition of group1 (distance component)
- **Keyword** group2: analogous to group1
- **Keyword** forceNoPBC: see definition of forceNoPBC (distance component)
- **Keyword** oneSiteTotalForce: see definition of oneSiteTotalForce (distance component)

4.2.6 distanceInv: mean distance between two groups of atoms.

The distanceInv {...} block defines a generalized mean distance between two groups of atoms 1 and 2, where each distance is taken to the power $-n$:

$$d_{1,2}^{[n]} = \left(\frac{1}{N_1 N_2} \sum_{i,j} d_{ij}^{-n} \right)^{-1/n} \quad (2)$$

where d_{ij} is the distance between atoms i and j in groups 1 and 2 respectively, and n is an even integer.

List of keywords (see also 4.14 for additional options):

- **Keyword** group1: see definition of group1 (distance component)
- **Keyword** group2: analogous to group1
- **Keyword** oneSiteTotalForce: see definition of oneSiteTotalForce (distance component)
- **Keyword** exponent <Exponent n in equation 2>

Context: distanceInv

Acceptable values: positive even integer

Default value: 6

Description: Defines the exponent to which the individual distances are elevated before averaging. The default value of 6 is useful for example to applying restraints based on NOE-measured distances.

This component returns a number ranging from 0 to the largest possible distance within the chosen boundary conditions.

4.3 Angles

4.3.1 angle: angle between three groups.

The angle {...} block defines an angle, and contains the three blocks group1, group2 and group3, defining the three groups. It returns an angle (in degrees) within the interval [0 : 180].

List of keywords (see also 4.14 for additional options):

- **Keyword** group1: see definition of group1 (distance component)
- **Keyword** group2: analogous to group1
- **Keyword** group3: analogous to group1
- **Keyword** forceNoPBC: see definition of forceNoPBC (distance component)
- **Keyword** oneSiteTotalForce: see definition of oneSiteTotalForce (distance component)

4.3.2 dipoleAngle: angle between two groups and dipole of a third group.

The dipoleAngle {...} block defines an angle, and contains the three blocks group1, group2 and group3, defining the three groups, being group1 the group where dipole is calculated. It returns an angle (in degrees) within the interval [0 : 180].

List of keywords (see also 4.14 for additional options):

- **Keyword** group1: see definition of group1 (distance component)
- **Keyword** group2: analogous to group1
- **Keyword** group3: analogous to group1
- **Keyword** forceNoPBC: see definition of forceNoPBC (distance component)
- **Keyword** oneSiteTotalForce: see definition of oneSiteTotalForce (distance component)

4.3.3 dihedral: torsional angle between four groups.

The dihedral {...} block defines a torsional angle, and contains the blocks group1, group2, group3 and group4, defining the four groups. It returns an angle (in degrees) within the interval [-180 : 180]. The Colvars module calculates all the distances between two angles taking into account periodicity. For instance, reference values for restraints or range boundaries can be defined by using any real number of choice.

List of keywords (see also 4.14 for additional options):

- **Keyword** group1: see definition of **group1** (distance component)
- **Keyword** group2: analogous to group1
- **Keyword** group3: analogous to group1
- **Keyword** group4: analogous to group1
- **Keyword** forceNoPBC: see definition of **forceNoPBC** (distance component)
- **Keyword** oneSiteTotalForce: see definition of **oneSiteTotalForce** (distance component)

4.3.4 polarTheta: polar angle in spherical coordinates.

The polarTheta { ... } block defines the polar angle in spherical coordinates, for the center of mass of a group of atoms described by the block atoms. It returns an angle (in degrees) within the interval [0 : 180]. To obtain spherical coordinates in a frame of reference tied to another group of atoms, use the fittingGroup (5.2) option within the atoms block. An example is provided in file examples/11_polar_angles.in of the Colvars public repository.

List of keywords (see also 4.14 for additional options):

- **Keyword** atoms {Atom group}
Context: polarPhi
Acceptable values: atoms { ... } block
Description: Defines the group of atoms for the COM of which the angle should be calculated.

4.3.5 polarPhi: azimuthal angle in spherical coordinates.

The polarPhi { ... } block defines the azimuthal angle in spherical coordinates, for the center of mass of a group of atoms described by the block atoms. It returns an angle (in degrees) within the interval [−180 : 180]. The Colvars module calculates all the distances between two angles taking into account periodicity. For instance, reference values for restraints or range boundaries can be defined by using any real number of choice. To obtain spherical coordinates in a frame of reference tied to another group of atoms, use the fittingGroup (5.2) option within the atoms block. An example is provided in file examples/11_polar_angles.in of the Colvars public repository.

List of keywords (see also 4.14 for additional options):

- **Keyword** atoms {Atom group}
Context: polarPhi
Acceptable values: atoms { ... } block
Description: Defines the group of atoms for the COM of which the angle should be calculated.

4.4 Contacts

4.4.1 coordNum: coordination number between two groups.

The coordNum {...} block defines a coordination number (or number of contacts), which calculates the function $(1 - (d/d_0)^n)/(1 - (d/d_0)^m)$, where d_0 is the “cutoff” distance, and n and m are exponents that can control its long range behavior and stiffness [3]. This function is summed over all pairs of atoms in group1 and group2:

$$C(\text{group1}, \text{group2}) = \sum_{i \in \text{group1}} \sum_{j \in \text{group2}} \frac{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^n}{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^m} \quad (3)$$

List of keywords (see also 4.14 for additional options):

- **Keyword group1:** see definition of **group1** (distance component)
- **Keyword group2:** analogous to group1
- **Keyword cutoff** \langle “Interaction” distance (*length unit*) \rangle
Context: coordNum
Acceptable values: positive decimal
Default value: 4.0 Å
Description: This number defines the switching distance to define an interatomic contact: for $d \ll d_0$, the switching function $(1 - (d/d_0)^n)/(1 - (d/d_0)^m)$ is close to 1, at $d = d_0$ it has a value of n/m ($1/2$ with the default n and m), and at $d \gg d_0$ it goes to zero approximately like d^{m-n} . Hence, for a proper behavior, m must be larger than n .
- **Keyword cutoff3** \langle Reference distance vector (*length unit*) \rangle
Context: coordNum
Acceptable values: “(x, y, z)” triplet of positive decimals
Default value: (4.0, 4.0, 4.0) Å
Description: The three components of this vector define three different cutoffs d_0 for each direction. This option is mutually exclusive with cutoff.
- **Keyword expNumer** \langle Numerator exponent \rangle
Context: coordNum
Acceptable values: positive even integer
Default value: 6
Description: This number defines the n exponent for the switching function.
- **Keyword expDenom** \langle Denominator exponent \rangle
Context: coordNum
Acceptable values: positive even integer
Default value: 12
Description: This number defines the m exponent for the switching function.

- **Keyword** `group2CenterOnly` ⟨Use only group2’s center of mass⟩
Context: `coordNum`
Acceptable values: `boolean`
Default value: `off`
Description: If this option is on, only contacts between each atoms in group1 and the center of mass of group2 are calculated (by default, the sum extends over all pairs of atoms in group1 and group2). If group2 is a dummyAtom, this option is set to yes by default.
- **Keyword** `tolerance` ⟨Pairlist control⟩
Context: `coordNum`
Acceptable values: `decimal`
Default value: `0.0`
Description: This controls the pairlist feature, dictating the minimum value for each summation element in Eq. 3 such that the pair that contributed the summation element is included in subsequent simulation timesteps until the next pairlist recalculation. For most applications, this value should be small (eg. 0.001) to avoid missing important contributions to the overall sum. Higher values will improve performance by reducing the number of pairs that contribute to the sum. Values above 1 will exclude all possible pair interactions. Similarly, values below 0 will never exclude a pair from consideration. To ensure continuous forces, Eq. 3 is further modified by subtracting the tolerance and then rescaling so that each pair covers the range $[0, 1]$.
- **Keyword** `pairListFrequency` ⟨Pairlist regeneration frequency⟩
Context: `coordNum`
Acceptable values: `positive integer`
Default value: `100`
Description: This controls the pairlist feature, dictating how many steps are taken between regenerating pairlists if the tolerance is greater than 0.

This component returns a dimensionless number, which ranges from approximately 0 (all interatomic distances are much larger than the cutoff) to $N_{\text{group1}} \times N_{\text{group2}}$ (all distances are less than the cutoff), or N_{group1} if `group2CenterOnly` is used. For performance reasons, at least one of group1 and group2 should be of limited size or `group2CenterOnly` should be used: the cost of the loop over all pairs grows as $N_{\text{group1}} \times N_{\text{group2}}$. Setting `tolerance > 0` ameliorates this to some degree, although every pair is still checked to regenerate the pairlist.

4.4.2 `selfCoordNum`: coordination number between atoms within a group.

The `selfCoordNum {...}` block defines a coordination number similarly to the component `coordNum`, but the function is summed over atom pairs within group1:

$$C(\text{group1}) = \sum_{i \in \text{group1}} \sum_{j > i} \frac{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^n}{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^m} \quad (4)$$

The keywords accepted by `selfCoordNum` are a subset of those accepted by `coordNum`, namely `group1` (here defining *all* of the atoms to be considered), `cutoff`, `expNumer`, and `expDenom`.

List of keywords (see also 4.14 for additional options):

- **Keyword** `group1`: see definition of `group1` (`coordNum` component)

- **Keyword** `cutoff`: see definition of `cutoff` (coordNum component)
- **Keyword** `cutoff3`: see definition of `cutoff3` (coordNum component)
- **Keyword** `expNumer`: see definition of `expNumer` (coordNum component)
- **Keyword** `expDenom`: see definition of `expDenom` (coordNum component)
- **Keyword** `tolerance`: see definition of `tolerance` (coordNum component)
- **Keyword** `pairListFrequency`: see definition of `pairListFrequency` (coordNum component)

This component returns a dimensionless number, which ranges from approximately 0 (all interatomic distances much larger than the cutoff) to $N_{\text{group1}} \times (N_{\text{group1}} - 1)/2$ (all distances within the cutoff). For performance reasons, group1 should be of limited size, because the cost of the loop over all pairs grows as N_{group1}^2 .

4.4.3 hBond: hydrogen bond between two atoms.

The `hBond { ... }` block defines a hydrogen bond, implemented as a coordination number (eq. 3) between the donor and the acceptor atoms. Therefore, it accepts the same options `cutoff` (with a different default value of 3.3 Å), `expNumer` (with a default value of 6) and `expDenom` (with a default value of 8). Unlike `coordNum`, it requires two atom numbers, acceptor and donor, to be defined. It returns an adimensional number, with values between 0 (acceptor and donor far outside the cutoff distance) and 1 (acceptor and donor much closer than the cutoff).

List of keywords (see also 4.14 for additional options):

- **Keyword** `acceptor` <Number of the acceptor atom>
Context: `hBond`
Acceptable values: positive integer
Description: Number that uses the same convention as `atomNumbers`.
- **Keyword** `donor`: analogous to `acceptor`
- **Keyword** `cutoff`: see definition of `cutoff` (coordNum component)
Note: default value is 3.3 Å.
- **Keyword** `expNumer`: see definition of `expNumer` (coordNum component)
Note: default value is 6.
- **Keyword** `expDenom`: see definition of `expDenom` (coordNum component)
Note: default value is 8.

4.5 Collective metrics

4.5.1 rmsd: root mean square displacement (RMSD) from reference positions.

The block `rmsd {...}` defines the root mean square replacement (RMSD) of a group of atoms with respect to a reference structure. For each set of coordinates $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$, the colvar component `rmsd` calculates the optimal rotation $U^{\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}}$ that best superimposes the coordinates $\{\mathbf{x}_i(t)\}$ onto a set of reference coordinates $\{\mathbf{x}_i^{(\text{ref})}\}$. Both the current and the reference coordinates are centered on their centers of geometry, $\mathbf{x}_{\text{cog}}(t)$ and $\mathbf{x}_{\text{cog}}^{(\text{ref})}$. The root mean square displacement is then defined as:

$$\text{RMSD}(\{\mathbf{x}_i(t)\}, \{\mathbf{x}_i^{(\text{ref})}\}) = \sqrt{\frac{1}{N} \sum_{i=1}^N \left| U(\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)) - (\mathbf{x}_i^{(\text{ref})} - \mathbf{x}_{\text{cog}}^{(\text{ref})}) \right|^2} \quad (5)$$

The optimal rotation $U^{\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}}$ is calculated within the formalism developed in reference [4], which guarantees a continuous dependence of $U^{\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}}$ with respect to $\{\mathbf{x}_i(t)\}$.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms` \langle Atom group \rangle
Context: `rmsd`
Acceptable values: `atoms {...}` block
Description: Defines the group of atoms of which the RMSD should be calculated. Optimal fit options (such as `refPositions` and `rotateToReference`) should typically NOT be set within this block. Exceptions to this rule are the special cases discussed in the *Advanced usage* paragraph below.
- **Keyword** `refPositions` \langle Reference coordinates \rangle
Context: `rmsd`
Acceptable values: space-separated list of (x, y, z) triplets
Description: This option (mutually exclusive with `refPositionsFile`) sets the reference coordinates for RMSD calculation, and uses these to compute the roto-translational fit. See the equivalent option `refPositions` within the atom group definition for details on acceptable formats and other features.
- **Keyword** `refPositionsFile` \langle Reference coordinates file \rangle
Context: `rmsd`
Acceptable values: UNIX filename
Description: This option (mutually exclusive with `refPositions`) sets the reference coordinates for RMSD calculation, and uses these to compute the roto-translational fit. See the equivalent option `refPositionsFile` within the atom group definition for details on acceptable file formats and other features.
- **Keyword** `atomPermutation` \langle Alternate ordering of atoms for RMSD computation \rangle
Context: `rmsd`
Acceptable values: List of atom numbers
Description: If defined, this parameter defines a re-ordering (permutation) of the 1-based atom numbers that can be used to compute the RMSD, typically due to molecular symmetry. This parameter can be specified multiple times, each one defining a new permutation: the returned RMSD value is the minimum over the set of permutations. For example, if the atoms making up the group are 6, 7,

8, 9, and atoms 7, 8, and 9 are invariant by circular permutation (as the hydrogens in a CH₃ group), a symmetry-adapted RMSD would be obtained by adding:

```
atomPermutation 6 8 9 7
```

```
atomPermutation 6 9 7 8
```

Note that *this does not affect the least-squares roto-translational fit*, which is done using the topology ordering of atoms, and the reference positions in the order provided. Therefore, this feature is mostly useful when using custom fitting parameters within the atom group, such as `fittingGroup`, or when fitting is disabled altogether.

This component returns a positive real number (in *length unit*).

4.5.2 Advanced usage of the rmsd component.

In the standard usage as described above, the rmsd component calculates a minimum RMSD, that is, current coordinates are optimally fitted onto the same reference coordinates that are used to compute the RMSD value. The fit itself is handled by the atom group object, whose parameters are automatically set by the rmsd component. For very specific applications, however, it may be useful to control the fitting process separately from the definition of the reference coordinates, to evaluate various types of non-minimal RMSD values. This can be achieved by setting the related options (`refPositions`, etc.) explicitly in the atom group block. This allows for the following non-standard cases:

1. applying the optimal translation, but no rotation (`rotateToReference off`), to bias or restrain the shape and orientation, but not the position of the atom group;
2. applying the optimal rotation, but no translation (`centerToReference off`), to bias or restrain the shape and position, but not the orientation of the atom group;
3. disabling the application of optimal roto-translations, which lets the RMSD component describe the deviation of atoms from fixed positions in the laboratory frame: this allows for custom positional restraints within the Colvars module;
4. fitting the atomic positions to different reference coordinates than those used in the RMSD calculation itself (by specifying `refPositions` or `refPositionsFile` within the atom group as well as within the rmsd block);
5. applying the optimal rotation and/or translation from a separate atom group, defined through `fittingGroup`: the RMSD then reflects the deviation from reference coordinates in a separate, moving reference frame (see example in the section on `fittingGroup`).

4.5.3 eigenvector: projection of the atomic coordinates on a vector.

The block `eigenvector { ... }` defines the projection of the coordinates of a group of atoms (or more precisely, their deviations from the reference coordinates) onto a vector in \mathbb{R}^{3n} , where n is the number of atoms in the group. The computed quantity is the total projection:

$$p(\{\mathbf{x}_i(t)\}, \{\mathbf{x}_i^{(\text{ref})}\}) = \sum_{i=1}^n \mathbf{v}_i \cdot \left(U(\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)) - (\mathbf{x}_i^{(\text{ref})} - \mathbf{x}_{\text{cog}}^{(\text{ref})}) \right), \quad (6)$$

where, as in the rmsd component, U is the optimal rotation matrix, $\mathbf{x}_{\text{cog}}(t)$ and $\mathbf{x}_{\text{cog}}^{(\text{ref})}$ are the centers of geometry of the current and reference positions respectively, and \mathbf{v}_i are the components of the vector for each atom. Example choices for (\mathbf{v}_i) are an eigenvector of the covariance matrix (essential mode), or a normal mode of the system. It is assumed that $\sum_i \mathbf{v}_i = 0$: otherwise, the Colvars module centers the \mathbf{v}_i automatically when reading them from the configuration.

List of keywords (see also 4.14 for additional options):

- **Keyword atoms:** see definition of **atoms** (rmsd component)
- **Keyword refPositions:** see definition of **refPositions** (rmsd component)
- **Keyword refPositionsFile:** see definition of **refPositionsFile** (rmsd component)
- **Keyword vector** \langle Vector components \rangle
Context: eigenvector
Acceptable values: space-separated list of (x, y, z) triplets
Description: This option (mutually exclusive with vectorFile) sets the values of the vector components.
- **Keyword vectorFile** \langle file containing vector components \rangle
Context: eigenvector
Acceptable values: UNIX filename
Description: This option (mutually exclusive with vector) sets the name of an XYZ (3.7.3) coordinate file containing the vector components. **Note:** *Reading data from a coordinate file may entail an automatic unit conversion if the length unit currently used by the MD engine is not Å. If this is not the desired behavior, this can be remedied using the normalizeVector option described below.*
- **Keyword normalizeVector** \langle Normalize the vector components when reading them? \rangle
Context: eigenvector
Acceptable values: boolean
Default value: off
Description: If this flag is activated, the norm of the vector $|\mathbf{v}| = \sqrt{\sum_i |\mathbf{v}_i|^2}$ is set equal to 1 by automatically rescaling all the components \mathbf{v}_i ; alternatively, the value of $|\mathbf{v}|$ is printed.
- **Keyword differenceVector** \langle The 3n-dimensional vector is the difference between vector and refPositions \rangle
Context: eigenvector
Acceptable values: boolean
Default value: off
Description: If this option is on, the numbers provided by vector are interpreted as another set of positions, \mathbf{x}'_i : the vector \mathbf{v}_i is then defined as $\mathbf{v}_i = (\mathbf{x}'_i - \mathbf{x}_i^{(\text{ref})})$. This allows to conveniently define a colvar ξ as a projection on the linear transformation between two sets of positions, “A” and “B”. If this flag is on, the vector is normalized so that $\xi = 0$ when the atoms are at the set of positions “A” and $\xi = 1$ at the set of positions “B”. Setting normalizeVector on overrides this behavior.

4.5.4 gyration: radius of gyration of a group of atoms.

The block `gyration` `{...}` defines the parameters for calculating the radius of gyration of a group of atomic positions $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$ with respect to their center of geometry, $\mathbf{x}_{\text{cog}}(t)$:

$$R_{\text{gyr}} = \sqrt{\frac{1}{N} \sum_{i=1}^N |\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)|^2} \quad (7)$$

This component must contain one `atoms` `{...}` block to define the atom group, and returns a positive number, expressed in *length unit*.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)

4.5.5 inertia: total moment of inertia of a group of atoms.

The block `inertia` `{...}` defines the parameters for calculating the total moment of inertia of a group of atomic positions $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$ with respect to their center of geometry, $\mathbf{x}_{\text{cog}}(t)$:

$$I = \sum_{i=1}^N |\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)|^2 \quad (8)$$

Note that all atomic masses are set to 1 for simplicity. This component must contain one `atoms` `{...}` block to define the atom group, and returns a positive number, expressed in *length unit*².

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)

4.5.6 dipoleMagnitude: dipole magnitude of a group of atoms.

The `dipoleMagnitude` `{...}` block defines the dipole magnitude of a group of atoms (norm of the dipole moment's vector), being `atoms` the group where dipole magnitude is calculated. It returns the magnitude in elementary charge e times *length unit*.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)

4.5.7 inertiaZ: total moment of inertia of a group of atoms around a chosen axis.

The block `inertiaZ {...}` defines the parameters for calculating the component along the axis \mathbf{e} of the moment of inertia of a group of atomic positions $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$ with respect to their center of geometry, $\mathbf{x}_{\text{cog}}(t)$:

$$I_{\mathbf{e}} = \sum_{i=1}^N ((\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)) \cdot \mathbf{e})^2 \quad (9)$$

Note that all atomic masses are set to 1 for simplicity. This component must contain one `atoms {...}` block to define the atom group, and returns a positive number, expressed in *length unit*².

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)
- **Keyword** `axis` \langle Projection axis \rangle
Context: `inertiaZ`
Acceptable values: (x, y, z) triplet
Default value: (0.0, 0.0, 1.0)
Description: The three components of this vector define (when normalized) the projection axis \mathbf{e} .

4.6 Rotations

The variables discussed in this section quantify the rotations of macromolecules (or other quasi-rigid objects) from a given set of reference coordinates to the current coordinates. Such rotations are computed following the same method used for best-fit RMSDs (see `rmsd` and `fittingGroup`). The underlying mathematical formalism is described in reference [4], and the implementation in reference [1].

The first of the functions described is the `orientation`, which describes the full rotation as a unit quaternion $q = (q_0, q_1, q_2, q_3)$, i.e. 4 numbers with one constraint (3 degrees of freedom). The quaternion q is *one of only two representations that are both complete and accurate*, the other being a 3×3 unit matrix with 3 independent parameters. Although q is used internally in the Colvars module for features such as the `rmsd` function and the `fittingGroup` option, its direct use as a collective variable is more difficult, and mostly limited to fixed or moving restraints.

The two functions `orientationAngle` and `orientationProj`, with the latter being the cosine of the former, represent the amplitude of the full rotation described by q , regardless of the direction of its axis. As one-dimensional scalar variables, both `orientationAngle` and `orientationProj` are a much reduced simplification of the full rotation. However, they can be used in a variety of methods including both restraints and PMF computations.

A slightly more complete parametrization is achieved by decomposing the full rotation into the two parameters, `tilt` and `spinAngle`. These quantify the amplitudes of two independent sub-rotations away from a certain axis \mathbf{e} , and around the same axis \mathbf{e} , respectively. The axis \mathbf{e} is chosen by the user, and is by default the Z axis: under that choice, `tilt` is equivalent to the sine of the Euler “pitch” angle θ , and `spinAngle` to the sum of the other two angles, $(\phi + \psi)$. This parameterization is mathematically well defined for almost all full rotations, including small ones when the current coordinates are almost completely aligned with the reference ones. However, a mathematical singularity prevents using the `spinAngle` function near configurations where the value of `tilt` is -1 (i.e. a -180° rotation around an axis orthogonal to \mathbf{e}). For these reasons, `tilt` and `spinAngle` are useful when the allowed rotations are known to have approximately the same axis, and differ only in the magnitude of the corresponding angle. In this use case, `spinAngle`

measures the angle of the sub-rotation around the chosen axis \mathbf{e} , whereas `tilt` measures the dot product between \mathbf{e} and the actual axis of the full rotation.

Lastly, the traditional Euler angles are also available as the functions `eulerPhi`, `eulerTheta` and `eulerPsi`. Altogether, these are sufficient to represent all three degrees of freedom of a full rotation. However, they also suffer from the potential “gimbal lock” problem, which emerges whenever $\theta \simeq \pm 90^\circ$, which includes also the case where the full rotation is small. Under such conditions, the angles ϕ and ψ are both ill-defined and cannot be used as collective variables. For these reasons, it is highly recommended that Euler angles are used only in simulations where their range of applicability is *known ahead of time*, and excludes configurations where $\theta \simeq \pm 90^\circ$ altogether.

4.6.1 orientation: orientation from reference coordinates.

The block `orientation` `{...}` returns the same optimal rotation used in the `rmsd` component to superimpose the coordinates $\{\mathbf{x}_i(t)\}$ onto a set of reference coordinates $\{\mathbf{x}_i^{(\text{ref})}\}$. Such component returns a four dimensional vector $\mathbf{q} = (q_0, q_1, q_2, q_3)$, with $\sum_i q_i^2 = 1$; this *quaternion* expresses the optimal rotation $\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}$ according to the formalism in reference [4]. The quaternion (q_0, q_1, q_2, q_3) can also be written as $(\cos(\theta/2), \sin(\theta/2)\mathbf{u})$, where θ is the angle and \mathbf{u} the normalized axis of rotation; for example, a rotation of 90° around the z axis is expressed as “(0.707, 0.0, 0.0, 0.707)”. The script `quaternion2rmatrix.tcl` provides Tcl functions for converting to and from a 4×4 rotation matrix in a format suitable for usage in VMD.

As for the component `rmsd`, the available options are `atoms`, `refPositionsFile` and `refPositions`.

Note: `refPositions` and `refPositionsFile` define the set of positions *from which* the optimal rotation is calculated, but this rotation is not applied to the coordinates of the atoms involved: it is used instead to define the variable itself.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (`rmsd` component)
- **Keyword** `refPositions`: see definition of `refPositions` (`rmsd` component)
- **Keyword** `refPositionsFile`: see definition of `refPositionsFile` (`rmsd` component)

- **Keyword** `closestToQuaternion` \langle Reference rotation \rangle

Context: `orientation`

Acceptable values: “(q0, q1, q2, q3)” quadruplet

Default value: (1.0, 0.0, 0.0, 0.0) (“null” rotation)

Description: Between the two equivalent quaternions (q_0, q_1, q_2, q_3) and $(-q_0, -q_1, -q_2, -q_3)$, the closer to (1.0, 0.0, 0.0, 0.0) is chosen. This simplifies the visualization of the colvar trajectory when sampled values are a smaller subset of all possible rotations. **Note:** *this only affects the output, never the dynamics.*

Tip: stopping the rotation of a protein. To stop the rotation of an elongated macromolecule in solution (and use an anisotropic box to save water molecules), it is possible to define a colvar with an `orientation` component, and restrain it through the harmonic bias around the identity rotation, (1.0, 0.0, 0.0, 0.0). Only the overall orientation of the macromolecule is affected, and *not* its internal degrees of freedom. The user should also take care that the macromolecule is composed by a single chain, or disable `wrapAll` otherwise.

4.6.2 orientationAngle: angle of rotation from reference coordinates.

The block `orientationAngle {...}` accepts the same base options as the component `orientation`: `atoms`, `refPositions`, `refPositionsFile`. The returned value is the angle of rotation θ between the current and the reference positions. This angle is expressed in degrees within the range $[0^\circ:180^\circ]$.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)
- **Keyword** `refPositions`: see definition of `refPositions` (rmsd component)
- **Keyword** `refPositionsFile`: see definition of `refPositionsFile` (rmsd component)

4.6.3 orientationProj: cosine of the angle of rotation from reference coordinates.

The block `orientationProj {...}` accepts the same base options as the component `orientation`: `atoms`, `refPositions`, `refPositionsFile`. The returned value is the cosine of the angle of rotation θ between the current and the reference positions. The range of values is $[-1:1]$.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)
- **Keyword** `refPositions`: see definition of `refPositions` (rmsd component)
- **Keyword** `refPositionsFile`: see definition of `refPositionsFile` (rmsd component)

4.6.4 spinAngle: angle of rotation around a given axis.

The complete rotation described by `orientation` can optionally be decomposed into two sub-rotations: one is a “*spin*” rotation around \mathbf{e} , and the other a “*tilt*” rotation around an axis orthogonal to \mathbf{e} . The component `spinAngle` measures the angle of the “*spin*” sub-rotation around \mathbf{e} .

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)
- **Keyword** `refPositions`: see definition of `refPositions` (rmsd component)
- **Keyword** `refPositionsFile`: see definition of `refPositionsFile` (rmsd component)

- **Keyword** `axis` ⟨Special rotation axis⟩

Context: `tilt`

Acceptable values: (x, y, z) triplet

Default value: (0.0, 0.0, 1.0)

Description: The three components of this vector define (when normalized) the special rotation axis used to calculate the `tilt` and `spinAngle` components.

The component `spinAngle` returns an angle (in degrees) within the periodic interval $[-180 : 180]$.

Note: the value of `spinAngle` is a continuous function almost everywhere, with the exception of configurations with the corresponding “tilt” angle equal to 180° (i.e. the `tilt` component is equal to -1): in those cases, `spinAngle` is undefined. If such configurations are expected, consider defining a `tilt` colvar using the same axis `e`, and restraining it with a lower wall away from -1 .

4.6.5 `tilt`: cosine of the rotation orthogonal to a given axis.

The component `tilt` measures the cosine of the angle of the “tilt” sub-rotation, which combined with the “spin” sub-rotation provides the complete rotation of a group of atoms. The cosine of the tilt angle rather than the tilt angle itself is implemented, because the latter is unevenly distributed even for an isotropic system: consider as an analogy the angle θ in the spherical coordinate system. The component `tilt` relies on the same options as `spinAngle`, including the definition of the axis `e`. The values of `tilt` are real numbers in the interval $[-1 : 1]$: the value 1 represents an orientation fully parallel to `e` (tilt angle = 0°), and the value -1 represents an anti-parallel orientation.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)
- **Keyword** `refPositions`: see definition of `refPositions` (rmsd component)
- **Keyword** `refPositionsFile`: see definition of `refPositionsFile` (rmsd component)
- **Keyword** `axis`: see definition of `axis` (`spinAngle` component)

4.6.6 `eulerPhi`: Roll angle from references coordinates.

Assuming the axes of the original frame are denoted as x, y, z and the axes of the rotated frame as X, Y, Z , the line of nodes, N , can be defined as the intersection of the plane xy and XY . The axis perpendicular to N and z is defined as P . In this case, as illustrated in the figure below, the complete rotation described by orientation can optionally be decomposed into three Euler angles:

- (i) the “roll” angle ϕ , i.e. the rotation angle from the x axis to the N axis;
- (ii) the “pitch” angle θ , i.e. the rotation angle from the P axis to the Z axis, and
- (iii) the “yaw” angle ψ , i.e. the rotation angle from the N axis to the X axis.



Although Euler angles are more straightforward to use than quaternions, they are also potentially subject to the “gimbal lock” problem:

https://en.wikipedia.org/wiki/Gimbal_lock

which arises whenever $\theta \simeq \pm 90^\circ$, including the common case when *the simulated coordinates are near the reference coordinates*. Therefore, a safe use of Euler angles as collective variables requires the use of restraints to avoid such singularities, such as done in reference [5] and in the [protein-ligand binding NAMD tutorial](#).

The eulerPhi component accepts exactly the same options as orientation, and measures the rotation angle from the x axis to the N axis. This angle is expressed in degrees within the periodic range $[-180^\circ : 180^\circ]$.

List of keywords (see also 4.14 for additional options):

- **Keyword** atoms: see definition of **atoms** (rmsd component)
- **Keyword** refPositions: see definition of **refPositions** (rmsd component)
- **Keyword** refPositionsFile: see definition of **refPositionsFile** (rmsd component)

4.6.7 eulerTheta: Pitch angle from references coordinates.

This component accepts exactly the same options as orientation, and measures the rotation angle from the P axis to the Z axis. This angle is expressed in degrees within the range $[-90^\circ : 90^\circ]$.

Warning: When this angle reaches -90° or 90° , the definition of orientation by euler angles suffers from the gimbal lock issue. You may need to apply a restraint to keep eulerTheta away from the singularities.

List of keywords (see also 4.14 for additional options):

- **Keyword** atoms: see definition of **atoms** (rmsd component)
- **Keyword** refPositions: see definition of **refPositions** (rmsd component)
- **Keyword** refPositionsFile: see definition of **refPositionsFile** (rmsd component)

4.6.8 eulerPsi: Yaw angle from references coordinates.

This component accepts exactly the same options as `orientation`, and measures the rotation angle from the N axis to the X axis. This angle is expressed in degrees within the periodic range $[-180^\circ : 180^\circ]$.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms`: see definition of `atoms` (rmsd component)
- **Keyword** `refPositions`: see definition of `refPositions` (rmsd component)
- **Keyword** `refPositionsFile`: see definition of `refPositionsFile` (rmsd component)

4.7 Raw data: building blocks for custom functions

4.7.1 cartesian: vector of atomic Cartesian coordinates.

The `cartesian { ... }` block defines a component returning a flat vector containing the Cartesian coordinates of all participating atoms, in the order $(x_1, y_1, z_1, \dots, x_n, y_n, z_n)$.

List of keywords (see also 4.14 for additional options):

- **Keyword** `atoms` \langle Group of atoms \rangle
Context: `cartesian`
Acceptable values: Block `atoms { ... }`
Description: Defines the atoms whose coordinates make up the value of the component. If `rotateToReference`, `centerToReference`, or `centerToOrigin` are defined, coordinates are evaluated within the moving frame of reference.

4.7.2 distancePairs: set of pairwise distances between two groups.

The `distancePairs { ... }` block defines a $N_1 \times N_2$ -dimensional variable that includes all mutual distances between the atoms of two groups. This can be useful, for example, to develop a new variable defined over two groups, by using the `scriptedFunction` feature.

List of keywords (see also 4.14 for additional options):

- **Keyword** `group1`: see definition of `group1` (distance component)
- **Keyword** `group2`: analogous to `group1`

- **Keyword** forceNoPBC: see definition of forceNoPBC (distance component)

This component returns a $N_1 \times N_2$ -dimensional vector of numbers, each ranging from 0 to the largest possible distance within the chosen boundary conditions.

4.8 Geometric path collective variables

The geometric path collective variables define the progress along a path, s , and the distance from the path, z . These CVs are proposed by Leines and Ensing[6], which differ from that[7] proposed by Branduardi et al., and utilize a set of geometric algorithms. The path is defined as a series of frames in the atomic Cartesian coordinate space or the CV space. s and z are computed as

$$s = \frac{m}{M} \pm \frac{1}{2M} \left(\frac{\sqrt{(\mathbf{v}_1 \cdot \mathbf{v}_3)^2 - |\mathbf{v}_3|^2(|\mathbf{v}_1|^2 - |\mathbf{v}_2|^2)} - (\mathbf{v}_1 \cdot \mathbf{v}_3)}{|\mathbf{v}_3|^2} - 1 \right) \quad (10)$$

$$z = \sqrt{\left(\mathbf{v}_1 + \frac{1}{2} \left(\frac{\sqrt{(\mathbf{v}_1 \cdot \mathbf{v}_3)^2 - |\mathbf{v}_3|^2(|\mathbf{v}_1|^2 - |\mathbf{v}_2|^2)} - (\mathbf{v}_1 \cdot \mathbf{v}_3)}{|\mathbf{v}_3|^2} - 1 \right) \mathbf{v}_4 \right)^2} \quad (11)$$

where $\mathbf{v}_1 = \mathbf{s}_m - \mathbf{z}$ is the vector connecting the current position to the closest frame, $\mathbf{v}_2 = \mathbf{z} - \mathbf{s}_{m-1}$ is the vector connecting the second closest frame to the current position, $\mathbf{v}_3 = \mathbf{s}_{m+1} - \mathbf{s}_m$ is the vector connecting the closest frame to the third closest frame, and $\mathbf{v}_4 = \mathbf{s}_m - \mathbf{s}_{m-1}$ is the vector connecting the second closest frame to the closest frame. m and M are the current index of the closest frame and the total number of frames, respectively. If the current position is on the left of the closest reference frame, the \pm in s turns to the positive sign. Otherwise it turns to the negative sign.

The equations above assume: (i) the frames are equidistant and (ii) the second and the third closest frames are neighbouring to the closest frame. When these assumptions are not satisfied, this set of path CV should be used carefully.

4.8.1 gspath: progress along a path defined in atomic Cartesian coordinate space.

In the gspath {...} and the gzpath {...} block all vectors, namely \mathbf{z} and \mathbf{s}_k are defined in atomic Cartesian coordinate space. More specifically, $\mathbf{z} = [\mathbf{r}_1, \dots, \mathbf{r}_n]$, where \mathbf{r}_i is the i -th atom specified in the atoms block. $\mathbf{s}_k = [\mathbf{r}_{k,1}, \dots, \mathbf{r}_{k,n}]$, where $\mathbf{r}_{k,i}$ means the i -th atom of the k -th reference frame.

List of keywords (see also 4.14 for additional options):

- **Keyword** atoms <Group of atoms>
Context: gspath and gzpath
Acceptable values: Block atoms {...}
Description: Defines the atoms whose coordinates make up the value of the component.
- **Keyword** refPositionsCol <PDB column containing atom flags>
Context: gspath and gzpath
Acceptable values: 0, B, X, Y, or Z
Description: If refPositionsFileN is a PDB file that contains all the atoms in the topology, this option may be provided to set which PDB field is used to flag the reference coordinates for atoms.

- **Keyword** `refPositionsFileN` \langle File containing the reference positions for fitting \rangle
Context: `gspath` and `gzpath`
Acceptable values: UNIX filename
Description: The path is defined by multiple `refPositionsFiles` which are similar to `refPositionsFile` in the `rmsd CV`. If your path consists of 10 nodes, you can list the coordinate file (in PDB or XYZ format) from `refPositionsFile1` to `refPositionsFile10`.
- **Keyword** `useSecondClosestFrame` \langle Define s_{m-1} as the second closest frame? \rangle
Context: `gspath` and `gzpath`
Acceptable values: boolean
Default value: on
Description: The definition assumes the second closest frame is neighbouring to the closest frame. This is not always true especially when the path is crooked. If this option is set to on (default), s_{m-1} is defined as the second closest frame. If this option is set to off, s_{m-1} is defined as the left or right neighbouring frame of the closest frame.
- **Keyword** `useThirdClosestFrame` \langle Define s_{m+1} as the third closest frame? \rangle
Context: `gspath` and `gzpath`
Acceptable values: boolean
Default value: off
Description: The definition assumes the third closest frame is neighbouring to the closest frame. This is not always true especially when the path is crooked. If this option is set to on, s_{m+1} is defined as the third closest frame. If this option is set to off (default), s_{m+1} is defined as the left or right neighbouring frame of the closest frame.
- **Keyword** `fittingAtoms` \langle The atoms that are used for alignment \rangle
Context: `gspath` and `gzpath`
Acceptable values: Group of atoms
Description: Before calculating v_1 , v_2 , v_3 and v_4 , the current frame need to be aligned to the corresponding reference frames. This option specifies which atoms are used to do alignment.

4.8.2 `gzpath`: distance from a path defined in atomic Cartesian coordinate space.

List of keywords (see also 4.14 for additional options):

- **Keyword** `useZsquare` \langle Compute z^2 instead of z \rangle
Context: `gzpath`
Acceptable values: boolean
Default value: off
Description: z is not differentiable when it is zero. This implementation workarounds it by setting the derivative of z to zero when $z = 0$. Another workaround is set this option to on, which computes z^2 instead of z , and then z^2 is differentiable when it is zero.

The usage of `gzpath` and `gspath` is illustrated below:

```

colvar {
  # Progress along the path
  name gs
  # The path is defined by 5 reference frames (from string-00.pdb to string-04.pdb)
  # Use atomic coordinate from atoms 1, 2 and 3 to compute the path
  gspath {
    atoms {atomnumbers { 1 2 3 }}
    refPositionsFile1 string-00.pdb
    refPositionsFile2 string-01.pdb
    refPositionsFile3 string-02.pdb
    refPositionsFile4 string-03.pdb
    refPositionsFile5 string-04.pdb
  }
}
colvar {
  # Distance from the path
  name gz
  # The path is defined by 5 reference frames (from string-00.pdb to string-04.pdb)
  # Use atomic coordinate from atoms 1, 2 and 3 to compute the path
  gzpath {
    atoms {atomnumbers { 1 2 3 }}
    refPositionsFile1 string-00.pdb
    refPositionsFile2 string-01.pdb
    refPositionsFile3 string-02.pdb
    refPositionsFile4 string-03.pdb
    refPositionsFile5 string-04.pdb
  }
}

```

4.8.3 linearCombination: Helper CV to define a linear combination of other CVs

This is a helper CV which can be defined as a linear combination of other CVs. It may be useful when you want to define the gspathCV {...} and the gzpathCV {...} as combinations of other CVs.

4.8.4 customColvar: Helper CV to define a mathematical expression as CV from other CVs

This is a helper CV which can be defined as a mathematical expression (see 4.15) of other CVs. It may be useful when you want to define the gspathCV {...}, the gzpathCV {...} and NeuralNetwork {...} as combinations of other CVs.

4.8.5 gspathCV: progress along a path defined in CV space.

In the `gspathCV { ... }` and the `gzpathCV { ... }` block all vectors, namely \mathbf{z} and \mathbf{s}_k are defined in CV space. More specifically, $\mathbf{z} = [\xi_1, \dots, \xi_n]$, where ξ_i is the i -th CV. $\mathbf{s}_k = [\xi_{k,1}, \dots, \xi_{k,n}]$, where $\xi_{k,i}$ means the i -th CV of the k -th reference frame. It should be note that these two CVs requires the `pathFile` option, which specifies a path file. Each line in the path file contains a set of space-seperated CV value of the reference frame. The sequence of reference frames matches the sequence of the lines.

List of keywords (see also 4.14 for additional options):

- **Keyword** `useSecondClosestFrame` \langle Define \mathbf{s}_{m-1} as the second closest frame? \rangle
Context: `gspathCV` and `gzpathCV`
Acceptable values: boolean
Default value: on
Description: The definition assumes the second closest frame is neighbouring to the closest frame. This is not always true especially when the path is crooked. If this option is set to on (default), \mathbf{s}_{m-1} is defined as the second closest frame. If this option is set to off, \mathbf{s}_{m-1} is defined as the left or right neighbouring frame of the closest frame.
- **Keyword** `useThirdClosestFrame` \langle Define \mathbf{s}_{m+1} as the third closest frame? \rangle
Context: `gspathCV` and `gzpathCV`
Acceptable values: boolean
Default value: off
Description: The definition assumes the third closest frame is neighbouring to the closest frame. This is not always true especially when the path is crooked. If this option is set to on, \mathbf{s}_{m+1} is defined as the third closest frame. If this option is set to off (default), \mathbf{s}_{m+1} is defined as the left or right neighbouring frame of the closest frame.
- **Keyword** `pathFile` \langle The file name of the path file. \rangle
Context: `gspathCV` and `gzpathCV`
Acceptable values: UNIX filename
Description: Defines the nodes or images that constitutes the path in CV space. The CVs of an image are listed in a line of `pathFile` using space-seperated format. Lines from top to button in `pathFile` corresponds images from initial to last.

4.8.6 gzpathCV: distance from a path defined in CV space.

List of keywords (see also 4.14 for additional options):

- **Keyword** `useZsquare` \langle Compute z^2 instead of z \rangle
Context: `gzpathCV`
Acceptable values: boolean
Default value: off
Description: z is not differentiable when it is zero. This implementation workarounds it by setting the derivative of z to zero when $z = 0$. Another workaround is set this option to on, which computes z^2 instead of z , and then z^2 is differentiable when it is zero.

The usage of gzpathCV and gspathCV is illustrated below:

```
colvar {  
  # Progress along the path  
  name gs  
  # Path defined by the CV space of two dihedral angles  
  gspathCV {  
    pathFile ./path.txt  
    dihedral {  
      name 001  
      group1 {atomNumbers {5}}  
      group2 {atomNumbers {7}}  
      group3 {atomNumbers {9}}  
      group4 {atomNumbers {15}}  
    }  
    dihedral {  
      name 002  
      group1 {atomNumbers {7}}  
      group2 {atomNumbers {9}}  
      group3 {atomNumbers {15}}  
      group4 {atomNumbers {17}}  
    }  
  }  
}  
  
colvar {  
  # Distance from the path  
  name gz  
  gzpathCV {  
    pathFile ./path.txt  
    dihedral {  
      name 001  
      group1 {atomNumbers {5}}  
      group2 {atomNumbers {7}}  
      group3 {atomNumbers {9}}  
      group4 {atomNumbers {15}}  
    }  
    dihedral {  
      name 002  
      group1 {atomNumbers {7}}  
      group2 {atomNumbers {9}}  
      group3 {atomNumbers {15}}  
      group4 {atomNumbers {17}}  
    }  
  }  
}
```

4.9 Arithmetic path collective variables

The arithmetic path collective variable in CV space uses a similar formula as the one proposed by Branduardi[7] et al., except that it computes s and z in CV space instead of RMSDs in Cartesian space. Moreover, this implementation allows different coefficients for each CV components as described in [8]. Assuming a path is composed of N reference frames and defined in an M -dimensional CV space, then the equations of s and z of the path are

$$s = \frac{1}{N-1} \frac{\sum_{i=0}^{N-1} \exp\left(-\lambda \sum_{j=1}^M c_j^2 (x_j - x_{i,j})^2\right)}{\sum_{i=0}^{N-1} \exp\left(-\lambda \sum_{j=1}^M c_j^2 (x_j - x_{i,j})^2\right)} \quad (12)$$

$$z = -\frac{1}{\lambda} \ln \left(\sum_{i=0}^{N-1} \exp \left(-\lambda \sum_{j=1}^M c_j^2 (x_j - x_{i,j})^2 \right) \right) \quad (13)$$

where c_j is the coefficient(weight) of the j -th CV, $x_{i,j}$ is the value of j -th CV of i -th reference frame and x_j is the value of j -th CV of current frame. λ is a parameter to smooth the variation of s and z . It should be noted that the index i ranges from 0 to $N-1$, and the definition of s is normalized by $1/(N-1)$. Consequently, the scope of s is $[0 : 1]$.

4.9.1 aspathCV: progress along a path defined in CV space.

This colvar component computes the s variable.

List of keywords (see also 4.14 for additional options):

- **Keyword weights** \langle Coefficients of the collective variables \rangle
Context: aspathCV and azpathCV
Acceptable values: space-separated numbers in a $\{ \dots \}$ block
Default value: $\{ 1.0 \dots \}$
Description: Define the coefficients. The j -th value in the $\{ \dots \}$ block corresponds the c_j in the equations.
- **Keyword lambda** \langle Smoothness of the variation of s and z \rangle
Context: aspathCV and azpathCV
Acceptable values: decimal
Default value: inverse of the mean square displacements of successive reference frames
Description: The value of λ in the equations.
- **Keyword pathFile** \langle The file name of the path file. \rangle
Context: aspathCV and azpathCV
Acceptable values: UNIX filename
Description: Defines the nodes or images that constitutes the path in CV space. The CVs of an image are listed in a line of pathFile using space-separated format. Lines from top to bottom in pathFile corresponds images from initial to last.

4.9.2 azpathCV: distance from a path defined in CV space.

This colvar component computes the z variable. Options are the same as in [4.9.1](#). The usage of azpathCV and aspathCV is illustrated below:

```
colvar {
  # Progress along the path
  name as
  # Path defined by the CV space of two dihedral angles
  aspathCV {
    pathFile ./path.txt
    weights {1.0 1.0}
    lambda 0.005
    dihedral {
      name 001
      group1 {atomNumbers {5}}
      group2 {atomNumbers {7}}
      group3 {atomNumbers {9}}
      group4 {atomNumbers {15}}
    }
    dihedral {
      name 002
      group1 {atomNumbers {7}}
      group2 {atomNumbers {9}}
      group3 {atomNumbers {15}}
      group4 {atomNumbers {17}}
    }
  }
}

colvar {
  # Distance from the path
  name az
  azpathCV {
    pathFile ./path.txt
    weights {1.0 1.0}
    lambda 0.005
    dihedral {
      name 001
      group1 {atomNumbers {5}}
      group2 {atomNumbers {7}}
      group3 {atomNumbers {9}}
      group4 {atomNumbers {15}}
    }
    dihedral {
```

```

name 002
group1 {atomNumbers {7}}
group2 {atomNumbers {9}}
group3 {atomNumbers {15}}
group4 {atomNumbers {17}}
}
}
}

```

4.10 Dense neural network in CV space

This colvar computes a non-linear combination of other scalar colvar components, and the non-linear transformation is defined by a dense neural network.[9]




Figure 1: Graphical representation of an example dense neural network with two hidden layers y_1 and y_2 . Both the input layer x and the output layer z have two nodes. The input nodes can be any existing scalar CVs.

The output of the j -th node of a k -th layer that has N_k nodes is computed by

$$y_{k,j} = f_k \left(\sum_{i=1}^{N_{k-1}} w_{(k,j),(k-1,i)} y_{k-1,i} + b_{k,j} \right), \quad (14)$$

where f_k is the activation function of the k -th layer, $w_{(k,j),(k-1,i)}$ is the weight of j -th node with respect to the i -th output of previous layer, and $b_{k,j}$ is the bias of j -th node of k -th layer.

List of keywords (see also 4.14 for additional options):

- **Keyword** `output_component` <The j -th node of the output or the last layer>
Context: NeuralNetwork
Acceptable values: integer starting from 0
Description: The value of this option specifies the output node to be used as the value of this CV.
- **Keyword** `layer i _WeightsFile` <The weights from layer $i - 1$ to layer i >
Context: NeuralNetwork
Acceptable values: UNIX filename
Description: The letter i in this option needs to be replaced with the indexing number starting from 1, for example, `layer1_WeightsFile` and `layer2_WeightsFile`. The value of this option specifies a

plain text file containing the weights from layer $i - 1$ to layer i . In the file, the number at k -th column and l -th row represents the weight from node k at layer $i - 1$ to node l at layer i .

- Keyword** `layeri_BiasesFile` ⟨The biases from layer $i - 1$ to layer i ⟩
Context: NeuralNetwork
Acceptable values: UNIX filename
Description: The letter i in this option needs to be replaced with the indexing number starting from 1, for example, `layer1_BiasesFile` and `layer2_BiasesFile`. The value of this option specifies a plain text file containing the weights of layer i . The file should have only one column, where the number at l -th row represents the bias of node l from layer $i - 1$ to layer i .
- Keyword** `layeri_activation` ⟨The activation function from layer $i - 1$ to layer i ⟩
Context: NeuralNetwork
Acceptable values: `tanh`, `sigmoid`, `linear`, `relu`, `lrelu100`, `elu`
Description: The letter i in this option needs to be replaced with the indexing number starting from 1, for example, `layer1_activation` and `layer2_activation`. The activation function from layer $i - 1$ to layer i . Available choices are `tanh`, `sigmoid`, `linear` (identity), `relu`, `lrelu100` (a leaky rely with coefficients 10^{-2} and 1), and `elu` (with coefficient 1).
- Keyword** `layeri_custom_activation` ⟨An alternative custom expression as the activation function from layer $i - 1$ to layer i ⟩
Context: NeuralNetwork
Acceptable values: string
Description: Mathematical expression to define the activation function from layer $i - 1$ to layer i . The input value must be written as `x`. For example, the ELU activation function can be expressed as `select(step(x), alpha*(exp(x)-1), x)`. It should be noted that this option is mutually exclusive with the option `colvar|NeuralNetwork|layeri_activation`

An example of configuration using NeuralNetwork is shown below:

```
colvar {
  # Define a neural network with 2 layers
  # The inputs are two torsion angles
  # and the first node at the output layer is used as the final CV
  name nn_output_1
  NeuralNetwork {
    output_component 0
    layer1_WeightsFile dense_1_weights.txt
    layer1_BiasesFile dense_1_biases.txt
    layer1_activation tanh
    layer2_WeightsFile dense_2_weights.txt
    layer2_BiasesFile dense_2_biases.txt
    layer2_activation tanh
    # The component coefficient is used for normalization
    componentCoeff 180.0
    dihedral {
      name 001
      # normalization factor 1.0/180.0
      componentCoeff 0.00555555555555555556
    }
  }
}
```

```

    group1 {atomNumbers {5}}
    group2 {atomNumbers {7}}
    group3 {atomNumbers {9}}
    group4 {atomNumbers {15}}
  }
  dihedral {
    name 002
    # normalization factor 1.0/180.0
    componentCoeff 0.0055555555555555556
    group1 {atomNumbers {7}}
    group2 {atomNumbers {9}}
    group3 {atomNumbers {15}}
    group4 {atomNumbers {17}}
  }
}

```

4.11 Shared keywords for all components

The following options can be used for any of the above colvar components in order to obtain a polynomial combination.

- Keyword name** <Name of this component>
Context: any component
Acceptable values: string
Default value: type of component + numeric id
Description: The name is a unique case-sensitive string which allows the Colvars module to identify this component. This is useful, for example, when combining multiple components via a scriptedFunction.
- Keyword scalable** <Attempt to calculate this component in parallel?>
Context: any component
Acceptable values: boolean
Default value: on, if available
Description: If set to on (default), the Colvars module will attempt to calculate this component in parallel to reduce overhead. Whether this option is available depends on the type of component: currently supported are distance, distanceZ, distanceXY, distanceVec, distanceDir, angle and dihedral. This flag influences computational cost, but does not affect numerical results: therefore, it should only be turned off for debugging or testing purposes.

4.12 Periodic components

The following components return real numbers that lie in a periodic interval:

- dihedral: torsional angle between four groups;

- `spinAngle`: angle of rotation around a predefined axis in the best-fit from a set of reference coordinates.

In certain conditions, `distanceZ` can also be periodic, namely when periodic boundary conditions (PBCs) are defined in the simulation and `distanceZ`'s axis is parallel to a unit cell vector.

In addition, a custom scalar colvar may be periodic depending on its user-defined expression. It will only be treated as such by the Colvars module if the period is specified using the `period` keyword, while `wrapAround` is optional.

The following keywords can be used within periodic components, or within custom variables (4.15)).

- **Keyword `period`** `<Period of the component>`
Context: `distanceZ`, custom colvars
Acceptable values: positive decimal
Default value: 0.0
Description: Setting this number enables the treatment of `distanceZ` as a periodic component: by default, `distanceZ` is not considered periodic. The keyword is supported, but irrelevant within dihedral or `spinAngle`, because their period is always 360 degrees.
- **Keyword `wrapAround`** `<Center of the wrapping interval for periodic variables>`
Context: `distanceZ`, dihedral, `spinAngle`, custom colvars
Acceptable values: decimal
Default value: 0.0
Description: By default, values of the periodic components are centered around zero, ranging from $-P/2$ to $P/2$, where P is the period. Setting this number centers the interval around this value. This can be useful for convenience of output, or to set the walls for a harmonicWalls in an order that would not otherwise be allowed.

Internally, all differences between two values of a periodic colvar follow the minimum image convention: they are calculated based on the two periodic images that are closest to each other.

Note: linear or polynomial combinations of periodic components (see 4.14) may become meaningless when components cross the periodic boundary. Use such combinations carefully: estimate the range of possible values of each component in a given simulation, and make use of `wrapAround` to limit this problem whenever possible.

4.13 Non-scalar components

When one of the following components are used, the defined colvar returns a value that is not a scalar number:

- `distanceVec`: 3-dimensional vector of the distance between two groups;
- `distanceDir`: 3-dimensional unit vector of the distance between two groups;
- `orientation`: 4-dimensional unit quaternion representing the best-fit rotation from a set of reference coordinates.

The distance between two 3-dimensional unit vectors is computed as the angle between them. The distance between two quaternions is computed as the angle between the two 4-dimensional unit vectors: because the orientation represented by q is the same as the one represented by $-q$, distances between two quaternions are computed considering the closest of the two symmetric images.

Non-scalar components carry the following restrictions:

- Calculation of total forces (outputTotalForce option) is currently not implemented.
- Each colvar can only contain one non-scalar component.
- Binning on a grid (abf, histogram and metadynamics with useGrids enabled) is currently not implemented for colvars based on such components.

Note: while these restrictions apply to individual colvars based on non-scalar components, no limit is set to the number of scalar colvars. To compute multi-dimensional histograms and PMFs, use sets of scalar colvars of arbitrary size.

4.13.1 Calculating total forces

In addition to the restrictions due to the type of value computed (scalar or non-scalar), a final restriction can arise when calculating total force (outputTotalForce option or application of a abf bias). total forces are available currently only for the following components: distance, distanceZ, distanceXY, angle, dihedral, rmsd, eigenvector and gyration.

4.14 Linear and polynomial combinations of components

To extend the set of possible definitions of colvars $\xi(\mathbf{r})$, multiple components $q_i(\mathbf{r})$ can be summed with the formula:

$$\xi(\mathbf{r}) = \sum_i c_i [q_i(\mathbf{r})]^{n_i} \quad (15)$$

where each component appears with a unique coefficient c_i (1.0 by default) the positive integer exponent n_i (1 by default).

Any set of components can be combined within a colvar, provided that they return the same type of values (scalar, unit vector, vector, or quaternion). By default, the colvar is the sum of its components. Linear or polynomial combinations (following equation (15)) can be obtained by setting the following parameters, which are common to all components:

- **Keyword** componentCoeff \langle Coefficient of this component in the colvar \rangle
Context: any component
Acceptable values: decimal
Default value: 1.0
Description: Defines the coefficient by which this component is multiplied (after being raised to componentExp) before being added to the sum.
- **Keyword** componentExp \langle Exponent of this component in the colvar \rangle
Context: any component
Acceptable values: integer
Default value: 1
Description: Defines the power at which the value of this component is raised before being added to the sum. When this exponent is different than 1 (non-linear sum), total forces and the Jacobian force are not available, making the colvar unsuitable for ABF calculations.

Example: To define the *average* of a colvar across different parts of the system, simply define within the same colvar block a series of components of the same type (applied to different atom groups), and assign to each component a componentCoeff of $1/N$.

4.15 Custom functions

Collective variables may be defined by specifying a custom function as an analytical expression. The expression is parsed by Lepton, the lightweight expression parser written by Peter Eastman (<https://simtk.org/projects/lepton>). Lepton produces efficient evaluation routines for the function and its derivatives.

- **Keyword** `customFunction` \langle Compute colvar as a custom function of its components \rangle
Context: colvar
Acceptable values: string
Description: Mathematical expression to define the colvar as a closed-form function of its colvar components. See below for the detailed syntax of Lepton expressions. Multiple mentions of this keyword can be used to define a vector variable (as in the example below).
- **Keyword** `customFunctionType` \langle Type of value returned by the scripted colvar \rangle
Context: colvar
Acceptable values: string
Default value: scalar
Description: With this flag, the user may specify whether the colvar is a scalar or one of the following vector types: `vector3` (a 3D vector), `unit_vector3` (a normalized 3D vector), or `unit_quaternion` (a normalized quaternion), or `vector`. Note that the scalar and vector cases are not necessary, as they are detected automatically.

The expression may use the collective variable components as variables, referred to by their user-defined name. Scalar elements of vector components may be accessed by appending a 1-based index to their name, as in the example below. When implementing generic functions of Cartesian coordinates rather than functions of existing components, the cartesian component may be particularly useful. A scalar-valued custom variable may be manually defined as periodic by providing the keyword `period`, and the optional keyword `wrapAround`, with the same meaning as in periodic components (see 4.12 for details). A vector variable may be defined by specifying the `customFunction` parameter several times: each expression defines one scalar element of the vector colvar, as in this example:

```
colvar {
  name custom

  # A 2-dimensional vector function of a scalar x and a 3-vector r
  customFunction cos(x) * (r1 + r2 + r3)
  customFunction sqrt(r1 * r2)

  distance {
    name x
    group1 { atomNumbers 1 }
    group2 { atomNumbers 50 }
  }
  distanceVec {
    name r
    group1 { atomNumbers 10 11 12 }
    group2 { atomNumbers 20 21 22 }
```

```
}
}
```

Numeric constants may be given in either decimal or exponential form (e.g. 3.12e-2). An expression may be followed by definitions for intermediate values that appear in the expression, separated by semicolons. For example, the expression:

$a^2 + a*b + b^2$; $a = a1 + a2$; $b = b1 + b2$

is exactly equivalent to:

$(a1 + a2)^2 + (a1 + a2) * (b1 + b2) + (b1 + b2)^2$.

The definition of an intermediate value may itself involve other intermediate values. All uses of a value must appear *before* that value's definition.

Lepton supports the usual arithmetic operators +, -, *, /, and ^ (power), as well as the following functions:

sqrt	Square root
exp	Exponential
log	Natural logarithm
erf	Error function
erfc	Complementary error function
sin	Sine (angle in radians)
cos	Cosine (angle in radians)
sec	Secant (angle in radians)
csc	Cosecant (angle in radians)
tan	Tangent (angle in radians)
cot	Cotangent (angle in radians)
asin	Inverse sine (in radians)
acos	Inverse cosine (in radians)
atan	Inverse tangent (in radians)
atan2	Two-argument inverse tangent (in radians)
sinh	Hyperbolic sine
cosh	Hyperbolic cosine
tanh	Hyperbolic tangent
abs	Absolute value
floor	Floor
ceil	Ceiling
min	Minimum of two values
max	Maximum of two values
delta	$\delta(x) = 1$ if $x = 0$, 0 otherwise
step	$\text{step}(x) = 0$ if $x < 0$, 1 if $x \geq 0$
select	$\text{select}(x, y, z) = z$ if $x = 0$, y otherwise

4.16 Defining grid parameters

Many algorithms require the definition of boundaries and/or characteristic spacings that can be used to define discrete “states” in the collective variable, or to combine variables with very different units. The parameters described below offer a way to specify these parameters only once for each variable, while using them multiple times in restraints, time-dependent biases or analysis methods.

- **Keyword** `width` ⟨Unit of the variable, or grid spacing⟩
Context: `colvar`
Acceptable values: positive decimal
Default value: 1.0
Description: This number defines the effective unit of measurement for the collective variable, and is used by the biasing methods for the following purposes. Harmonic (6.5), harmonic walls (6.7) and linear restraints (6.8) use it to set the physical unit of the force constant, which is useful for multidimensional restraints involving multiple variables with very different units (for examples, *length unit* or degrees °) with a single, scaled force constant. The values of the scaled force constant in the units of each variable are printed at initialization time. Histograms (6.10), ABF (6.2) and metadynamics (6.4) all use this number as the initial choice for the grid spacing along this variable: for this reason, `width` should generally be no larger than the standard deviation of the `colvar` in an unbiased simulation. Unless it is required to control the spacing, it is usually simplest to keep the default value of 1, so that restraint force constants are provided with their full physical unit.
- **Keyword** `lowerBoundary` ⟨Lower boundary of the `colvar`⟩
Context: `colvar`
Acceptable values: decimal
Default value: natural boundary of the function
Description: Defines the lowest end of the interval of “relevant” values for the variable. This number can be, for example, a true physical boundary imposed by the choice of function (e.g. the distance function is always larger than zero): if this is the case, and only one function is used to define the variable, the default value of this number is set to the lowest end of the range of values of that function, if available (see Section 4.1). Alternatively, this value may be provided by the user, to represent for example the left-most point of a PMF calculation along this variable. In the latter case, it is the user’s responsibility to either (a) ensure the variable does not go significantly beyond the boundary (for example by adding a `harmonicWalls` restraint, 6.7), or (b) instruct the code that this is a true physical boundary by setting `hardLowerBoundary`.
- **Keyword** `upperBoundary` ⟨Upper boundary of the `colvar`⟩
Context: `colvar`
Acceptable values: decimal
Default value: natural boundary of the function
Description: Similarly to `lowerBoundary`, defines the highest of the “relevant” values of the variable.
- **Keyword** `hardLowerBoundary` ⟨Whether the lower boundary is the physical lower limit⟩
Context: `colvar`
Acceptable values: boolean
Default value: provided by the component
Description: When the `colvar` has a “natural” boundary (for example, a distance `colvar` cannot go below 0) this flag is automatically enabled. For more complex variable definitions, or when `lowerBoundary` is provided directly by the user, it may be useful to set this flag explicitly. This option does not affect simulation results, but enables some internal optimizations by letting the code know that the variable is unable to cross the lower boundary, regardless of whether restraints are applied to it.
- **Keyword** `hardUpperBoundary` ⟨Whether the upper boundary is the physical upper limit of the `colvar`’s values⟩

Context: colvar

Acceptable values: boolean

Default value: provided by the component

Description: Analogous to hardLowerBoundary.

- **Keyword** `expandBoundaries` \langle Allow to expand the two boundaries if needed \rangle

Context: colvar

Acceptable values: boolean

Default value: off

Description: If defined, lowerBoundary and upperBoundary may be automatically expanded to accommodate colvar values that do not fit in the initial range. Currently, this option is used by the metadynamics bias (6.4) to keep all of its hills fully within the grid. This option cannot be used when the initial boundaries already span the full period of a periodic colvar.

4.17 Trajectory output

- **Keyword** `outputValue` \langle Output a trajectory for this colvar \rangle

Context: colvar

Acceptable values: boolean

Default value: on

Description: If colvarsTrajFrequency is non-zero, the value of this colvar is written to the trajectory file every colvarsTrajFrequency steps in the column labeled “<name>”.

- **Keyword** `outputVelocity` \langle Output a velocity trajectory for this colvar \rangle

Context: colvar

Acceptable values: boolean

Default value: off

Description: If colvarsTrajFrequency is defined, the finite-difference calculated velocity of this colvar are written to the trajectory file under the label “v_<name>”.

- **Keyword** `outputEnergy` \langle Output an energy trajectory for this colvar \rangle

Context: colvar

Acceptable values: boolean

Default value: off

Description: This option applies only to extended Lagrangian colvars. If colvarsTrajFrequency is defined, the kinetic energy of the extended degree and freedom and the potential energy of the restraining spring are written to the trajectory file under the labels “Ek_<name>” and “Ep_<name>”.

- **Keyword** `outputTotalForce` \langle Output a total force trajectory for this colvar \rangle

Context: colvar

Acceptable values: boolean

Default value: off

Description: If colvarsTrajFrequency is defined, the total force on this colvar (i.e. the projection of all atomic total forces onto this colvar — see equation (20) in section 6.2) are written to the trajectory file under the label “fs_<name>”. For extended Lagrangian colvars, the “total force” felt by the extended degree of freedom is simply the force from the harmonic spring. This force includes all of

the forces computed by the force field, as well as any forces by fixes defined prior to fix colvars. **Note:** not all components support this option. The physical unit for this force is *energy unit/(colvar unit)*.

- **Keyword** outputAppliedForce \langle Output an applied force trajectory for this colvar \rangle
Context: colvar
Acceptable values: boolean
Default value: off
Description: If colvarsTrajFrequency is defined, the total force applied on this colvar by Colvars biases are written to the trajectory under the label “fa_<name>”. For extended Lagrangian colvars, this force is actually applied to the extended degree of freedom rather than the geometric colvar itself. The physical unit for this force is *energy unit/(colvar unit)*.

4.18 Extended Lagrangian

The following options enable extended-system dynamics, where a colvar is coupled to an additional degree of freedom (fictitious particle) by a harmonic spring. This extended coordinate masks the colvar and replaces it transparently from the perspective of biasing and analysis methods. Biasing forces are then applied to the extended degree of freedom, and the actual geometric colvar (function of Cartesian coordinates) only feels the force from the harmonic spring. This is particularly useful when combined with an **abf** bias to perform eABF simulations ().

Note that for some biases (**harmonicWalls**, **histogram**), this masking behavior is controlled by the keyword **bypassExtendedLagrangian**. Specifically for harmonicWalls, the default behavior is to bypass extended Lagrangian coordinates and act directly on the actual colvars.

- **Keyword** extendedLagrangian \langle Add extended degree of freedom \rangle
Context: colvar
Acceptable values: boolean
Default value: off
Description: Adds a fictitious particle to be coupled to the colvar by a harmonic spring. The fictitious mass and the force constant of the coupling potential are derived from the parameters extendedTimeConstant and extendedFluctuation, described below. Biasing forces on the colvar are applied to this fictitious particle, rather than to the atoms directly. This implements the extended Lagrangian formalism used in some metadynamics simulations [3]. .
- **Keyword** extendedFluctuation \langle Standard deviation between the colvar and the fictitious particle (colvar unit) \rangle
Context: colvar
Acceptable values: positive decimal
Description: Defines the spring stiffness for the extendedLagrangian mode, by setting the typical deviation between the colvar and the extended degree of freedom due to thermal fluctuation. The spring force constant is calculated internally as $k_B T / \sigma^2$, where σ is the value of extendedFluctuation.
- **Keyword** extendedTimeConstant \langle Oscillation period of the fictitious particle (fs) \rangle
Context: colvar
Acceptable values: positive decimal

Default value: 200

Description: Defines the inertial mass of the fictitious particle, by setting the oscillation period of the harmonic oscillator formed by the fictitious particle and the spring. The period should be much larger than the MD time step to ensure accurate integration of the extended particle's equation of motion. The fictitious mass is calculated internally as $k_B T (\tau / 2\pi\sigma)^2$, where τ is the period and σ is the typical fluctuation (see above).

- **Keyword** `extendedTemp` \langle Temperature for the extended degree of freedom (K) \rangle
Context: colvar
Acceptable values: positive decimal
Default value: thermostat temperature
Description: Temperature used for calculating the coupling force constant of the extended variable (see `extendedFluctuation`) and, if needed, as a target temperature for extended Langevin dynamics (see `extendedLangevinDamping`). This should normally be left at its default value.
- **Keyword** `extendedLangevinDamping` \langle Damping factor for extended Langevin dynamics (ps^{-1}) \rangle
Context: colvar
Acceptable values: positive decimal
Default value: 1.0
Description: If this is non-zero, the extended degree of freedom undergoes Langevin dynamics at temperature `extendedTemp`. The friction force is minus `extendedLangevinDamping` times the velocity. This is useful because the extended dynamics coordinate may heat up in the transient non-equilibrium regime of ABF. Use moderate damping values, to limit viscous friction (potentially slowing down diffusive sampling) and stochastic noise (increasing the variance of statistical measurements). In doubt, use the default value.

4.19 Multiple time-step variables

- **Keyword** `timeStepFactor` \langle Compute this colvar once in a certain number of timesteps \rangle
Context: colvar
Acceptable values: positive integer
Default value: 1
Description: Instructs this colvar to activate at a time interval equal to the base (MD) timestep times `timeStepFactor`.^[10] At other time steps, the value of the variable is not updated, and no biasing forces are applied. Any forces exerted by biases are accumulated over the given time interval, then applied as an impulse at the next update.

4.20 Backward-compatibility

- **Keyword** `subtractAppliedForce` \langle Do not include biasing forces in the total force for this colvar \rangle
Context: colvar
Acceptable values: boolean
Default value: off

Description: If the colvar supports total force calculation (see 4.13.1), all forces applied to this colvar by biases will be removed from the total force. This keyword allows to recover some of the “system force” calculation available in the Colvars module before version 2016-08-10. Please note that removal of *all* other external forces (including biasing forces applied to a different colvar) is *no longer supported*, due to changes in the underlying simulation engines (primarily NAMD). This option may be useful when continuing a previous simulation where the removal of external/applied forces is essential. *For all new simulations, the use of this option is not recommended.*

4.21 Statistical analysis

Run-time calculations of statistical properties that depend explicitly on time can be performed for individual collective variables. Currently, several types of time correlation functions, running averages and running standard deviations are implemented. For run-time computation of histograms, please see the histogram bias (6.10).

- Keyword** `corrFunc` \langle Calculate a time correlation function? \rangle
Context: colvar
Acceptable values: boolean
Default value: off
Description: Whether or not a time correlation function should be calculated for this colvar.
- Keyword** `corrFuncWithColvar` \langle Colvar name for the correlation function \rangle
Context: colvar
Acceptable values: string
Description: By default, the auto-correlation function (ACF) of this colvar, ξ_i , is calculated. When this option is specified, the correlation function is calculated instead with another colvar, ξ_j , which must be of the same type (scalar, vector, or quaternion) as ξ_i .
- Keyword** `corrFuncType` \langle Type of the correlation function \rangle
Context: colvar
Acceptable values: velocity, coordinate or coordinate_p2
Default value: velocity
Description: With coordinate or velocity, the correlation function $C_{i,j}(t) = \langle \Pi(\xi_i(t_0), \xi_j(t_0 + t)) \rangle$ is calculated between the variables ξ_i and ξ_j , or their velocities. $\Pi(\xi_i, \xi_j)$ is the scalar product when calculated between scalar or vector values, whereas for quaternions it is the cosine between the two corresponding rotation axes. With coordinate_p2, the second order Legendre polynomial, $(3\cos(\theta)^2 - 1)/2$, is used instead of the cosine.
- Keyword** `corrFuncNormalize` \langle Normalize the time correlation function? \rangle
Context: colvar
Acceptable values: boolean
Default value: on
Description: If enabled, the value of the correlation function at $t = 0$ is normalized to 1; otherwise, it equals to $\langle O(\xi_i, \xi_j) \rangle$.
- Keyword** `corrFuncLength` \langle Length of the time correlation function \rangle
Context: colvar

Acceptable values: positive integer

Default value: 1000

Description: Length (in number of points) of the time correlation function.

- **Keyword** `corrFuncStride` \langle Stride of the time correlation function \rangle
Context: colvar
Acceptable values: positive integer
Default value: 1
Description: Number of steps between two values of the time correlation function.
- **Keyword** `corrFuncOffset` \langle Offset of the time correlation function \rangle
Context: colvar
Acceptable values: positive integer
Default value: 0
Description: The starting time (in number of steps) of the time correlation function (default: $t = 0$).
Note: *the value at $t = 0$ is always used for the normalization.*
- **Keyword** `corrFuncOutputFile` \langle Output file for the time correlation function \rangle
Context: colvar
Acceptable values: UNIX filename
Default value: `output.<name>.corrfunc.dat`
Description: The time correlation function is saved in this file.
- **Keyword** `runAve` \langle Calculate the running average and standard deviation \rangle
Context: colvar
Acceptable values: boolean
Default value: off
Description: Whether or not the running average and standard deviation should be calculated for this colvar.
- **Keyword** `runAveLength` \langle Length of the running average window \rangle
Context: colvar
Acceptable values: positive integer
Default value: 1000
Description: Length (in number of points) of the running average window.
- **Keyword** `runAveStride` \langle Stride of the running average window values \rangle
Context: colvar
Acceptable values: positive integer
Default value: 1
Description: Number of steps between two values within the running average window.
- **Keyword** `runAveOutputFile` \langle Output file for the running average and standard deviation \rangle
Context: colvar
Acceptable values: UNIX filename
Default value: `output.<name>.runave.traj`
Description: The running average and standard deviation are saved in this file.

5 Selecting atoms

To define collective variables, atoms are usually selected as groups. Each group is defined using an identifier that is unique in the context of the specific colvar component (e.g. for a distance component, the two groups are group1 and group2). The identifier is followed by a brace-delimited block containing selection keywords and other parameters, including an optional name:

- **Keyword** `name` \langle Unique name for the atom group \rangle
Context: atom group
Acceptable values: string
Description: This parameter defines a unique name for this atom group, which can be referred to in the definition of other atom groups (including in other colvars) by invoking `atomsOfGroup` as a selection keyword.

5.1 Atom selection keywords

Selection keywords may be used individually or in combination with each other, and each can be repeated any number of times. Selection is incremental: each keyword adds the corresponding atoms to the selection, so that different sets of atoms can be combined. However, atoms included by multiple keywords are only counted once. Below is an example configuration for an atom group called “atoms”. **Note:** *this is an unusually varied combination of selection keywords, demonstrating how they can be combined together: most simulations only use one of them.*

```
atoms {  
  
    # add atoms 1 and 3 to this group (note: first atom in the system is 1)  
    atomNumbers {  
        1 3  
    }  
  
    # add atoms starting from 20 up to and including 50  
    atomNumbersRange 20-50  
    # add index group (requires a .ndx file to be provided globally)  
    indexGroup Water  
}
```

The resulting selection includes atoms 1 and 3, those between 20 and 50, and those in the index group called “Water”. The indices of this group are read from the file provided by the global keyword `indexFile`. The complete list of selection keywords available in LAMMPS is:

- **Keyword** `atomNumbers` \langle List of atom numbers \rangle
Context: atom group

Acceptable values: space-separated list of positive integers

Description: This option adds to the group all the atoms whose numbers are in the list. *The number of the first atom in the system is 1: to convert from a VMD selection, use “atomselect get serial”.*

- **Keyword** `indexGroup` \langle Name of index group to be used (GROMACS format) \rangle
Context: atom group
Acceptable values: string
Description: If the name of an index file has been provided by `indexFile`, this option allows to select one index group from that file: the atoms from that index group will be used to define the current group.
- **Keyword** `atomsOfGroup` \langle Name of group defined previously \rangle
Context: atom group
Acceptable values: string
Description: Refers to a group defined previously using its user-defined name. This adds all atoms of that named group to the current group.
- **Keyword** `atomNumbersRange` \langle Atoms within a number range \rangle
Context: atom group
Acceptable values: \langle Starting number \rangle - \langle Ending number \rangle
Description: This option includes in the group all atoms whose numbers are within the range specified. *The number of the first atom in the system is 1.*
- **Keyword** `dummyAtom` \langle Dummy atom position (*length unit*) \rangle
Context: atom group
Acceptable values: (x, y, z) triplet
Description: Instead of selecting any atom, this option makes the group a virtual particle at a fixed position in space. This is useful e.g. to replace a group’s center of geometry with a user-defined position.

5.2 Moving frame of reference.

The following options define an automatic calculation of an optimal translation (`centerToReference`) or optimal rotation (`rotateToReference`), that superimposes the positions of this group to a provided set of reference coordinates. Alternately, `centerToOrigin` applies a translation to place the geometric center of the group at (0, 0, 0). This can allow, for example, to effectively remove from certain colvars the effects of molecular tumbling and of diffusion. Given the set of atomic positions \mathbf{x}_i , the colvar ξ can be defined on a set of roto-translated positions $\mathbf{x}'_i = R(\mathbf{x}_i - \mathbf{x}^C) + \mathbf{x}^{\text{ref}}$. \mathbf{x}^C is the geometric center of the \mathbf{x}_i , R is the optimal rotation matrix to the reference positions and \mathbf{x}^{ref} is the geometric center of the reference positions.

Components that are defined based on pairwise distances are naturally invariant under global roto-translations. Other components are instead affected by global rotations or translations: however, they can be made invariant if they are expressed in the frame of reference of a chosen group of atoms, using the `centerToReference` and `rotateToReference` options. Finally, a few components are defined by convention using a roto-translated frame (e.g. the minimal RMSD): for these components, `centerToReference` and `rotateToReference` are enabled by default. In typical applications, the default settings result in the expected behavior.

Warning on rotating frames of reference and periodic boundary conditions. `rotateToReference` affects coordinates that depend on minimum-image distances in periodic boundary conditions (PBC). After rotation of the coordinates, the periodic cell vectors become irrelevant: the rotated system is effectively non-periodic. A safe way to handle this is to ensure that the relevant inter-group distance vectors remain smaller than the half-size of the periodic cell. If this is not desirable, one should avoid the rotating frame of reference, and apply orientational restraints to the reference group instead, in order to keep the orientation of the reference group consistent with the orientation of the periodic cell.

Warning on rotating frames of reference and ABF. Note that `centerToReference` and `rotateToReference` may affect the Jacobian derivative of colvar components in a way that is not taken into account by default. Be careful when using these options in ABF simulations or when using total force values.

- **Keyword** `centerToReference` \langle Implicitly remove translations for this group \rangle
Context: atom group
Acceptable values: boolean
Default value: off
Description: If this option is on, the center of geometry of the group will be aligned with that of the reference positions provided by `refPositions` or `refPositionsFile`. Colvar components will only have access to the aligned positions. **Note:** unless otherwise specified, `rmsd` and `eigenvector` set this option to on *by default*.
- **Keyword** `centerToOrigin` \langle Implicitly remove translations for this group by keeping its center at the origin \rangle
Context: atom group
Acceptable values: boolean
Default value: off
Description: This option implies `centerToReference`. If this option is on, coordinates from the group will be translated so that the center of geometry of the group remains at (0, 0, 0), except if `fittingGroup` is enabled. In that case, the translation applied is the translation that brings the center of geometry of the fitting group to (0, 0, 0).
- **Keyword** `rotateToReference` \langle Implicitly remove rotations for this group \rangle
Context: atom group
Acceptable values: boolean
Default value: off
Description: If this option is on, the coordinates of this group will be optimally superimposed to the reference positions provided by `refPositions` or `refPositionsFile`. The rotation will be performed around the center of geometry if `centerToReference` is on, or around the origin otherwise. The algorithm used is the same employed by the orientation colvar component [4]. Forces applied to the atoms of this group will also be implicitly rotated back to the original frame. **Note:** unless otherwise specified, `rmsd` and `eigenvector` set this option to on *by default*.
- **Keyword** `refPositions` \langle Reference positions for fitting (*length unit*) \rangle
Context: atom group
Acceptable values: space-separated list of (x, y, z) triplets
Description: This option provides a list of reference coordinates for `centerToReference` and/or `rotateToReference`, and is mutually exclusive with `refPositionsFile`. If only `centerToReference` is on, the list may contain a single (x, y, z) triplet; if also `rotateToReference` is on, the list should be as long as the atom group, and *its order must match the order in which atoms were defined*.

- **Keyword** `refPositionsFile` ⟨File containing the reference positions for fitting⟩
Context: atom group
Acceptable values: UNIX filename
Description: This option provides a list of reference coordinates for `centerToReference` and/or `rotateToReference`, and is mutually exclusive with `refPositions`. The acceptable file format is XYZ (3.7.3), which is read in double precision. .
- **Keyword** `fittingGroup` ⟨Use an alternate set of atoms to define the roto-translation⟩
Context: atom group
Acceptable values: Block `fittingGroup { ... }`
Default value: This atom group itself
Description: If either `centerToReference` or `rotateToReference` is defined, this keyword defines an alternate atom group to calculate the optimal roto-translation. Use this option to define a continuous rotation if the structure of the group involved changes significantly (a typical symptom would be the message “Warning: discontinuous rotation!”). **Performance considerations:** note that enabling this option will result in projecting each of the atomic gradients of the colvar (e.g. the RMSD) onto each of the gradients of the roto-translation, which may be a computationally expensive operation: see the closely related `enableFitGradients` for details.

The following two options have default values appropriate for the vast majority of applications, and are only provided to support rare, special cases.

- **Keyword** `enableFitGradients` ⟨Include the roto-translational contribution to colvar gradients⟩
Context: atom group
Acceptable values: boolean
Default value: on
Description: When either `centerToReference` or `rotateToReference` is on, the gradients of some colvars include terms proportional to $\partial R / \partial \mathbf{x}_i$ (rotational gradients) and $\partial \mathbf{x}^C / \partial \mathbf{x}_i$ (translational gradients). By default, these terms are calculated and included in the total gradients; if this option is set to off, they are neglected. In the case of a minimum RMSD component, this flag is automatically disabled because the contributions of those derivatives to the gradients cancel out; other types of variable will require projecting each of the gradients of the variable onto each of the gradients of the roto-translation (i.e. a $O(N^2)$ loop). When `fittingGroup` is enabled, the computation is a $O(N \times M)$ loop for all variables, including RMSDs.

5.3 Treatment of periodic boundary conditions.

In simulations with periodic boundary conditions, many of the implemented colvar components rely on the fact that each position within a group of atoms is at the nearest periodic image from the center of geometry of the group itself. However, due to the internal wrapping of individual atomic positions done by LAMMPS, this assumption is broken if the group straddles one of the unit cell’s boundaries. For this reason, within the Colvars module all coordinates are unwrapped by default to avoid discontinuities (see `unwrap` keyword in 3.2).

The user should determine whether maintaining the default value of `unwrap`, depending on the specifics of each system. In general, internal coordinate wrapping by LAMMPS does not affect the calculation of colvars if each atom group satisfies one or more of the following:

- i)* it is composed by only one atom;
- ii)* it is used by a colvar component which does not make use of its center of geometry, but only of pairwise distances (`distanceInv`, `coordNum`, `hBond`, `alpha`, `dihedralPC`);
- iii)* it is used by a colvar component that ignores the ill-defined Cartesian components of its center of mass (such as the *x* and *y* components of a membrane's center of mass modeled with `distanceZ`).

5.4 Performance of a Colvars calculation based on group size.

In simulations performed with message-passing programs (such as NAMD or LAMMPS), the calculation of energy and forces is distributed (i.e., parallelized) across multiple nodes, as well as over the processor cores of each node. When Colvars is enabled, certain atomic coordinates are collected on a single node, where the calculation of collective variables and of their biases is executed. This means that for simulations over large numbers of nodes, a Colvars calculation may produce a significant overhead, coming from the costs of transmitting atomic coordinates to one node and of processing them.

Performance can be improved in multiple ways:

- The calculation of variables, components and biases can be distributed over the processor cores of the node where the Colvars module is executed. Currently, an equal weight is assigned to each colvar, or to each component of those colvars that include more than one component. The performance of simulations that use many colvars or components is improved automatically. For simulations that use a single large colvar, it may be advisable to partition it in multiple components, which will be then distributed across the available cores. In LAMMPS, this feature is supported automatically when LAMMPS is compiled with OpenMP support. If printed, the message “SMP parallelism is available.” indicates the availability of the option. If available, the option is turned on by default, but may be disabled using the keyword `smp` if required for debugging.
- As a general rule, the size of atom groups should be kept relatively small (up to a few thousands of atoms, depending on the size of the entire system in comparison). To gain an estimate of the computational cost of a large colvar, one can use a test calculation of the same colvar in VMD (hint: use the `time Tcl` command to measure the cost of running `cv update`).

6 Biasing and analysis methods

A biasing or analysis method can be applied to existing collective variables by using the following configuration:

```
<biastype> {  
  name <name>  
  colvars <xi1> <xi2> ...  
  <parameters>  
}
```

The keyword <biastype> indicates the method of choice. There can be multiple instances of the same method, e.g. using multiple harmonic blocks allows defining multiple restraints.

All biasing and analysis methods implemented recognize the following options:

- **Keyword name** ⟨Identifier for the bias⟩
Context: colvar bias
Acceptable values: string
Default value: <type of bias><bias index>
Description: This string is used to identify the bias or analysis method in the output, and to name some output files. **Tip:** although a unique name is assigned automatically, *you are strongly encouraged to give a name to a bias that you may want to analyze later*. For example, processing a trajectory file (sec. 3.7.5) containing a harmonic restraint is much simpler with a clearly identifiable name, like “smd” for a moving harmonic restraint or “us” for a static one.
- **Keyword colvars** ⟨Collective variables involved⟩
Context: colvar bias
Acceptable values: space-separated list of colvar names
Description: This option selects by name all the variables to which this bias or analysis will be applied.
- **Keyword outputEnergy** ⟨Write the current bias energy to the trajectory file⟩
Context: colvar bias
Acceptable values: boolean
Default value: off
Description: If this option is chosen and colvarsTrajFrequency is not zero, the current value of the biasing energy will be written to the trajectory file during the simulation.
- **Keyword outputFreq** ⟨Frequency (number of steps) at which output files are written⟩
Context: colvar bias
Acceptable values: positive integer
Default value: colvarsRestartFrequency
Description: If this bias produces aggregated data that needs to be written to disk (for example, a PMF), this number specifies the number of steps after which these data are written to files. A value of zero disables writing files for this bias during the simulation (except for outputEnergy, which is controlled by colvarsTrajFrequency). All output files are also written at the end of a simulation run, regardless of the value of this number.

- **Keyword** `bypassExtendedLagrangian` \langle Apply bias to actual colvars, bypassing extended coordinates \rangle
Context: colvar bias
Acceptable values: boolean
Default value: off
Description: This option is implemented by the `harmonicWalls` and `histogram` biases. It is only relevant if the bias is applied to one or several extended-Lagrangian colvars (4.18), for example within an eABF () simulation. Usually, biases use the value of the extended coordinate as a proxy for the actual colvar, and their biasing forces are applied to the extended coordinates as well. If `bypassExtendedLagrangian` is enabled, the bias behaves as if there were no extended coordinates, and accesses the value of the underlying colvars, applying any biasing forces along the gradients of those variables.
- **Keyword** `stepZeroData` \langle Accumulate data starting at step 0 of a simulation run \rangle
Context: colvar bias
Acceptable values: boolean
Default value: off
Description: This option is meaningful for biases that record and accumulate data during a simulation, such as ABF (6.2), metadynamics (6.4), histograms (6.10) and in general any bias that accumulates free-energy samples with thermodynamic integration, or TI (6.1). When this option is disabled (default), data will only be recorded into the bias after the first coordinate update: this is generally the correct choice in simulation runs. Biasing energy and forces will always be computed for all active biases, regardless of this option.
- **Keyword** `scaledBiasingForce` \langle Scale biasing force by a factor in an external histogram \rangle
Context: colvar bias
Acceptable values: boolean
Default value: off
Description: If this option is set to on, the biasing force at each step will be scaled by a factor provided in the grid of an external histogram. The histogram file is provided by `scaledBiasingForceFactorsGrid`. It ought to be noted that if the variables are not in any grids of the histogram, the scaling factor is 1.0.
- **Keyword** `scaledBiasingForceFactorsGrid` \langle A histogram file with the scaling factor of biasing force in each bin \rangle
Context: colvar bias
Acceptable values: string
Description: If `scaledBiasingForce` is set to on, this option accepts the filename of the histogram file that contains the scaling factors. The format of the histogram file is the same as the PMF files written by metadynamics or ABF.

6.1 Thermodynamic integration

The methods implemented here provide a variety of estimators of conformational free-energies. These are carried out at run-time, or with the use of post-processing tools over the generated output files. The specifics of each estimator are discussed in the documentation of each biasing or analysis method.

A special case is the traditional thermodynamic integration (TI) method, used for example to compute potentials of mean force (PMFs). Most types of restraints (6.5, 6.7, 6.8, ...) as well as metadynamics (6.4) can optionally use TI alongside their own estimator, based on the keywords documented below.

- **Keyword** `writeTIPMF` \langle Write the PMF computed by thermodynamic integration \rangle

Context: colvar bias

Acceptable values: boolean

Default value: off

Description: If the bias is applied to a variable that supports the calculation of total forces (see `outputtotalForce` and 4.13.1), this option allows calculating the corresponding PMF by thermodynamic integration, and writing it to the file `output.<name>.ti.pmf`, where `<name>` is the name of the bias and the contents of the file are in multicolumn text format (3.7.4). The total force includes the forces applied to the variable by all bias, except those from this bias itself. If any bias applies time-dependent forces besides the one using this option, an error is raised.

- **Keyword** `writeTISamples` \langle Write the free-energy gradient samples \rangle

Context: colvar bias

Acceptable values: boolean

Default value: off

Description: This option allows to compute total forces for use with thermodynamic integration as done by the keyword `writeTIPMF`. The names of the files containing the variables' histogram and mean thermodynamic forces are `output.<name>.ti.count` and `output.<name>.ti.force`, respectively: these can be used by `abf_integrate` (see 6.2.4) or similar utility. Note that because the `.force` file contains mean forces instead of free-energy gradients, `abf_integrate <filename> -s -1.0` should be used. This option is on by default when `writeTIPMF` is on, but can be enabled separately if the bias is applied to more than one variable, making not possible the direct integration of the PMF at runtime. If any bias applies time-dependent forces besides the one using this option, an error is raised.

In adaptive biasing force (ABF) (6.2) the above keywords are not recognized, because their functionality is either included already (conventional ABF) or not available (extended-system ABF).

6.2 Adaptive Biasing Force

For a full description of the Adaptive Biasing Force method, see reference [11]. For details about this implementation, see references [12] and [13]. **When publishing research that makes use of this functionality, please cite references [11] and [13].**

An alternate usage of this feature is the application of custom tabulated biasing potentials to one or more colvars. See `inputPrefix` and `updateBias` below.

Combining ABF with the extended Lagrangian feature (4.18) of the variables produces the extended-system ABF variant of the method ().

ABF is based on the thermodynamic integration (TI) scheme for computing free energy profiles. The free energy as a function of a set of collective variables $\xi = (\xi_i)_{i \in [1,n]}$ is defined from the canonical distribution of ξ , $\mathcal{P}(\xi)$:

$$A(\xi) = -\frac{1}{\beta} \ln \mathcal{P}(\xi) + A_0 \quad (16)$$

In the TI formalism, the free energy is obtained from its gradient, which is generally calculated in the form of the average of a force F_ξ exerted on ξ , taken over an iso- ξ surface:

$$\nabla_\xi A(\xi) = \langle -F_\xi \rangle_\xi \quad (17)$$

Several formulae that take the form of (17) have been proposed. This implementation relies partly on the classic formulation [14], and partly on a more versatile scheme originating in a work by Ruiz-Montero et al. [15], generalized by den Otter [16] and extended to multiple variables by Ciccotti et al. [17]. Consider a system subject to constraints of the form $\sigma_k(x) = 0$. Let $(v_i)_{i \in [1,n]}$ be arbitrarily chosen vector fields ($\mathbb{R}^{3N} \rightarrow \mathbb{R}^{3N}$) verifying, for all i, j , and k :

$$v_i \cdot \nabla_x \xi_j = \delta_{ij} \quad (18)$$

$$v_i \cdot \nabla_x \sigma_k = 0 \quad (19)$$

then the following holds [17]:

$$\frac{\partial A}{\partial \xi_i} = \langle v_i \cdot \nabla_x V - k_B T \nabla_x \cdot v_i \rangle_{\xi} \quad (20)$$

where V is the potential energy function. v_i can be interpreted as the direction along which the force acting on variable ξ_i is measured, whereas the second term in the average corresponds to the geometric entropy contribution that appears as a Jacobian correction in the classic formalism [14]. Condition (18) states that the direction along which the total force on ξ_i is measured is orthogonal to the gradient of ξ_j , which means that the force measured on ξ_i does not act on ξ_j .

Equation (19) implies that constraint forces are orthogonal to the directions along which the free energy gradient is measured, so that the measurement is effectively performed on unconstrained degrees of freedom.

In the framework of ABF, \mathbf{F}_{ξ} is accumulated in bins of finite size $\delta\xi$, thereby providing an estimate of the free energy gradient according to equation (17). The biasing force applied along the collective variables to overcome free energy barriers is calculated as:

$$\mathbf{F}^{\text{ABF}} = \alpha(N_{\xi}) \times \nabla_x \tilde{A}(\xi) \quad (21)$$

where $\nabla_x \tilde{A}$ denotes the current estimate of the free energy gradient at the current point ξ in the collective variable subspace, and $\alpha(N_{\xi})$ is a scaling factor that is ramped from 0 to 1 as the local number of samples N_{ξ} increases to prevent nonequilibrium effects in the early phase of the simulation, when the gradient estimate has a large variance. See the `fullSamples` parameter below for details.

As sampling of the phase space proceeds, the estimate $\nabla_x \tilde{A}$ is progressively refined. The biasing force introduced in the equations of motion guarantees that in the bin centered around ξ , the forces acting along the selected collective variables average to zero over time. Eventually, as the underlying free energy surface is canceled by the adaptive bias, evolution of the system along ξ is governed mainly by diffusion. Although this implementation of ABF can in principle be used in arbitrary dimension, a higher-dimension collective variable space is likely to be difficult to sample and visualize. Most commonly, the number of variables is one or two, sometimes three.

6.2.1 ABF requirements on collective variables

The following conditions must be met for an ABF simulation to be possible and to produce an accurate estimate of the free energy profile. Note that these requirements do not apply when using the extended-system ABF method ().

1. *Only linear combinations* of colvar components can be used in ABF calculations.

2. *Availability of total forces* is necessary. The following colvar components can be used in ABF calculations: distance, distance_xy, distance_z, angle, dihedral, gyration, rmsd and eigenvector. Atom groups may not be replaced by dummy atoms, unless they are excluded from the force measurement by specifying oneSiteTotalForce, if available.
3. *Mutual orthogonality of colvars*. In a multidimensional ABF calculation, equation (18) must be satisfied for any two colvars ξ_i and ξ_j . Various cases fulfill this orthogonality condition:
 - ξ_i and ξ_j are based on non-overlapping sets of atoms.
 - atoms involved in the force measurement on ξ_i do not participate in the definition of ξ_j . This can be obtained using the option oneSiteTotalForce of the distance, angle, and dihedral components (example: Ramachandran angles ϕ , ψ).
 - ξ_i and ξ_j are orthogonal by construction. Useful cases are the sum and difference of two components, or distance_z and distance_xy using the same axis.
4. *Mutual orthogonality of components*: when several components are combined into a colvar, it is assumed that their vectors v_i (equation (20)) are mutually orthogonal. The cases described for colvars in the previous paragraph apply.
5. *Orthogonality of colvars and constraints*: equation 19 can be satisfied in two simple ways, if either no constrained atoms are involved in the force measurement (see point 3 above) or pairs of atoms joined by a constrained bond are part of an *atom group* which only intervenes through its center (center of mass or geometric center) in the force measurement. In the latter case, the contributions of the two atoms to the left-hand side of equation 19 cancel out. For example, all atoms of a rigid TIP3P water molecule can safely be included in an atom group used in a distance component.

6.2.2 Parameters for ABF

ABF depends on parameters from collective variables to define the grid on which free energy gradients are computed. In the direction of each colvar, the grid ranges from lowerBoundary to upperBoundary, and the bin width (grid spacing) is set by the width parameter. The following specific parameters can be set in the ABF configuration block:

- **Keyword name:** see definition of name (biasing and analysis methods)
- **Keyword colvars:** see definition of colvars (biasing and analysis methods)
- **Keyword outputEnergy:** see definition of outputEnergy (biasing and analysis methods)
- **Keyword outputFreq:** see definition of outputFreq (biasing and analysis methods)
- **Keyword stepZeroData:** see definition of stepZeroData (biasing and analysis methods)
- **Keyword fullSamples** ⟨Number of samples in a bin prior to application of the ABF⟩
Context: abf
Acceptable values: positive integer
Default value: 200

Description: To avoid nonequilibrium effects due to large fluctuations of the force exerted along the colvars, it is recommended to apply a biasing force only after the estimate has started converging. If `fullSamples` is non-zero, the applied biasing force is scaled by a factor $\alpha(N_\xi)$ between 0 and 1. If the number of samples N_ξ in the current bin is higher than `fullSamples`, the factor is one. If it is less than half of `fullSamples`, the factor is zero and no bias is applied. Between those two thresholds, the factor follows a linear ramp from 0 to 1: $\alpha(N_\xi) = (2N_\xi/\text{fullSamples}) - 1$.

- Keyword** `maxForce` \langle Maximum magnitude of the ABF force \rangle
Context: `abf`
Acceptable values: positive decimals (one per colvar)
Default value: disabled
Description: This option enforces a cap on the magnitude of the biasing force effectively applied by this ABF bias on each colvar. This can be useful in the presence of singularities in the PMF such as hard walls, where the discretization of the average force becomes very inaccurate, causing the colvar's diffusion to get “stuck” at the singularity. To enable this cap, provide one non-negative value for each colvar. The unit of force is *energy unit/(colvar unit)*.
- Keyword** `hideJacobian` \langle Remove geometric entropy term from calculated free energy gradient? \rangle
Context: `abf`
Acceptable values: boolean
Default value: no
Description: In a few special cases, most notably distance-based variables, an alternate definition of the potential of mean force is traditionally used, which excludes the Jacobian term describing the effect of geometric entropy on the distribution of the variable. This results, for example, in particle-particle potentials of mean force being flat at large separations. The Jacobian term is exactly represented in equation (20) by the second term of the average, $-k_B T \nabla_x \cdot v_i$. Enabling the `hideJacobian` option causes the output data to follow the traditional potential of mean force convention, by omitting this contribution from the measured free energy gradients. To ensure uniform sampling despite the incomplete description of the free energy, an additional biasing force that counteracts the Jacobian force is applied internally by the colvar.
- Keyword** `historyFreq` \langle Frequency (in timesteps) at which ABF history files are accumulated \rangle
Context: `abf`
Acceptable values: positive integer
Default value: 0
Description: If this number is non-zero, the free energy gradient estimate and sampling histogram (and the PMF in one-dimensional calculations) are written to files on disk at the given time interval. History file names use the same prefix as output files, with “.hist” appended (*output.hist.pmf*). `historyFreq` must be a multiple of `outputFreq`.
- Keyword** `inputPrefix` \langle Filename prefix for reading ABF data \rangle
Context: `abf`
Acceptable values: list of strings
Description: If this parameter is set, for each item in the list, ABF tries to read a gradient and a sampling files named `<inputPrefix>.grad` and `<inputPrefix>.count`. This is done at startup and sets the initial state of the ABF algorithm. The data from all provided files is combined appropriately. Also, the grid definition (min and max values, width) need not be the same that for the current run. This command is useful to piece together data from simulations in different regions of collective variable space, or change the colvar boundary values and widths. Note that it is not recommended

to use it to switch to a smaller width, as that will leave some bins empty in the finer data grid. This option is NOT compatible with reading the data from a restart file (input keyword of the `fix ID group-ID colvars` command).

- **Keyword** `applyBias` ⟨Apply the ABF bias?⟩

Context: `abf`

Acceptable values: `boolean`

Default value: `yes`

Description: If this is set to no, the calculation proceeds normally but the adaptive biasing force is not applied. Data is still collected to compute the free energy gradient. This is mostly intended for testing purposes, and should not be used in routine simulations.

- **Keyword** `updateBias` ⟨Update the ABF bias?⟩

Context: `abf`

Acceptable values: `boolean`

Default value: `yes`

Description: If this is set to no, the initial biasing force (e.g. read from a restart file or through `inputPrefix`) is not updated during the simulation. As a result, a constant bias is applied. This can be used to apply a custom, tabulated biasing potential to any combination of colvars. To that effect, one should prepare a gradient file containing the gradient of the potential to be applied (negative of the bias force), and a count file containing only values greater than `fullSamples`. These files must match the grid parameters of the colvars.

6.2.3 Output files

The ABF bias produces the following files, all in multicolumn text format (3.7.4):

- `output.grad`: current estimate of the free energy gradient (grid), in multicolumn;
- `output.count`: histogram of samples collected, on the same grid;
- `output.pmf`: integrated free energy profile or PMF (for dimension 1, and dimension 2 or 3).

Also in the case of one-dimensional calculations, the ABF bias can report its current energy via `outputEnergy`; in higher dimensions, such computation is not implemented and the energy reported is zero.

If several ABF biases are defined concurrently, their name is inserted to produce unique filenames for output, as in `output.abf1.grad`. This should not be done routinely and could lead to meaningless results: only do it if you know what you are doing!

If the colvar space has been partitioned into sections (*windows*) in which independent ABF simulations have been run, the resulting data can be merged using the `inputPrefix` option described above (a run of 0 steps is enough).

6.2.4 Multidimensional free energy surfaces

The ABF method only produces an estimate of the free energy *gradient*. The free energy surface itself can be computed depending on the value of `integrate` and related options.

- **Keyword** `integrate` \langle Integrate free energy surface from ABF gradients (dim < 3) \rangle

Context: `abf`

Acceptable values: `boolean`

Default value: `yes`

Description: This option, active by default when the dimension of the colvar space is 3 or less, enables the calculation of an integrated free energy surface every time ABF output files are written. In dimension 2 or 3, integration is performed by solving a Poisson equation: [18]

$$\nabla^2 A_t = \nabla \cdot G_t \quad (22)$$

where G_t is the estimated gradient at time t , and A_t is corresponding free energy surface. The free energy surface is written under the file name `<outputName>.pmf`, in a plain text format (see 3.7.4) that can be read by most data plotting and analysis programs (e.g. Gnuplot). Periodic boundary conditions are applied to periodic coordinates, and Neumann boundary conditions otherwise (imposed free energy gradient at the boundary of the domain). The grid used for free energy discretization is extended by one point along non-periodic coordinates, but not along periodic coordinates. See ref. [18] for details.

- **Keyword** `integrateTol` \langle Tolerance for free energy integration \rangle

Context: `abf`

Acceptable values: `positive decimal`

Default value: `1e-6`

Description: The conjugate gradients algorithm used to integrate the free energy surface is stopped when the RMS error reaches `integrateTol`.

- **Keyword** `integrateMaxIterations` \langle Maximum iterations for free energy integration \rangle

Context: `abf`

Acceptable values: `integer`

Default value: `10000`

Description: The conjugate gradients algorithm used to integrate the free energy surface is stopped when the number of iterations reaches `integrateMaxIterations`, unless the RMS error has reached `integrateTol` before.

In dimension 4 or greater, integrating the discretized gradient becomes non-trivial. The standalone utility `abf_integrate` is provided to perform that task. Because 4D ABF calculations are uncommon, **this tool is practically deprecated** by the Poisson integration described above.

`abf_integrate` reads the gradient data and uses it to perform a Monte-Carlo (M-C) simulation in discretized collective variable space (specifically, on the same grid used by ABF to discretize the free energy gradient). By default, a history-dependent bias (similar in spirit to metadynamics) is used: at each M-C step, the bias at the current position is incremented by a preset amount (the *hill height*). Upon convergence, this bias counteracts optimally the underlying gradient; it is negated to obtain the estimate of the free energy surface.

`abf_integrate` is invoked using the command-line:

```
abf_integrate <gradient_file> [-n <nsteps>] [-t <temp>] [-m (0|1)] [-h <hill_height>] [-f <factor>]
```

The gradient file name is provided first, followed by other parameters in any order. They are described below, with their default value in square brackets:

- `-n`: number of M-C steps to be performed; by default, a minimal number of steps is chosen based on the size of the grid, and the integration runs until a convergence criterion is satisfied (based on the RMSD between the target gradient and the real PMF gradient)

- -t: temperature for M-C sampling; expects kcal/mol as energy unit (unrelated to the simulation temperature) [500 K]
- -s: scaling factor for the gradients; when using a histogram of total forces obtained from `outputTotalForce` or the `.force` file written by `writeTISamples`, a scaling factor of -1 should be used [1.0]
- -m: use metadynamics-like biased sampling? (0 = false) [1]
- -h: increment for the history-dependent bias (“hill height”) [0.01 kcal/mol]
- -f: if non-zero, this factor is used to scale the increment stepwise in the second half of the M-C sampling to refine the free energy estimate [0.5]

Using the default values of all parameters should give reasonable results in most cases.

`abf_integrate` produces the following output files:

- `<gradient_file>.pmf`: computed free energy surface
- `<gradient_file>.histo`: histogram of M-C sampling (not usable in a straightforward way if the history-dependent bias has been applied)
- `<gradient_file>.est`: estimated gradient of the calculated free energy surface (from finite differences)
- `<gradient_file>.dev`: deviation between the user-provided numerical gradient and the actual gradient of the calculated free energy surface. The RMS norm of this vector field is used as a convergence criterion and output periodically during integration.

Note: Typically, the “deviation” vector field does not vanish as the integration converges. This happens because the numerical estimate of the gradient does not exactly derive from a potential, due to numerical approximations used to obtain it (finite sampling and discretization on a grid). See Ref.[18] for details.

6.3 Extended-system Adaptive Biasing Force (eABF)

Extended-system ABF (eABF) is a variant of ABF (6.2) where the bias is not applied directly to the collective variable, but to an extended coordinate (“fictitious variable”) λ that evolves dynamically according to Newtonian or Langevin dynamics. Such an extended coordinate is enabled for a given colvar using the `extendedLagrangian` and associated keywords (4.18). The theory of eABF and the present implementation are documented in detail in reference [19].

Defining an ABF bias on a colvar wherein the `extendedLagrangian` option is active will perform eABF automatically; there is no dedicated option.

The extended variable λ is coupled to the colvar $z = \xi(q)$ by the harmonic potential $(k/2)(z - \lambda)^2$. Under eABF dynamics, the adaptive bias on λ is the running estimate of the average spring force:

$$F^{\text{bias}}(\lambda^*) = \langle k(\lambda - z) \rangle_{\lambda^*} \quad (23)$$

where the angle brackets indicate a canonical average conditioned by $\lambda = \lambda^*$. At long simulation times, eABF produces a flat histogram of the extended variable λ , and a flattened histogram of ξ , whose exact shape depends on the strength of the coupling as defined by `extendedFluctuation` in the colvar. Coupling should be somewhat loose for faster exploration and convergence, but strong enough that the bias does help overcome barriers along the colvar ξ . [19] Distribution of the colvar may be assessed by plotting its histogram, which is written to the `output.zcount` file in every eABF simulation. Note that a histogram

bias (6.10) applied to an extended-Lagrangian colvar will access the extended degree of freedom λ , not the original colvar ξ ; however, the joint histogram may be explicitly requested by listing the name of the colvar twice in a row within the colvars parameter of the histogram block.

The eABF PMF is that of the coordinate λ , it is not exactly the free energy profile of ξ . That quantity can be calculated based on the CZAR estimator.

6.3.1 CZAR estimator of the free energy

The *corrected z-averaged restraint* (CZAR) estimator is described in detail in reference [19]. It is computed automatically in eABF simulations, regardless of the number of colvars involved. Note that ABF may also be applied on a combination of extended and non-extended colvars; in that case, CZAR still provides an unbiased estimate of the free energy gradient.

CZAR estimates the free energy gradient as:

$$A'(z) = -\frac{1}{\beta} \frac{d \ln \tilde{\rho}(z)}{dz} + k(\langle \lambda \rangle_z - z). \quad (24)$$

where $z = \xi(q)$ is the colvar, λ is the extended variable harmonically coupled to z with a force constant k , and $\tilde{\rho}(z)$ is the observed distribution (histogram) of z , affected by the eABF bias.

Parameters for the CZAR estimator are:

- Keyword** `CZARestimator` \langle Calculate CZAR estimator of the free energy? \rangle
Context: abf
Acceptable values: boolean
Default value: yes
Description: This option is only available when ABF is performed on extended-Lagrangian colvars. When enabled, it triggers calculation of the free energy following the CZAR estimator.
- Keyword** `writeCZARwindowFile` \langle Write internal data from CZAR to a separate file? \rangle
Context: abf
Acceptable values: boolean
Default value: no
Description: When this option is enabled, eABF simulations will write a file containing the z-averaged restraint force under the name `output.zgrad`. The same information is always included in the colvars state file, which is sufficient for restarting an eABF simulation. These separate file is only useful when joining adjacent windows from a stratified eABF simulation, either to continue the simulation in a broader window or to compute a CZAR estimate of the PMF over the full range of the coordinate(s). **Important warning.** Unbiased free-energy estimators from eABF dynamics rely on some form of sampling histogram. When running stratified (windowed) calculations this histogram becomes discontinuous, and as a result the free energy gradient estimated by CZAR is inaccurate at the window boundary, resulting in visible "blips" in the PMF. As a workaround, we recommend manually replacing the two free energy gradient values at the boundary, either with the ABF values from `.grad` files (accurate in the limit of tight coupling), or with values interpolated for the neighboring values of the CZAR gradient.

Similar to ABF, the CZAR estimator produces two output files in multicolumn text format (3.7.4):

- `output.czar.grad`: current estimate of the free energy gradient (grid), in multicolumn;

- `output.czar.pmf`: only for one-dimensional calculations, integrated free energy profile or PMF.

The sampling histogram associated with the CZAR estimator is the z -histogram, which is written in the file `output.zcount`.

6.4 Metadynamics

The metadynamics method uses a history-dependent potential [20] that generalizes to any type of colvars the conformational flooding [21] and local elevation [22] methods, originally formulated to use as colvars the principal components of a covariance matrix or a set of dihedral angles, respectively. The metadynamics potential on the colvars $\xi = (\xi_1, \xi_2, \dots, \xi_{N_{cv}})$ is defined as:

$$V_{\text{meta}}(\xi(t)) = \sum_{t'=\delta t, 2\delta t, \dots}^{t'<t} W \prod_{i=1}^{N_{cv}} \exp\left(-\frac{(\xi_i(t) - \xi_i(t'))^2}{2\sigma_{\xi_i}^2}\right), \quad (25)$$

where V_{meta} is the history-dependent potential acting on the *current* values of the colvars ξ , and depends only parametrically on the *previous* values of the colvars. V_{meta} is constructed as a sum of N_{cv} -dimensional repulsive Gaussian “hills”, whose height is a chosen energy constant W , and whose centers are the previously explored configurations $(\xi(\delta t), \xi(2\delta t), \dots)$.

During the simulation, the system evolves towards the nearest minimum of the “effective” potential of mean force $\tilde{A}(\xi)$, which is the sum of the “real” underlying potential of mean force $A(\xi)$ and the the metadynamics potential, $V_{\text{meta}}(\xi)$. Therefore, at any given time the probability of observing the configuration ξ^* is proportional to $\exp(-\tilde{A}(\xi^*)/\kappa_B T)$: this is also the probability that a new Gaussian “hill” is added at that configuration. If the simulation is run for a sufficiently long time, each local minimum is canceled out by the sum of the Gaussian “hills”. At that stage the “effective” potential of mean force $\tilde{A}(\xi)$ is constant, and $-V_{\text{meta}}(\xi)$ is an estimator of the “real” potential of mean force $A(\xi)$, save for an additive constant:

$$A(\xi) \simeq -V_{\text{meta}}(\xi) + K \quad (26)$$

Such estimate of the free energy can be provided by enabling `writeFreeEnergyFile`. Assuming that the set of collective variables includes all relevant degrees of freedom, the predicted error of the estimate is a simple function of the correlation times of the colvars τ_{ξ_i} , and of the user-defined parameters W , σ_{ξ_i} and δt [23]. In typical applications, a good rule of thumb can be to choose the ratio $W/\delta t$ much smaller than $\kappa_B T/\tau_{\xi}$, where τ_{ξ} is the longest among ξ ’s correlation times: σ_{ξ_i} then dictates the resolution of the calculated PMF.

If the metadynamics parameters are chosen correctly, after an equilibration time, t_e , the estimator provided by eq. 26 oscillates on time around the “real” free energy, thereby a better estimate of the latter can be obtained as the time average of the bias potential after t_e [24, 25]:

$$A(\xi) = -\frac{1}{t_{\text{tot}} - t_e} \int_{t_e}^{t_{\text{tot}}} V_{\text{meta}}(\xi, t) dt \quad (27)$$

where t_e is the time after which the bias potential grows (approximately) evenly during the simulation and t_{tot} is the total simulation time. The free energy calculated according to eq. 27 can thus be obtained averaging on time multiple time-dependent free energy estimates, that can be printed out through the keyword `keepFreeEnergyFiles`. An alternative is to obtain the free energy profiles by summing the hills added during the simulation; the hills trajectory can be printed out by enabling the option `writeHillsTrajectory`.

6.4.1 Treatment of the PMF boundaries

In typical scenarios the Gaussian hills of a metadynamics potential are interpolated and summed together onto a grid, which is much more efficient than computing each hill independently at every step (the keyword `useGrids` is on by default). This numerical approximation typically yields negligible errors in the resulting PMF [1]. However, due to the finite thickness of the Gaussian function, the metadynamics potential would suddenly vanish each time a variable exceeds its grid boundaries.

To avoid such discontinuity the Colvars metadynamics code will keep an explicit copy of each hill that straddles a grid's boundary, and will use it to compute metadynamics forces outside the grid. This measure is taken to protect the accuracy and stability of a metadynamics simulation, except in cases of “natural” boundaries (for example, the $[0 : 180]$ interval of an angle colvar) or when the flags `hardLowerBoundary` and `hardUpperBoundary` are explicitly set by the user. Unfortunately, processing explicit hills alongside the potential and force grids could easily become inefficient, slowing down the simulation and increasing the state file's size.

In general, it is a good idea to *define a repulsive potential to avoid hills from coming too close to the grid's boundaries*, for example as a `harmonicWalls` restraint (see 6.7).

Example: Using harmonic walls to protect the grid's boundaries.

```
colvar {  
  name r  
  distance { ... }  
  upperBoundary 15.0  
  width 0.2  
}  
  
metadynamics {  
  name meta_r  
  colvars r  
  hillWeight 0.001  
  hillWidth 2.0  
}  
  
harmonicWalls {  
  name wall_r  
  colvars r  
  upperWalls 13.0  
  upperWallConstant 2.0  
}
```

In the colvar `r`, the distance function used has a `lowerBoundary` automatically set to 0 by default, thus the keyword `lowerBoundary` itself is not mandatory and `hardLowerBoundary` is set to yes internally. However, `upperBoundary` does not have such a “natural” choice of value. The metadynamics potential `meta_r` will individually process any hill whose center is too close to the `upperBoundary`, more precisely within fewer grid points than 6 times the Gaussian σ parameter plus one. It goes without saying that if the colvar `r` represents a distance between two freely-moving molecules, it will cross this “threshold” rather frequently.

In this example, where the value of `hillWidth` (2σ) amounts to 2 grid points, the threshold is $6+1 = 7$ grid points away from `upperBoundary`. In explicit units, the width of r is $w_r = 0.2 \text{ \AA}$, and the threshold is $15.0 - 7 \times 0.2 = 13.6 \text{ \AA}$.

The `wall_r` restraint included in the example prevents this: the position of its `upperWall` is 13 \AA , i.e. 3 grid points below the buffer's threshold (13.6 \AA). For the chosen value of `upperWallConstant`, the energy of the `wall_r` bias at $r = r_{\text{upper}} = 13.6 \text{ \AA}$ is:

$$E^* = \frac{1}{2}k \left(\frac{r - r_{\text{upper}}}{w_r} \right)^2 = \frac{1}{2}2.0(-3)^2 = 9 \text{ kcal/mol}$$

which results in a relative probability $\exp(-E^*/\kappa_B T) \simeq 3 \times 10^{-7}$ that r crosses the threshold. The probability that r exceeds `upperBoundary`, which is further away, has also become vanishingly small. At that point, you may want to set `hardUpperBoundary` to yes for r , and let `meta_r` know that no special treatment near the grid's boundaries will be needed.

What is the impact of the wall restraint onto the PMF? Not a very complicated one: the PMF reconstructed by metadynamics will simply show a sharp increase in free-energy where the wall potential kicks in ($r > 13 \text{ \AA}$). You may then choose between using the PMF only up until that point and discard the rest, or subtracting the energy of the `harmonicWalls` restraint from the PMF itself. Keep in mind, however, that the statistical convergence of metadynamics may be less accurate where the wall potential is strong.

In summary, although it would be simpler to set the wall's position `upperWall` and the grid's boundary `upperBoundary` to the same number, the finite width of the Gaussian hills calls for setting the former strictly within the latter.

6.4.2 Basic configuration keywords

To enable a metadynamics calculation, a `metadynamics {...}` block must be defined in the Colvars configuration file. Its mandatory keywords are `colvars`, the variables involved, `hillWeight`, the weight parameter W , and the widths 2σ of the Gaussian hills in each dimension given by the single dimensionless parameter `hillWidth`, or more explicitly by the `gaussianSigmas`.

- **Keyword** `name`: see definition of `name` (biasing and analysis methods)
- **Keyword** `colvars`: see definition of `colvars` (biasing and analysis methods)
- **Keyword** `outputEnergy`: see definition of `outputEnergy` (biasing and analysis methods)
- **Keyword** `outputFreq`: see definition of `outputFreq` (biasing and analysis methods)
- **Keyword** `writeTIPMF`: see definition of `writeTIPMF` (biasing and analysis methods)
- **Keyword** `writeTISamples`: see definition of `writeTISamples` (biasing and analysis methods)
- **Keyword** `stepZeroData`: see definition of `stepZeroData` (biasing and analysis methods)
- **Keyword** `hillWeight` $\langle \text{Height of each hill (energy unit)} \rangle$
Context: metadynamics
Acceptable values: positive decimal
Description: This option sets the height W of the Gaussian hills that are added during this run.

Lower values provide more accurate sampling of the system's degrees of freedom at the price of longer simulation times to complete a PMF calculation based on metadynamics.

- **Keyword** `hillWidth` \langle Width 2σ of a Gaussian hill, measured in number of grid points \rangle
Context: metadynamics
Acceptable values: positive decimal
Description: This keyword sets the Gaussian width $2\sigma_{\xi_i}$ for all colvars, expressed in *number of grid points*, with the grid spacing along each colvar ξ determined by the respective value of `width`. Values between 1 and 3 are recommended for this option: smaller numbers will fail to adequately interpolate each Gaussian function [1], while larger values may be unable to account for steep free-energy gradients. The values of each half-width σ_{ξ_i} in the physical units of ξ_i are also printed by LAMMPS at initialization time; alternatively, they may be set explicitly via `gaussianSigmas`.
- **Keyword** `gaussianSigmas` \langle Half-widths σ of the Gaussian hill (one for each colvar) \rangle
Context: metadynamics
Acceptable values: space-separated list of decimals
Description: This option sets the parameters σ_{ξ_i} of the Gaussian hills along each colvar ξ_i , expressed in *the same unit of ξ_i* . No restrictions are placed on each value, but a warning will be printed if `useGrids` is on and the Gaussian width $2\sigma_{\xi_i}$ is smaller than the corresponding grid spacing, `width(ξ_i)`. If not given, default values will be computed from the dimensionless number `hillWidth`.
- **Keyword** `newHillFrequency` \langle Frequency of hill creation \rangle
Context: metadynamics
Acceptable values: positive integer
Default value: 1000
Description: This option sets the number of steps after which a new Gaussian hill is added to the metadynamics potential. The product of this number and the integration time-step defines the parameter δt in eq. 25. Higher values provide more accurate statistical sampling, at the price of longer simulation times to complete a PMF calculation.

6.4.3 Output files

When interpolating grids are enabled (default behavior), the PMF is written by default every `colvarsRestartFrequency` steps to the file `output.pmf` in multicolumn text format (3.7.4). The following two options allow to disable or control this behavior and to track statistical convergence:

- **Keyword** `writeFreeEnergyFile` \langle Periodically write the PMF for visualization \rangle
Context: metadynamics
Acceptable values: boolean
Default value: on
Description: When `useGrids` and this option are on, the PMF is written every `outputFreq` steps.
- **Keyword** `keepFreeEnergyFiles` \langle Keep all the PMF files \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: When `writeFreeEnergyFile` and this option are on, the step number is included in

the file name, thus generating a series of PMF files. Activating this option can be useful to follow more closely the convergence of the simulation, by comparing PMFs separated by short times.

- **Keyword** `writeHillsTrajectory` \langle Write a log of new hills \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: If this option is on, a file containing the Gaussian hills written by the metadynamics bias, with the name:
`"output.colvars.<name>.hills.traj"`,
which can be useful to post-process the time series of the Gaussian hills. Each line is written every `newHillFrequency`, regardless of the value of `outputFreq`. When `multipleReplicas` is on, its name is changed to:
`"output.colvars.<name>.<replicaID>.hills.traj"`.
The columns of this file are the centers of the hills, $\xi_i(t')$, followed by the half-widths, σ_{ξ_i} , and the weight, W . **Note:** prior to version 2020-02-24, the full-width 2σ of the Gaussian was reported in lieu of σ .

6.4.4 Performance optimization

The following options control the computational cost of metadynamics calculations, but do not affect results. Default values are chosen to minimize such cost with no loss of accuracy.

- **Keyword** `useGrids` \langle Interpolate the hills with grids \rangle
Context: metadynamics
Acceptable values: boolean
Default value: on
Description: This option discretizes all hills for improved performance, accumulating their energy and their gradients on two separate grids of equal spacing. Grids are defined by the values of `lowerBoundary`, `upperBoundary` and `width` for each colvar. Currently, this option is implemented for all types of variables except the non-scalar types (`distanceDir` or `orientation`). If `expandBoundaries` is defined in one of the colvars, grids are automatically expanded along the direction of that colvar.
- **Keyword** `rebinGrids` \langle Recompute the grids when reading a state file \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: When restarting from a state file, the grid's parameters (boundaries and widths) saved in the state file override those in the configuration file. Enabling this option forces the grids to match those in the current configuration file.
- **Keyword** `keepHills` \langle Write each individual hill to the state file \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: When `useGrids` and this option are on, all hills are saved to the state file in their

analytic form, alongside their grids. This makes it possible to later use exact analytic Gaussians for rebinGrids. To only keep track of the history of the added hills, writeHillsTrajectory is preferable.

6.4.5 Ensemble-Biased Metadynamics

The ensemble-biased metadynamics (EBMetaD) approach [26] is designed to reproduce a target probability distribution along selected collective variables. Standard metadynamics can be seen as a special case of EBMetaD with a flat distribution as target. This is achieved by weighing the Gaussian functions used in the metadynamics approach by the inverse of the target probability distribution:

$$V_{\text{EBMetaD}}(\xi(t)) = \sum_{t'=\delta t, 2\delta t, \dots}^{t' \leq t} \frac{W}{\exp(S_p) \rho_{\text{exp}}(\xi(t'))} \prod_{i=1}^{N_{\text{cv}}} \exp\left(-\frac{(\xi_i(t) - \xi_i(t'))^2}{2\sigma_{\xi_i}^2}\right), \quad (28)$$

where $\rho_{\text{exp}}(\xi)$ is the target probability distribution and $S_p = -\int \rho_{\text{exp}}(\xi) \log \rho_{\text{exp}}(\xi) d\xi$ its corresponding differential entropy. The method is designed so that during the simulation the resulting distribution of the collective variable ξ converges to $\rho_{\text{exp}}(\xi)$. A practical application of EBMetaD is to reproduce an “experimental” probability distribution, for example the distance distribution between spectroscopic labels inferred from Förster resonance energy transfer (FRET) or double electron-electron resonance (DEER) experiments [26].

The PMF along ξ can be estimated from the bias potential and the target distribution [26]:

$$A(\xi) \simeq -V_{\text{EBMetaD}}(\xi) - \kappa_B T \log \rho_{\text{exp}}(\xi) \quad (29)$$

and obtained by enabling writeFreeEnergyFile. Similarly to eq. 27, a more accurate estimate of the free energy can be obtained by averaging (after an equilibration time) multiple time-dependent free energy estimates (see keepFreeEnergyFiles).

The following additional options define the configuration for the ensemble-biased metadynamics approach:

- Keyword ebMeta** (Perform ensemble-biased metadynamics)
 - Context:** metadynamics
 - Acceptable values:** boolean
 - Default value:** off
 - Description:** If enabled, this flag activates the ensemble-biased metadynamics as described by Marinelli et al.[26]. The target distribution file, targetdistfile, is then required. The keywords lowerBoundary, upperBoundary and width for the respective variables are also needed to set the binning (grid) of the target distribution file.
- Keyword targetDistFile** (Target probability distribution file for ensemble-biased metadynamics)
 - Context:** metadynamics
 - Acceptable values:** multicolumn text file
 - Description:** This file provides the target probability distribution, $\rho_{\text{exp}}(\xi)$, reported in eq. 28. The latter distribution must be a tabulated function provided in a multicolumn text format (see 3.7.4). The provided distribution is then normalized.

- **Keyword** `ebMetaEquilSteps` ⟨Number of equilibration steps for ensemble-biased metadynamics⟩

Context: metadynamics

Acceptable values: positive integer

Description: The EBMetaD approach may introduce large hills in regions with small values of the target probability distribution (eq. 28). This happens, for example, if the probability distribution sampled by a conventional molecular dynamics simulation is significantly different from the target distribution. This may lead to instabilities at the beginning of the simulation related to large biasing forces. In this case, it is useful to introduce an equilibration stage in which the bias potential gradually switches from standard metadynamics (eq. 25) to EBMetaD (eq. 28) as $\lambda V_{\text{meta}}(\xi) + (1 - \lambda)V_{\text{EBmetaD}}(\xi)$, where $\lambda = (\text{ebMetaEquilSteps} - \text{step})/\text{ebMetaEquilSteps}$ and `step` is the current simulation step number.

- **Keyword** `targetDistMinVal` ⟨Minimum value of the target distribution in reference to its maximum value⟩

Context: metadynamics

Acceptable values: positive decimal

Description: It is useful to set a minimum value of the target probability distribution to avoid values of the latter that are nearly zero, leading to very large hills. This parameter sets the minimum value of the target probability distribution that is expressed as a fraction of its maximum value: `minimum value = maximum value X targetDistMinVal`. This implies that $0 < \text{targetDistMinVal} < 1$ and its default value is set to 1/1000000. To avoid divisions by zero (see eq. 28), if `targetDistMinVal` is set as zero, values of ρ_{exp} equal to zero are replaced by the smallest positive value read in the same file.

As with standard metadynamics, multidimensional probability distributions can be targeted using a single metadynamics block using multiple colvars and a multidimensional target distribution file (see 3.7.4). Instead, multiple probability distributions on different variables can be targeted separately in the same simulation by introducing multiple metadynamics blocks with the `ebMeta` option.

Example: EBMetaD configuration for a single variable.

```
colvar {
  name r
  distance {
    group1 { atomNumbers 991 992 }
    group2 { atomNumbers 1762 1763 }
  }
  upperBoundary 100.0
  width 0.1
}

metadynamics {
  name ebmeta
  colvars r
  hillWeight 0.01
  hillWidth 3.0
  ebMeta on
  targetDistFile targetdist1.dat
}
```



```
ebMetaEquilSteps 500000
}
```

where `targetdist1.dat` is a text file in “multicolumn” format (3.7.4) with the same width as the variable `r` (0.1 in this case):

```
# 1
# 0.0 0.1 1000 0

0.05 0.0012
0.15 0.0014
... ...
99.95 0.0010
```

Tip: Besides setting a meaningful value for `targetDistMinVal`, the exploration of unphysically low values of the target distribution (which would lead to very large hills and possibly numerical instabilities) can be also prevented by restricting sampling to a given interval, using e.g. `harmonicWalls` restraint (6.7).

6.4.6 Well-tempered metadynamics

The following options define the configuration for the “well-tempered” metadynamics approach [27]:

- **Keyword** `wellTempered` \langle Perform well-tempered metadynamics \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: If enabled, this flag causes well-tempered metadynamics as described by Barducci et al.[27] to be performed, rather than standard metadynamics. The parameter `biasTemperature` is then required. This feature was contributed by Li Li (Luthey-Schulten group, Department of Chemistry, UIUC).
- **Keyword** `biasTemperature` \langle Temperature bias for well-tempered metadynamics \rangle
Context: metadynamics
Acceptable values: positive decimal
Description: When running metadynamics in the long time limit, collective variable space is sampled to a modified temperature $T + \Delta T$. In conventional metadynamics, the temperature “boost” ΔT would constantly increases with time. Instead, in well-tempered metadynamics ΔT must be defined by the user via `biasTemperature`. The written PMF includes the scaling factor $(T + \Delta T)/\Delta T$ [27]. A careful choice of ΔT determines the sampling and convergence rate, and is hence crucial to the success of a well-tempered metadynamics simulation.

6.4.7 Multiple-walker metadynamics

Metadynamics calculations can be performed concurrently by multiple replicas that share a common history. This variant of the method is called multiple-walker metadynamics [28]: the Gaussian hills of all replicas are periodically combined into a single biasing potential, intended to converge to a single PMF.

In the implementation here described [1], replicas communicate through files. This arrangement allows launching the replicas either (1) as a bundle (i.e. a single job in a cluster's queueing system) or (2) as fully independent runs (i.e. as separate jobs for the queueing system). One advantage of the use case (1) is that an identical Colvars configuration can be used for all replicas (otherwise, replicaID needs to be manually set to a different string for each replica). However, the use case (2) is less demanding in terms of high-performance computing resources: a typical scenario would be a computer cluster (including virtual servers from a cloud provider) where not all nodes are connected to each other at high speed, and thus each replica runs on a small group of nodes or a single node.

Whichever way the replicas are started (coupled or not), a shared filesystem is needed so that each replica can read the files created by the others: paths to these files are stored in the shared file replicasRegistry. This file, and those listed in it, are read every replicaUpdateFrequency steps. Each time the Colvars state file is written (for example, colvarsRestartFrequency steps), the file named:

output.colvars.name.replicaID.state

is written as well; this file contains only the state of the metadynamics bias, which the other replicas will read in turn. In between the times when this file is modified/replaced, new hills are also temporarily written to the file named:

output.colvars.name.replicaID.hills

Both files are only used for communication, and may be deleted after the replica begins writing files with a new *output*.

Example: Multiple-walker metadynamics with file-based communication.

```
metadynamics {
  name mymtd
  colvars x
  hillWeight 0.001
  newHillFrequency 1000
  hillWidth 3.0

  multipleReplicas on
  replicasRegistry /shared-folder/mymtd-replicas.txt
  replicaUpdateFrequency 50000 # Best if larger than newHillFrequency
}
```

The following are the multiple-walkers related options:

- **Keyword** `multipleReplicas` \langle Enable multiple-walker metadynamics \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: This option turns on multiple-walker communication between replicas.
- **Keyword** `replicasRegistry` \langle Multiple replicas database file \rangle
Context: metadynamics
Acceptable values: UNIX filename
Description: If `multipleReplicas` is on, this option sets the path to the replicas' shared database file. It is best to use an absolute path (especially when running individual replicas in separate folders).

- **Keyword** `replicaUpdateFrequency` 〈How often hills are shared between replicas〉
Context: metadynamics
Acceptable values: positive integer
Description: If `multipleReplicas` is on, this option sets the number of steps after which each replica tries to read the other replicas' files. On a networked file system, it is best to use a number of steps that corresponds to at least a minute of wall time.
- **Keyword** `replicaID` 〈Set the identifier for this replica〉
Context: metadynamics
Acceptable values: string
Default value: replica index (only if a shared communicator is used)
Description: If `multipleReplicas` is on, this option sets a unique identifier for this replicas. When the replicas are launched in a single command (i.e. they share a parallel communicator and are tightly synchronized) this value is optional, and defaults to the replica's numeric index (starting at zero). However, when the replicas are launched as independent runs this option is required.
- **Keyword** `writePartialFreeEnergyFile` 〈Periodically write the contribution to the PMF from this replica〉
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: If `multipleReplicas` is on, enabling this option produces an additional file `output.partial.pmf`, which can be useful to monitor the contribution of each replica to the total PMF (which is written to the file `output.pmf`). **Note:** the name of this file is chosen for consistency and convenience, *but its content is not a PMF* and it is not expected to converge, even if the total PMF does.

6.5 Harmonic restraints

The harmonic biasing method may be used to enforce fixed or moving restraints, including variants of Steered and Targeted MD. Within energy minimization runs, it allows for restrained minimization, e.g. to calculate relaxed potential energy surfaces. In the context of the Colvars module, harmonic potentials are meant according to their textbook definition:

$$V(\xi) = \frac{1}{2}k \left(\frac{\xi - \xi_0}{w_\xi} \right)^2 \quad (30)$$

There are two noteworthy aspects of this expression:

1. Because the *standard coefficient* of $1/2$ of the harmonic potential is included, this expression differs from harmonic bond and angle potentials historically used in common force fields, where the factor was typically omitted resulting in a non-standard definition of the force constant.
2. The variable ξ is not only centered at ξ_0 , but is also *scaled by its characteristic length scale* w_ξ (keyword `width`). The resulting dimensionless variable $z = (\xi - \xi_0)/w_\xi$ is typically easier to treat numerically: for example, when the forces typically experienced by ξ are much smaller than k/w_ξ and k is chosen equal to $\kappa_B T$ (thermal energy), the resulting probability distribution of z is approximately a Gaussian with mean equal to 0 and standard deviation equal to 1.

This property can be used for setting the force constant in umbrella-sampling ensemble runs: if the restraint centers are chosen in increments of w_ξ , the resulting distributions of ξ are most often optimally overlapped. In regions where the underlying free-energy landscape induces highly skewed distributions of ξ , additional windows may be added as needed, with spacings finer than w_ξ .

Beyond one dimension, the use of a scaled harmonic potential also allows a standard definition of a multi-dimensional restraint with a unified force constant:

$$V(\xi_1, \dots, \xi_M) = \frac{1}{2} k \sum_{i=1}^M \left(\frac{\xi_i - \xi_0}{w_\xi} \right)^2 \quad (31)$$

If one-dimensional or homogeneous multi-dimensional restraints are defined, and there are no other uses for the parameter w_ξ , width *can be left at its default value of 1*.

A harmonic restraint is defined by a harmonic `{...}` block, which may contain the following keywords:

- **Keyword** `name`: see definition of `name` (biasing and analysis methods)
- **Keyword** `colvars`: see definition of `colvars` (biasing and analysis methods)
- **Keyword** `outputEnergy`: see definition of `outputEnergy` (biasing and analysis methods)
- **Keyword** `writeTIPMF`: see definition of `writeTIPMF` (biasing and analysis methods)
- **Keyword** `writeTISamples`: see definition of `writeTISamples` (biasing and analysis methods)
- **Keyword** `stepZeroData`: see definition of `stepZeroData` (biasing and analysis methods)
- **Keyword** `forceConstant` \langle Scaled force constant (*energy unit*) \rangle

Context: harmonic

Acceptable values: positive decimal

Default value: 1.0

Description: This option defines a *scaled* force constant k for the harmonic potential (eq. 31). To ensure consistency for multidimensional restraints, it is divided internally by the square of the specific width of each variable (which is 1 by default). This makes all values effectively dimensionless and of commensurate size. For instance, if this force constant is set to the thermal energy $\kappa_B T$ (equal to RT if molar units are used), then the amplitude of the thermal fluctuations of each variable ξ will be on the order of its width, w_ξ . This can be used to estimate the optimal spacing of umbrella-sampling windows (under the assumption that the force constant is larger than the curvature of the underlying free energy). *The values of the actual force constants k/w_ξ^2 are always printed when the restraint is defined.*

- **Keyword** `centers` \langle Initial harmonic restraint centers \rangle

Context: harmonic

Acceptable values: space-separated list of colvar values

Description: The centers (equilibrium values) of the restraint, ξ_0 , are entered here. The number of values must be the number of requested colvars. Each value is a decimal number if the corresponding colvar returns a scalar, a “(x, y, z)” triplet if it returns a unit vector or a vector, and a “(q0, q1, q2, q3)” quadruplet if it returns a rotational quaternion. If a colvar has periodicities or symmetries, its closest image to the restraint center is considered when calculating the harmonic potential.

Tip: A complex set of restraints can be applied to a system, by defining several colvars, and applying one or more harmonic restraints to different groups of colvars. In some cases, dozens of colvars can be defined, but their value may not be relevant: to limit the size of the colvars trajectory file, it may be wise to disable `outputValue` for such “ancillary” variables, and leave it enabled only for “relevant” ones.

6.5.1 Moving restraints: steered molecular dynamics

The following options allow to change gradually the centers of the harmonic restraints during a simulation. When the centers are changed continuously, a steered MD in a collective variable space is carried out.

- **Keyword** `targetCenters` `<Steer the restraint centers towards these targets>`
Context: harmonic
Acceptable values: space-separated list of colvar values
Description: When defined, the current centers will be moved towards these values during the simulation. By default, the centers are moved over a total of `targetNumSteps` steps by a linear interpolation, in the spirit of Steered MD. If `targetNumStages` is set to a nonzero value, the change is performed in discrete stages, lasting `targetNumSteps` steps *each*. This second mode may be used to sample successive windows in the context of an Umbrella Sampling simulation. When continuing a simulation run, the centers specified in the configuration file `<colvarsConfig>` are overridden by those saved in the restart file `<colvarsInput>`. To perform Steered MD in an arbitrary space of colvars, it is sufficient to use this option and enable `outputAccumulatedWork` and/or `outputAppliedForce` within each of the colvars involved.
- **Keyword** `targetNumSteps` `<Number of steps for steering>`
Context: harmonic
Acceptable values: positive integer
Description: In single-stage (continuous) transformations, defines the number of MD steps required to move the restraint centers (or force constant) towards the values specified with `targetCenters` or `targetForceConstant`. After the target values have been reached, the centers (resp. force constant) are kept fixed. In multi-stage transformations, this sets the number of MD steps *per stage*.
- **Keyword** `outputCenters` `<Write the current centers to the trajectory file>`
Context: harmonic
Acceptable values: boolean
Default value: off
Description: If this option is chosen and `colvarsTrajFrequency` is not zero, the positions of the restraint centers will be written to the trajectory file during the simulation. This option allows to conveniently extract the PMF from the colvars trajectory files in a steered MD calculation.

Note on restarting moving restraint simulations: Information about the current step and stage of a simulation with moving restraints is stored in the restart file (state file). Thus, such simulations can be run in several chunks, and restarted directly using the same colvars configuration file. In case of a restart, the values of parameters such as `targetCenters`, `targetNumSteps`, etc. should not be changed manually.

6.5.2 Moving restraints: umbrella sampling

The centers of the harmonic restraints can also be changed in discrete stages: in this case a one-dimensional umbrella sampling simulation is performed. The sampling windows in simulation are calculated in sequence. The colvars trajectory file may then be used both to evaluate the correlation times between consecutive windows, and to calculate the frequency distribution of the colvar of interest in each window. Furthermore, frequency distributions on a predefined grid can be automatically obtained by using the histogram bias (see 6.10).

To activate an umbrella sampling simulation, the same keywords as in the previous section can be used, with the addition of the following:

- **Keyword** `targetNumStages` \langle Number of stages for steering \rangle
Context: harmonic
Acceptable values: non-negative integer
Default value: 0
Description: If non-zero, sets the number of stages in which the restraint centers or force constant are changed to their target values. If zero, the change is continuous. Each stage lasts `targetNumSteps` MD steps. To sample both ends of the transformation, the simulation should be run for `targetNumSteps` \times (`targetNumStages` + 1).

6.5.3 Changing force constant

The force constant of the harmonic restraint may also be changed to equilibrate [29].

- **Keyword** `targetForceConstant` \langle Change the force constant towards this value \rangle
Context: harmonic
Acceptable values: positive decimal
Description: When defined, the current `forceConstant` will be moved towards this value during the simulation. Time evolution of the force constant is dictated by the `targetForceExponent` parameter (see below). By default, the force constant is changed smoothly over a total of `targetNumSteps` steps. This is useful to introduce or remove restraints in a progressive manner. If `targetNumStages` is set to a nonzero value, the change is performed in discrete stages, lasting `targetNumSteps` steps *each*. This second mode may be used to compute the conformational free energy change associated with the restraint, within the FEP or TI formalisms. For convenience, the code provides an estimate of the free energy derivative for use in TI, with the format:
`colvars: Lambda= ***,** dA/dLambda= ***,**`
A more complete free energy calculation (particularly with regard to convergence analysis), while not handled by the Colvars module, can be performed by post-processing the colvars trajectory, if `colvarsTrajFrequency` is set to a suitably small value. It should be noted, however, that restraint free energy calculations may be handled more efficiently by an indirect route, through the determination of a PMF for the restrained coordinate.[29]
- **Keyword** `targetForceExponent` \langle Exponent in the time-dependence of the force constant \rangle
Context: harmonic
Acceptable values: decimal equal to or greater than 1.0

Default value: 1.0

Description: Sets the exponent, α , in the function used to vary the force constant as a function of time. The force is varied according to a coupling parameter λ , raised to the power α : $k_\lambda = k_0 + \lambda^\alpha(k_1 - k_0)$, where k_0 , k_λ , and k_1 are the initial, current, and final values of the force constant. The parameter λ evolves linearly from 0 to 1, either smoothly, or in `targetNumStages` equally spaced discrete stages, or according to an arbitrary schedule set with `lambdaSchedule`. When the initial value of the force constant is zero, an exponent greater than 1.0 distributes the effects of introducing the restraint more smoothly over time than a linear dependence, and ensures that there is no singularity in the derivative of the restraint free energy with respect to `lambda`. A value of 4 has been found to give good results in some tests.

- **Keyword** `targetEquilSteps` \langle Number of steps discarded from TI estimate \rangle

Context: harmonic

Acceptable values: positive integer

Description: Defines the number of steps within each stage that are considered equilibration and discarded from the restraint free energy derivative estimate reported in the output.

- **Keyword** `lambdaSchedule` \langle Schedule of lambda-points for changing force constant \rangle

Context: harmonic

Acceptable values: list of real numbers between 0 and 1

Description: If specified together with `targetForceConstant`, sets the sequence of discrete λ values that will be used for different stages.

6.6 Computing the work of a changing restraint

If the restraint centers or force constant are changed continuously (`targetNumStages` undefined) it is possible to record the net work performed by the changing restraint:

- **Keyword** `outputAccumulatedWork` \langle Write the accumulated work of the changing restraint to the Colvars trajectory file \rangle

Context: harmonic

Acceptable values: boolean

Default value: off

Description: If `targetCenters` or `targetForceConstant` are defined and this option is enabled, the accumulated work from the beginning of the simulation will be written to the trajectory file (`colvarsTrajFrequency` must be non-zero). When the simulation is continued from a state file, the previously accumulated work is included in the integral. This option allows to conveniently extract the estimated PMF of a steered MD calculation (when `targetCenters` is used), or of other simulation protocols.

6.7 Harmonic wall restraints

The `harmonicWalls {...}` bias is closely related to the harmonic bias (see 6.5), with the following two differences: (i) instead of a center a *lower wall* and/or an *upper wall* are defined, outside of which the bias

implements a half-harmonic potential;

$$V(\xi) = \begin{cases} \frac{1}{2}k \left(\frac{\xi - \xi_{\text{upper}}}{w_{\xi}} \right)^2 & \text{if } \xi > \xi_{\text{upper}} \\ 0 & \text{if } \xi_{\text{lower}} \leq \xi \leq \xi_{\text{upper}} \\ \frac{1}{2}k \left(\frac{\xi - \xi_{\text{lower}}}{w_{\xi}} \right)^2 & \text{if } \xi < \xi_{\text{lower}} \end{cases} \quad (32)$$

where ξ_{lower} and ξ_{upper} are the lower and upper wall thresholds, respectively; (ii) because an interval between two walls is defined, only scalar variables can be used (but any number of variables can be defined, and the wall bias is intrinsically multi-dimensional).

Note: this bias replaces the keywords `lowerWall`, `lowerWallConstant`, `upperWall` and `upperWallConstant` defined in the `colvar` context. Those keywords are deprecated.

The `harmonicWalls` bias implements the following options:

- **Keyword** `name`: see definition of `name` (biasing and analysis methods)
- **Keyword** `colvars`: see definition of `colvars` (biasing and analysis methods)
- **Keyword** `outputEnergy`: see definition of `outputEnergy` (biasing and analysis methods)
- **Keyword** `writeTIPMF`: see definition of `writeTIPMF` (biasing and analysis methods)
- **Keyword** `writeTISamples`: see definition of `writeTISamples` (biasing and analysis methods)
- **Keyword** `stepZeroData`: see definition of `stepZeroData` (biasing and analysis methods)
- **Keyword** `lowerWalls` \langle Position of the lower wall(s) \rangle
Context: `colvar`
Acceptable values: Space-separated list of decimals
Description: Defines the values ξ_{lower} below which a confining restraint on the colvar is applied to each colvar ξ .
- **Keyword** `upperWalls` \langle Position of the upper wall(s) \rangle
Context: `colvar`
Acceptable values: Space-separated list of decimals
Description: Defines the values ξ_{upper} above which a confining restraint on the colvar is applied to each colvar ξ .
- **Keyword** `forceConstant`: see definition of `forceConstant` (Harmonic restraints)
- **Keyword** `lowerWallConstant` \langle Force constant for the lower wall \rangle
Context: `harmonicWalls`
Acceptable values: positive decimal
Default value: `forceConstant`
Description: When both sets of walls are defined (lower and upper), this keyword allows setting different force constants for them. As with `forceConstant`, the specified constant is divided internally by the square of the specific width of each variable (see also the equivalent keyword for the harmonic restraint, `forceConstant`). The force constant reported in the output as “*k*”, and used in the change of force constant scheme, is the geometric mean of `upperWallConstant` and `lowerWallConstant`.

- **Keyword** `upperWallConstant`: analogous to `lowerWallConstant`
- **Keyword** `targetForceConstant`: see definition of `targetForceConstant` (harmonic restraints)
- **Keyword** `targetForceConstant` \langle Change the force constant(s) towards this value \rangle
Context: `harmonicWalls`
Acceptable values: positive decimal
Description: This keyword allows changing either one or both of the wall force constants over time. In the case that `lowerWallConstant` and `upperWallConstant` have the same value, the behavior of this keyword is identical to the corresponding keyword in the harmonic restraint; otherwise, the change schedule is applied to the geometric mean of the two constant. When only one set of walls is defined (`lowerWall` or `upperWalls`), only the respective force constant is changed. **Note:** if only one of the two force constants is meant to change over time, it is possible to use two instances of `harmonicWalls`, and apply the changing schedule only to one of them.
- **Keyword** `targetNumSteps`: see definition of `targetNumSteps` (harmonic restraints)
- **Keyword** `targetForceExponent`: see definition of `targetForceExponent` (harmonic restraints)
- **Keyword** `targetEquilSteps`: see definition of `targetEquilSteps` (harmonic restraints)
- **Keyword** `targetNumStages`: see definition of `targetNumStages` (harmonic restraints)
- **Keyword** `lambdaSchedule`: see definition of `lambdaSchedule` (harmonic restraints)
- **Keyword** `outputAccumulatedWork`: see definition of `outputAccumulatedWork` (harmonic restraints)
- **Keyword** `bypassExtendedLagrangian` \langle Apply bias to actual colvars, bypassing extended coordinates \rangle
Context: `harmonicWalls`
Acceptable values: boolean
Default value: on
Description: This option behaves as `bypassExtendedLagrangian` for other biases, but **it defaults to on, unlike in the general case**. Thus, by default, the `harmonicWalls` bias applies to the actual colvars, so that the distribution of the colvar between the walls is unaffected by the bias, which then applies a flat-bottom potential *as a function of the colvar value*. This bias will affect the extended coordinate distribution near the walls. If `bypassExtendedLagrangian` is disabled, `harmonicWalls` applies a flat-bottom potential *as a function of the extended coordinate*. Conversely, this bias will then modify the distribution of the actual colvar value near the walls.

Example 1: harmonic walls for one variable with two different force constants.

```
harmonicWalls {
  name mywalls
  colvars dist
  lowerWalls 22.0
  upperWalls 38.0
  lowerWallConstant 2.0
```



```
upperWallConstant 10.0
}
```

Example 2: harmonic walls for two variables with a single force constant.

```
harmonicWalls {
  name mywalls
  colvars phi psi
  lowerWalls -180.0 0.0
  upperWalls 0.0 180.0
  forceConstant 5.0
}
```

6.8 Linear restraints

The linear keyword defines a linear potential:

$$V(\xi) = k \left(\frac{\xi - \xi_0}{w_\xi} \right) \quad (33)$$

whose force is simply given by the constant k/w_ξ itself:

$$f(\xi) = k/w_\xi \quad (34)$$

This type of bias is therefore most useful in situations where a *constant force* is desired. As all other restraints, it can be defined on one or more CVs, with each contribution added to the total potential and the parameters w_ξ determining the relative magnitude for each.

Example: A possible use case of the linear bias is mimicking a constant electric field acting on a specific particle, or the center of mass of many particles. In the following example, a linear restraint is applied on a distanceZ variable (4.2.2), generating a constant force parallel to the Z axis of magnitude 5 *energy unit/length unit*:

```
colvar {
  name z
  distanceZ {
    ...
  }
}

linear {
  colvars z
  forceConstant 5.0
  centers 0.0
}
```

Another useful application of a linear restraint is to enforce experimental constraints in a simulation, with a lower non-equilibrium work than e.g. harmonic restraints [30]. There is generally a unique strength of bias for each CV center, which means you must know the bias force constant specifically for the center of the CV. This force constant may be found by using experiment directed simulation described in section 6.9.

- **Keyword** name: see definition of **name** (biasing and analysis methods)
- **Keyword** colvars: see definition of **colvars** (biasing and analysis methods)
- **Keyword** outputEnergy: see definition of **outputEnergy** (biasing and analysis methods)
- **Keyword** forceConstant \langle Scaled force constant (*energy unit*) \rangle
Context: linear
Acceptable values: Decimal
Default value: 1.0
Description: This option defines a *scaled* force constant k for the linear bias. To ensure consistency for multidimensional restraints, it is divided internally by the specific width of each variable (which is 1 by default), so that all variables are effectively dimensionless and of commensurate size. *The values of k/w_ξ are always printed when the restraint is defined.*
- **Keyword** centers \langle Initial linear restraint centers \rangle
Context: linear
Acceptable values: space-separated list of colvar values
Description: These are analogous to the **centers** keyword of the harmonic restraint. Although they do not affect dynamics, they are here necessary to ensure a well-defined energy for the linear bias.
- **Keyword** writeTIPMF: see definition of **writeTIPMF** (biasing and analysis methods)
- **Keyword** writeTISamples: see definition of **writeTISamples** (biasing and analysis methods)
- **Keyword** targetForceConstant: see definition of **targetForceConstant** (Harmonic restraints)
- **Keyword** targetNumSteps: see definition of **targetNumSteps** (Harmonic restraints)
- **Keyword** targetForceExponent: see definition of **targetForceExponent** (Harmonic restraints)
- **Keyword** targetEquilSteps: see definition of **targetEquilSteps** (Harmonic restraints)
- **Keyword** targetNumStages: see definition of **targetNumStages** (Harmonic restraints)
- **Keyword** lambdaSchedule: see definition of **lambdaSchedule** (Harmonic restraints)
- **Keyword** outputAccumulatedWork: see definition of **outputAccumulatedWork** (Harmonic restraints)

6.9 Adaptive Linear Bias/Experiment Directed Simulation

Experiment directed simulation applies a linear bias with a changing force constant. Please cite White and Voth [31] when using this feature. As opposed to that reference, the force constant here is scaled by the width corresponding to the biased colvar. In White and Voth, each force constant is scaled by the colvars set center. The bias converges to a linear bias, after which it will be the minimal possible bias. You may also stop the simulation, take the median of the force constants (ForceConst) found in the colvars trajectory

file, and then apply a linear bias with that constant. All the notes about units described in sections 6.8 and 6.5 apply here as well. **This is not a valid simulation of any particular statistical ensemble and is only an optimization algorithm until the bias has converged.**

- **Keyword name:** see definition of **name** (biasing and analysis methods)
- **Keyword colvars:** see definition of **colvars** (biasing and analysis methods)
- **Keyword centers** \langle Collective variable centers \rangle
Context: alb
Acceptable values: space-separated list of colvar values
Description: The desired center (equilibrium values) which will be sought during the adaptive linear biasing. The number of values must be the number of requested colvars. Each value is a decimal number if the corresponding colvar returns a scalar, a “(x, y, z)” triplet if it returns a unit vector or a vector, and a “q0, q1, q2, q3)” quadruplet if it returns a rotational quaternion. If a colvar has periodicities or symmetries, its closest image to the restraint center is considered when calculating the linear potential.
- **Keyword updateFrequency** \langle The duration of updates \rangle
Context: alb
Acceptable values: An integer
Description: This is, N , the number of simulation steps to use for each update to the bias. This determines how long the system requires to equilibrate after a change in force constant ($N/2$), how long statistics are collected for an iteration ($N/2$), and how quickly energy is added to the system (at most, $A/2N$, where A is the forceRange). Until the force constant has converged, the method as described is an optimization procedure and not an integration of a particular statistical ensemble. It is important that each step should be uncorrelated from the last so that iterations are independent. Therefore, N should be at least twice the autocorrelation time of the collective variable. The system should also be able to dissipate energy as fast as $N/2$, which can be done by adjusting thermostat parameters. Practically, N has been tested successfully at significantly shorter than the autocorrelation time of the collective variables being biased and still converge correctly.
- **Keyword forceRange** \langle The expected range of the force constant in units of energy \rangle
Context: alb
Acceptable values: A space-separated list of decimal numbers
Default value: $3 k_b T$
Description: This is largest magnitude of the force constant which one expects. If this parameter is too low, the simulation will not converge. If it is too high the simulation will waste time exploring values that are too large. A value of $3 k_b T$ has worked well in the systems presented as a first choice. This parameter is dynamically adjusted over the course of a simulation. The benefit is that a bad guess for the forceRange can be corrected. However, this can lead to large amounts of energy being added over time to the system. To prevent this dynamic update, add **hardForceRange yes** as a parameter
- **Keyword rateMax** \langle The maximum rate of change of force constant \rangle
Context: alb
Acceptable values: A list of space-separated real numbers
Description: This optional parameter controls how much energy is added to the system from this bias. Tuning this separately from the updateFrequency and forceRange can allow for large bias changes but with a low rateMax prevents large energy changes that can lead to instability in the simulation.

6.10 Multidimensional histograms

The histogram feature is used to record the distribution of a set of collective variables in the form of a N-dimensional histogram. Defining such a histogram is generally useful for analysis purposes, but it has no effect on the simulation.

Example 1: the two-dimensional histogram of a distance and an angle can be generated using the configuration below. The histogram code requires that each variable is a scalar number that is confined within a pre-defined interval. The interval's boundaries may be specified by hand (e.g. through `lowerBoundary` and `upperBoundary` in the variable definition), or auto-detected based on the type of function. In this example, the lower boundary for the distance variable “r” is automatically set to zero, and interval for the three-body angle “theta” is 0° and 180° : however, that an upper boundary for the distance “r” still needs to be specified manually. The grid spacings for the two variables are 0.2 length unit and 3.0° , respectively.

```
colvar {
  name r
  width 0.2
  upperBoundary 20.0
  distance { ... }
}

colvar {
  name theta
  width 3.0
  dihedral { ... }
}

histogram {
  name hist2d
  colvars r theta
}
```

Example 2: This example is similar to the previous one, but with the important difference that the parameters for the histogram's grid are defined *explicitly* for this histogram instance. Therefore, this histogram's grid may differ from the one defined from parameters embedded in the `colvar { ... }` block (for example, narrower intervals and finer grid spacings may be selected).

```
colvar {
  name r
  upperBoundary 20.0
  distance { ... }
}

colvar {
  name theta
  dihedral { ... }
```

```

}

histogram {
  name hist2d
  colvars r theta
  histogramGrid {
    widths 0.1 1.0
    lowerBoundaries 2.0 30.0
    upperBoundaries 10.0 90.0
  }
}

```

The standard keywords below are used to control the histogram's computation and to select the variables that are sampled. See also [6.10.1](#) for keywords used to define the grid, [6.10.2](#) for output parameters and [6.10.3](#) for more advanced keywords.

- **Keyword** `name`: see definition of `name` (biasing and analysis methods)
- **Keyword** `colvars`: see definition of `colvars` (biasing and analysis methods)
- **Keyword** `stepZeroData`: see definition of `stepZeroData` (biasing and analysis methods)

6.10.1 Defining grids for multidimensional histograms

Grid parameters for the histogram may be provided at the level of the individual variables, or via a dedicated configuration block `histogramGrid { . . . }` inside the configuration of this histogram. The options supported *inside this block* are:

- **Keyword** `lowerBoundaries` \langle Lower boundaries of the grid \rangle
Context: `histogramGrid`
Acceptable values: list of space-separated decimals
Description: This option defines the lower boundaries of the grid, overriding any values defined by the `lowerBoundary` keyword of each colvar. Note that when `gatherVectorColvars` is on, each vector variable is automatically treated as a scalar, and a single value should be provided for it.
- **Keyword** `upperBoundaries`: analogous to `lowerBoundaries`
- **Keyword** `widths`: analogous to `lowerBoundaries`

6.10.2 Output options for multi-dimensional histograms

The accumulated histogram is written in the Colvars state file, allowing for its accumulation across continued runs. Additionally, the following files are written depending on the histogram's dimensionality:

- **Keyword** `outputFreq`: see definition of `outputFreq` (biasing and analysis methods)
- **Keyword** `outputFile` ⟨Write the histogram to a file⟩
Context: histogram
Acceptable values: UNIX filename
Default value: `output.<name>.dat`
Description: Name of the file containing histogram data (multicolumn format, see 3.7.4), which is written every `outputFreq` steps. If `outputFile` is set to none, the file is not written.
- **Keyword** `outputFileDX` ⟨Write the histogram to a file⟩
Context: histogram
Acceptable values: UNIX filename
Default value: `output.<name>.dx`
Description: Name of the file containing histogram data (OpenDX format), which is written every `outputFreq` steps. For the special case of 3 variables, VMD may be used to visualize this file. This file is written by default if the dimension is 3 or more (you cannot visualize it easily for dimensions 4 and above, but the DX format is still more compact than the multicolumn format). If `outputFileDX` is set to none, the file is not written.

6.10.3 Histogramming vector variables

- **Keyword** `gatherVectorColvars` ⟨Treat vector variables as multiple observations of a scalar variable?⟩
Context: histogram
Acceptable values: UNIX filename
Default value: off
Description: When this is set to on, the components of a multi-dimensional colvar (e.g. one based on cartesian, distancePairs, or a vector of scalar numbers given by `scriptedFunction`) are treated as multiple observations of a scalar variable. This results in the histogram being accumulated multiple times for each simulation step). When multiple vector variables are included in histogram, these must have the same length because their components are accumulated together. For example, if ξ , λ and τ are three variables of dimensions 5, 5 and 1, respectively, for each iteration 5 triplets (ξ_i, λ_i, τ) ($i = 1, \dots, 5$) are accumulated into a 3-dimensional histogram.
- **Keyword** `weights` ⟨Relative contributions of each vector component to the histogram⟩
Context: histogram
Acceptable values: list of space-separated decimals
Default value: all weights equal to 1
Description: When `gatherVectorColvars` is on, the components of each multi-dimensional colvar are accumulated with a different weight. For example, if x and y are two distinct cartesian variables defined on the same group of atoms, the corresponding 2D histogram can be weighted on a per-atom basis in the definition of histogram.

As with any other biasing and analysis method, when a histogram is applied to an extended-system colvar (4.18), it accesses the value of the extended coordinate rather than that of the actual colvar. This can be overridden by enabling the `bypassExtendedLagrangian` option. A *joint histogram* of the actual colvar and the extended coordinate may be collected by specifying the colvar name twice in a row in the `colvars`

parameter (e.g. `colvars myColvar myColvar`): the first instance will be understood as the actual colvar, and the second, as the extended coordinate.

- **Keyword** `bypassExtendedLagrangian`: see definition of `bypassExtendedLagrangian` (biasing and analysis methods)

6.11 Probability distribution-restraints

The `histogramRestraint` bias implements a continuous potential of many variables (or of a single high-dimensional variable) aimed at reproducing a one-dimensional statistical distribution that is provided by the user. The M variables (ξ_1, \dots, ξ_M) are interpreted as multiple observations of a random variable ξ with unknown probability distribution. The potential is minimized when the histogram $h(\xi)$, estimated as a sum of Gaussian functions centered at (ξ_1, \dots, ξ_M) , is equal to the reference histogram $h_0(\xi)$:

$$V(\xi_1, \dots, \xi_M) = \frac{1}{2}k \int (h(\xi) - h_0(\xi))^2 d\xi \quad (35)$$

$$h(\xi) = \frac{1}{M\sqrt{2\pi\sigma^2}} \sum_{i=1}^M \exp\left(-\frac{(\xi - \xi_i)^2}{2\sigma^2}\right) \quad (36)$$

When used in combination with a `distancePairs` multi-dimensional variable, this bias implements the refinement algorithm against ESR/DEER experiments published by Shen *et al* [32].

This bias behaves similarly to the `histogram` bias with the `gatherVectorColvars` option, with the important difference that *all* variables are gathered, resulting in a one-dimensional histogram. Future versions will include support for multi-dimensional histograms.

The list of options is as follows:

- **Keyword** `name`: see definition of `name` (biasing and analysis methods)
- **Keyword** `colvars`: see definition of `colvars` (biasing and analysis methods)
- **Keyword** `outputEnergy`: see definition of `outputEnergy` (biasing and analysis methods)
- **Keyword** `lowerBoundary` \langle Lower boundary of the colvar grid \rangle
Context: `histogramRestraint`
Acceptable values: decimal
Description: Defines the lowest end of the interval where the reference distribution $h_0(\xi)$ is defined. Exactly one value must be provided, because only one-dimensional histograms are supported by the current version.
- **Keyword** `upperBoundary`: analogous to `lowerBoundary`
- **Keyword** `width` \langle Width of the colvar grid \rangle
Context: `histogramRestraint`
Acceptable values: positive decimal
Description: Defines the spacing of the grid where the reference distribution $h_0(\xi)$ is defined.

- **Keyword** gaussianSigma \langle Standard deviation of the approximating Gaussian \rangle
Context: histogramRestraint
Acceptable values: positive decimal
Default value: $2 \times \text{width}$
Description: Defines the parameter σ in eq. 36.
- **Keyword** forceConstant \langle Force constant (*energy unit*) \rangle
Context: histogramRestraint
Acceptable values: positive decimal
Default value: 1.0
Description: Defines the parameter k in eq. 35.
- **Keyword** refHistogram \langle Reference histogram $h_0(\xi)$ \rangle
Context: histogramRestraint
Acceptable values: space-separated list of M positive decimals
Description: Provides the values of $h_0(\xi)$ consecutively. The mid-point convention is used, i.e. the first point that should be included is for $\xi = \text{lowerBoundary} + \text{width}/2$. If the integral of $h_0(\xi)$ is not normalized to 1, $h_0(\xi)$ is rescaled automatically before use.
- **Keyword** refHistogramFile \langle Reference histogram $h_0(\xi)$ \rangle
Context: histogramRestraint
Acceptable values: UNIX file name
Description: Provides the values of $h_0(\xi)$ as contents of the corresponding file (mutually exclusive with refHistogram). The format is that of a text file, with each line containing the space-separated values of ξ and $h_0(\xi)$. The same numerical conventions as refHistogram are used.
- **Keyword** writeHistogram \langle Periodically write the instantaneous histogram $h(\xi)$ \rangle
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: If on, the histogram $h(\xi)$ is written every colvarsRestartFrequency steps to a file with the name *output*.<name>.hist.dat This is useful to diagnose the convergence of $h(\xi)$ against $h_0(\xi)$.

7 Syntax changes from older versions

The following is a list of syntax changes in Colvars since its first release. Many of the older keywords are still recognized by the current code, thanks to specific compatibility code. *This is not a list of new features:* its primary purpose is to make you aware of those improvements that affect the use of old configuration files with new versions of the code.

Note: if you are using any of the NAMD and VMD tutorials:

<https://www.ks.uiuc.edu/Training/Tutorials/>

please be aware that *several of these tutorials are not actively maintained*: for those cases, this list will help you reconcile any inconsistencies.

- **Colvars version 2016-06-09 or later (LAMMPS version 15Sep2016 or later).**
The legacy keyword `refPositionsGroup` has been renamed `fittingGroup` for clarity (the legacy version is still supported).
- **Colvars version 2016-08-10 or later (LAMMPS version 15Sep2016 or later).**
“System forces” have been replaced by “total forces” (see for example `outputTotalForce`). See the following page for more information:
<https://colvars.github.io/README-totalforce.html>
- **Colvars version 2017-01-09 or later (LAMMPS version 10Mar2017 or later).**
A new type of restraint, `harmonicWalls` (see 6.7), replaces and improves upon the legacy keywords `lowerWall` and `upperWall`: these are still supported as short-hands.
- **Colvars version 2018-11-15 or later (LAMMPS version 23Nov2018 or later).**
The global analysis keyword has been discontinued: specific analysis tasks are controlled directly by the keywords `corrFunc` and `runAve`, which continue to remain off by default.
- **Colvars version 2020-02-25 or later.**
The parameter `hillWidth`, expressing the Gaussian width 2σ in relative units (number of grid points), does not have a default value any more. A new alternative parameter `gaussianSigmas` allows setting the σ parameters explicitly for each variable if needed.
Furthermore, to facilitate the use of other analysis tools such as for example `sum_hills`:
https://www.plumed.org/doc-v2.6/user-doc/html/sum_hills.html
the format of the file written by `writeHillsTrajectory` has also been changed to use σ instead of 2σ . This change does not affect how the biasing potential is written in the state file, or the simulated trajectory.
- **Colvars version 2020-02-25 or later.**
The legacy keywords `lowerWall` and `upperWall` of a colvar definition block do not have default values any longer, and need to be set explicitly, preferably as part of the `harmonicWalls` restraint. When using an ABF bias, it is recommended to set the two walls equal to `lowerBoundary` and `upperBoundary`, respectively. When using a metadynamics bias, it is recommended to set the two walls strictly *within* `lowerBoundary` and `upperBoundary`; see 6.4.1 for details.
- **Colvars version 2020-11-09 or later.**
The legacy keyword `disableForces` for atom groups is now deprecated and will be discontinued in

a future release. Atom groups now have an automated way to save computation if forces are not used, and enabling this option otherwise would lead to incorrect behavior.

Up-to-date documentation can always be accessed at:

<https://colvars.github.io/colvars-refman-lammps/colvars-refman-lammps.html>

8 Compilation notes

The Colvars module is typically built using the recipes of each supported software package: for this reason, no installation instructions are needed, and the vast majority of the features described in this manual are supported in the most common builds of each package.

This section lists the few cases where the choice of compilation settings affects the availability features in the Colvars module.

- Scripting commands using the Tcl language (<https://www.tcl.tk>) are supported in VMD and NAMD. All precompiled builds of each code include Tcl, and it is highly recommended to enable Tcl support in any custom build, using precompiled Tcl libraries from the UIUC website.
- The Lepton library (<https://simtk.org/projects/lepton>) used to implement the `customFunction` feature is currently included only in NAMD (always on) and in LAMMPS (on by default). For VMD, a [patch](#) that allows to link Lepton is available. There is no mechanism to link Lepton to GROMACS at this moment.
- Some features require compilation using the C++11 language standard, which is either supported (VMD) or required (GROMACS, LAMMPS, NAMD) by the engines. However, for historic reasons many of the VMD official builds do not conform to the C++11 standard yet, which results in restricted functionality. Detailed and up-to-date information can be found at: <https://colvars.github.io/README-c++11.html>

Index

NeuralNetwork

- layer*i*_BiasesFile, 43
- layer*i*_WeightsFile, 42
- layer*i*_activation, 43
- layer*i*_custom_activation, 43
- output_component, 42

abf

- CZARestimator, 69
- applyBias, 66
- colvars, 64
- fullSamples, 64
- hideJacobian, 65
- historyFreq, 65
- inputPrefix, 65
- integrateMaxIterations, 67
- integrateTol, 67
- integrate, 67
- maxForce, 65
- name, 64
- outputEnergy, 64
- outputFreq, 64
- stepZeroData, 64
- updateBias, 66
- writeCZARwindowFile, 69

alb

- centers, 88
- colvars, 88
- forceRange, 88
- name, 88
- rateMax, 88
- updateFrequency, 88

angle

- forceNoPBC, 20
- group1, 20
- group2, 20
- group3, 20
- oneSiteTotalForce, 20

angle, dipoleAngle, dihedral

- oneSiteTotalForce, 16

aspathCV and azpathCV

- lambda, 40
- pathFile, 40
- weights, 40

cartesian

- atoms, 34

colvar

- corrFuncLength, 53
- corrFuncNormalize, 53
- corrFuncOffset, 54
- corrFuncOutputFile, 54
- corrFuncStride, 54
- corrFuncType, 53
- corrFuncWithColvar, 53
- corrFunc, 53
- customFunctionType, 47
- customFunction, 47
- expandBoundaries, 50
- extendedFluctuation, 51
- extendedLagrangian, 51
- extendedLangevinDamping, 52
- extendedTemp, 52
- extendedTimeConstant, 51
- hardLowerBoundary, 49
- hardUpperBoundary, 49
- lowerBoundary, 49
- lowerWalls, 84
- name, 14
- outputAppliedForce, 51
- outputEnergy, 50
- outputTotalForce, 50
- outputValue, 50
- outputVelocity, 50
- runAveLength, 54
- runAveOutputFile, 54
- runAveStride, 54
- runAve, 54
- subtractAppliedForce, 52
- timeStepFactor, 52
- upperBoundary, 49
- upperWalls, 84
- width, 49

coordNum

- cutoff3, 22
- cutoff, 22
- expDenom, 22
- expNumer, 22
- group1, 22
- group2CenterOnly, 23
- group2, 22
- pairListFrequency, 23

- tolerance, 23
- dihedral
 - forceNoPBC, 21
 - group1, 21
 - group2, 21
 - group3, 21
 - group4, 21
 - oneSiteTotalForce, 21
- dipoleAngle
 - forceNoPBC, 20
 - group1, 20
 - group2, 20
 - group3, 20
 - oneSiteTotalForce, 20
- dipoleMagnitude
 - atoms, 28
- distanceDir
 - forceNoPBC, 19
 - group1, 19
 - group2, 19
 - oneSiteTotalForce, 19
- distanceInv
 - exponent, 19
 - group1, 19
 - group2, 19
 - oneSiteTotalForce, 19
- distancePairs
 - forceNoPBC, 35
 - group1, 34
 - group2, 34
- distanceVec
 - forceNoPBC, 18
 - group1, 18
 - group2, 18
 - oneSiteTotalForce, 18
- distanceXY
 - axis, 18
 - forceNoPBC, 18
 - main, 18
 - ref2, 18
 - ref, 18
- distanceZ
 - axis, 17
 - forceNoPBC, 17
 - main, 17
 - oneSiteTotalForce, 17, 18
 - ref2, 17
 - ref, 17

- distanceZ, dihedral, spinAngle, custom colvars
 - wrapAround, 45
- distanceZ, custom colvars
 - period, 45
- distance
 - forceNoPBC, 16
 - group1, 16
 - group2, 16
- eigenvector
 - atoms, 27
 - differenceVector, 27
 - normalizeVector, 27
 - refPositionsFile, 27
 - refPositions, 27
 - vectorFile, 27
 - vector, 27
- eulerPhi
 - atoms, 33
 - refPositionsFile, 33
 - refPositions, 33
- eulerPsi
 - atoms, 34
 - refPositionsFile, 34
 - refPositions, 34
- eulerTheta
 - atoms, 33
 - refPositionsFile, 33
 - refPositions, 33
- gspathCV and gzpathCV
 - pathFile, 38
 - useSecondClosestFrame, 38
 - useThirdClosestFrame, 38
- gspath and gspath
 - fittingAtoms, 36
- gspath and gzpath
 - atoms, 35
 - refPositionsCol, 35
 - refPositionsFileN, 36
 - useSecondClosestFrame, 36
 - useThirdClosestFrame, 36
- gyration
 - atoms, 28
- gzpathCV
 - useZsquare, 38
- gzpath
 - useZsquare, 36
- hBond
 - acceptor, 24

- cutoff, 24
- donor, 24
- expDenom, 24
- expNumer, 24
- harmonicWalls
 - bypassExtendedLagrangian, 85
 - colvars, 84
 - forceConstant, 84
 - lambdaSchedule, 85
 - lowerWallConstant, 84
 - name, 84
 - outputAccumulatedWork, 85
 - outputEnergy, 84
 - stepZeroData, 84
 - targetEquilSteps, 85
 - targetForceConstant, 85
 - targetForceExponent, 85
 - targetNumStages, 85
 - targetNumSteps, 85
 - upperWallConstant, 85
 - writeTIPMF, 84
 - writeTISamples, 84
- harmonic
 - centers, 80
 - colvars, 80
 - forceConstant, 80
 - lambdaSchedule, 83
 - name, 80
 - outputAccumulatedWork, 83
 - outputCenters, 81
 - outputEnergy, 80
 - stepZeroData, 80
 - targetCenters, 81
 - targetEquilSteps, 83
 - targetForceConstant, 82
 - targetForceExponent, 82
 - targetNumStages, 82
 - targetNumSteps, 81
 - writeTIPMF, 80
 - writeTISamples, 80
- histogramGrid
 - lowerBoundaries, 90
 - upperBoundaries, 90
 - widths, 90
- histogramRestraint
 - colvars, 92
 - forceConstant, 93
 - gaussianSigma, 93
 - lowerBoundary, 92
 - name, 92
 - outputEnergy, 92
 - refHistogramFile, 93
 - refHistogram, 93
 - upperBoundary, 92
 - width, 92
- histogram
 - bypassExtendedLagrangian, 92
 - colvars, 90
 - gatherVectorColvars, 91
 - name, 90
 - outputFileDX, 91
 - outputFile, 91
 - outputFreq, 91
 - stepZeroData, 90
 - weights, 91
- inertiaZ
 - atoms, 29
 - axis, 29
- inertia
 - atoms, 28
- linear
 - centers, 87
 - colvars, 87
 - forceConstant, 87
 - lambdaSchedule, 87
 - name, 87
 - outputAccumulatedWork, 87
 - outputEnergy, 87
 - targetEquilSteps, 87
 - targetForceConstant, 87
 - targetForceExponent, 87
 - targetNumStages, 87
 - targetNumSteps, 87
 - writeTIPMF, 87
 - writeTISamples, 87
- metadynamics
 - biasTemperature, 77
 - colvars, 72
 - ebMetaEquilSteps, 76
 - ebMeta, 75
 - gaussianSigmas, 73
 - hillWeight, 72
 - hillWidth, 73
 - keepFreeEnergyFiles, 73
 - keepHills, 74
 - multipleReplicas, 78

- name, 72
- newHillFrequency, 73
- outputEnergy, 72
- outputFreq, 72
- rebinGrids, 74
- replicaID, 79
- replicaUpdateFrequency, 79
- replicasRegistry, 78
- stepZeroData, 72
- targetDistFile, 75
- targetDistMinVal, 76
- useGrids, 74
- wellTempered, 77
- writeFreeEnergyFile, 73
- writeHillsTrajectory, 74
- writeHistogram, 93
- writePartialFreeEnergyFile, 79
- writeTIPMF, 72
- writeTISamples, 72
- orientationAngle
 - atoms, 31
 - refPositionsFile, 31
 - refPositions, 31
- orientationProj
 - atoms, 31
 - refPositionsFile, 31
 - refPositions, 31
- orientation
 - atoms, 30
 - closestToQuaternion, 30
 - refPositionsFile, 30
 - refPositions, 30
- polarPhi
 - atoms, 21
- rmsd
 - atomPermutation, 25
 - atoms, 25
 - refPositionsFile, 25
 - refPositions, 25
- selfCoordNum
 - cutoff3, 24
 - cutoff, 24
 - expDenom, 24
 - expNumer, 24
 - group1, 23
 - pairListFrequency, 24
 - tolerance, 24
- spinAngle
 - atoms, 31
 - refPositionsFile, 31
 - refPositions, 31
- tilt
 - atoms, 32
 - axis, 32
 - refPositionsFile, 32
 - refPositions, 32
- any component
 - componentCoeff, 46
 - componentExp, 46
 - name, 44
 - scalable, 44
- atom group
 - atomNumbersRange, 56
 - atomNumbers, 55
 - atomsOfGroup, 56
 - centerToOrigin, 57
 - centerToReference, 57
 - dummyAtom, 56
 - enableFitGradients, 58
 - fittingGroup, 58
 - indexGroup, 56
 - name, 55
 - refPositionsFile, 58
 - refPositions, 57
 - rotateToReference, 57
- colvar bias
 - bypassExtendedLagrangian, 61
 - colvars, 60
 - name, 60
 - outputEnergy, 60
 - outputFreq, 60
 - scaledBiasingForceFactorsGrid, 61
 - scaledBiasingForce, 61
 - stepZeroData, 61
 - writeTIPMF, 62
 - writeTISamples, 62
- global
 - colvarsRestartFrequency, 7
 - colvarsTrajFrequency, 7
 - indexFile, 8
 - smp, 8
 - units, 5

Keyword of the fix colvars command

- input, 6
- output, 6
- seed, 6
- tstat, 6
- keyword of the fix colvars command
- unwrap, 6

References

- [1] Giacomo Fiorin, Michael L. Klein, and Jérôme Hénin. Using collective variables to drive molecular dynamics simulations. *Mol. Phys.*, 111(22-23):3345–3362, 2013. 4, 31, 73, 75, 80
- [2] Steve Plimpton. Fast parallel algorithms for short-range molecular dynamics. *J. Comp. Phys.*, 117(1):1–19, 1995. 7
- [3] M. Iannuzzi, A. Laio, and M. Parrinello. Efficient exploration of reactive potential energy surfaces using car-parrinello molecular dynamics. *Phys. Rev. Lett.*, 90(23):238302, 2003. 24, 53
- [4] E A Coutias, C Seok, and K A Dill. Using quaternions to calculate RMSD. *J. Comput. Chem.*, 25(15):1849–1857, 2004. 27, 31, 32, 59
- [5] Haohao Fu, Wensheng Cai, Jérôme Hénin, Benoît Roux, and Christophe Chipot. New coarse variables for the accurate determination of standard binding free energies. *J. Chem. Theory Comput.*, 13(11):5173–5178, 2017. 35
- [6] G. D. Leines and B. Ensing. Path finding on high-dimensional free energy landscapes. *Phys. Rev. Lett.*, 109:020601, 2012. 37
- [7] Davide Branduardi, Francesco Luigi Gervasio, and Michele Parrinello. From a to b in free energy space. *J Chem Phys*, 126(5):054103, 2007. 37, 42
- [8] F. Comitani L. Hovan and F. L. Gervasio. Defining an optimal metric for the path collective variables. *J. Chem. Theory Comput.*, 15:25–32, 2019. 42
- [9] Haochuan Chen, Han Liu, Heying Feng, Haohao Fu, Wensheng Cai, Xueguang Shao, and Christophe Chipot. Mlcv: Bridging Machine-Learning-Based Dimensionality Reduction and Free-Energy Calculation. *J. Chem. Inf. Model.*, 62(1):1–8, 2022. 44
- [10] Marco Jacopo Ferrarotti, Sandro Bottaro, Andrea Pérez-Villa, and Giovanni Bussi. Accurate multiple time step in biased molecular simulations. *Journal of chemical theory and computation*, 11:139–146, 2015. 54
- [11] Eric Darve, David Rodríguez-Gómez, and Andrew Pohorille. Adaptive biasing force method for scalar and vector free energy calculations. *J. Chem. Phys.*, 128(14):144120, 2008. 64
- [12] J. Hénin and C. Chipot. Overcoming free energy barriers using unconstrained molecular dynamics simulations. *J. Chem. Phys.*, 121:2904–2914, 2004. 64
- [13] Jérôme Hénin, Giacomo Fiorin, Christophe Chipot, and Michael L. Klein. Exploring multidimensional free energy landscapes using time-dependent biases on collective variables. *J. Chem. Theory Comput.*, 6(1):35–47, 2010. 64
- [14] A. Carter, E. G. Ciccotti, J. T. Hynes, and R. Kapral. Constrained reaction coordinate dynamics for the simulation of rare events. *Chem. Phys. Lett.*, 156:472–477, 1989. 65
- [15] M. J. Ruiz-Montero, D. Frenkel, and J. J. Brey. Efficient schemes to compute diffusive barrier crossing rates. *Mol. Phys.*, 90:925–941, 1997. 65
- [16] W. K. den Otter. Thermodynamic integration of the free energy along a reaction coordinate in cartesian coordinates. *J. Chem. Phys.*, 112:7283–7292, 2000. 65

- [17] Giovanni Ciccotti, Raymond Kapral, and Eric Vanden-Eijnden. Blue moon sampling, vectorial reaction coordinates, and unbiased constrained dynamics. *ChemPhysChem*, 6(9):1809–1814, 2005. 65
- [18] J. Hénin. Fast and accurate multidimensional free energy integration. *J. Chem. Theory Comput.*, 2021. 69, 70
- [19] Adrien Lesage, Tony Lelièvre, Gabriel Stoltz, and Jérôme Hénin. Smoothed biasing forces yield unbiased free energies with the extended-system adaptive biasing force method. *J. Phys. Chem. B*, 121(15):3676–3685, 2017. 70, 71
- [20] A. Laio and M. Parrinello. Escaping free-energy minima. *Proc. Natl. Acad. Sci. USA*, 99(20):12562–12566, 2002. 72
- [21] Helmut Grubmüller. Predicting slow structural transitions in macromolecular systems: Conformational flooding. *Phys. Rev. E*, 52(3):2893–2906, Sep 1995. 72
- [22] T. Huber, A. E. Torda, and W.F. van Gunsteren. Local elevation - A method for improving the searching properties of molecular-dynamics simulation. *Journal of Computer-Aided Molecular Design*, 8(6):695–708, DEC 1994. 72
- [23] G. Bussi, A. Laio, and M. Parrinello. Equilibrium free energies from nonequilibrium metadynamics. *Phys. Rev. Lett.*, 96(9):090601, 2006. 72
- [24] Fabrizio Marinelli, Fabio Pietrucci, Alessandro Laio, and Stefano Piana. A kinetic model of trp-cage folding from multiple biased molecular dynamics simulations. *PLOS Computational Biology*, 5(8):1–18, 2009. 72
- [25] Yanier Crespo, Fabrizio Marinelli, Fabio Pietrucci, and Alessandro Laio. Metadynamics convergence law in a multidimensional system. *Phys. Rev. E*, 81:055701, May 2010. 72
- [26] Fabrizio Marinelli and José D. Faraldo-Gómez. Ensemble-biased metadynamics: A molecular simulation method to sample experimental distributions. *Biophys. J.*, 108(12):2779–2782, 2015. 77
- [27] Alessandro Barducci, Giovanni Bussi, and Michele Parrinello. Well-tempered metadynamics: A smoothly converging and tunable free-energy method. *Phys. Rev. Lett.*, 100:020603, 2008. 79
- [28] P. Raiteri, A. Laio, F. L. Gervasio, C. Micheletti, and M. Parrinello. Efficient reconstruction of complex free energy landscapes by multiple walkers metadynamics. *J. Phys. Chem. B*, 110(8):3533–9, 2006. 79
- [29] Yuqing Deng and Benoît Roux. Computations of standard binding free energies with molecular dynamics simulations. *J. Phys. Chem. B*, 113(8):2234–2246, 2009. 84
- [30] Jed W. Pitera and John D. Chodera. On the use of experimental observations to bias simulated ensembles. *J. Chem. Theory Comput.*, 8:3445–3451, 2012. 88
- [31] Andrew D. White and Gregory A. Voth. Efficient and minimal method to bias molecular simulations with experimental data. *J. Chem. Theory Comput.*, 10(8):3023–3030, 2014. 89
- [32] Rong Shen, Wei Han, Giacomo Fiorin, Shahidul M Islam, Klaus Schulten, and Benoît Roux. Structural refinement of proteins by restrained molecular dynamics simulations with non-interacting molecular fragments. *PLoS Comput. Biol.*, 11(10):e1004368, 2015. 94